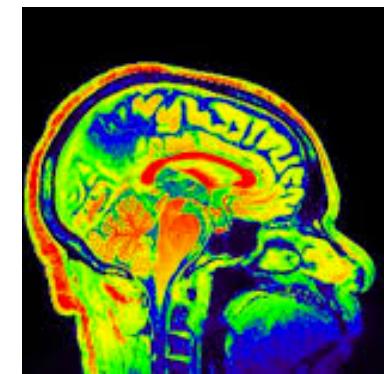
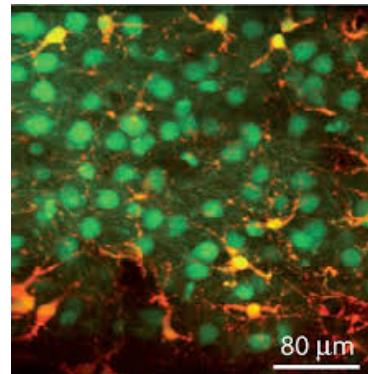


Statistical Models for Neural Data: from Regression / GLMs to Latent Variables

Jonathan Pillow

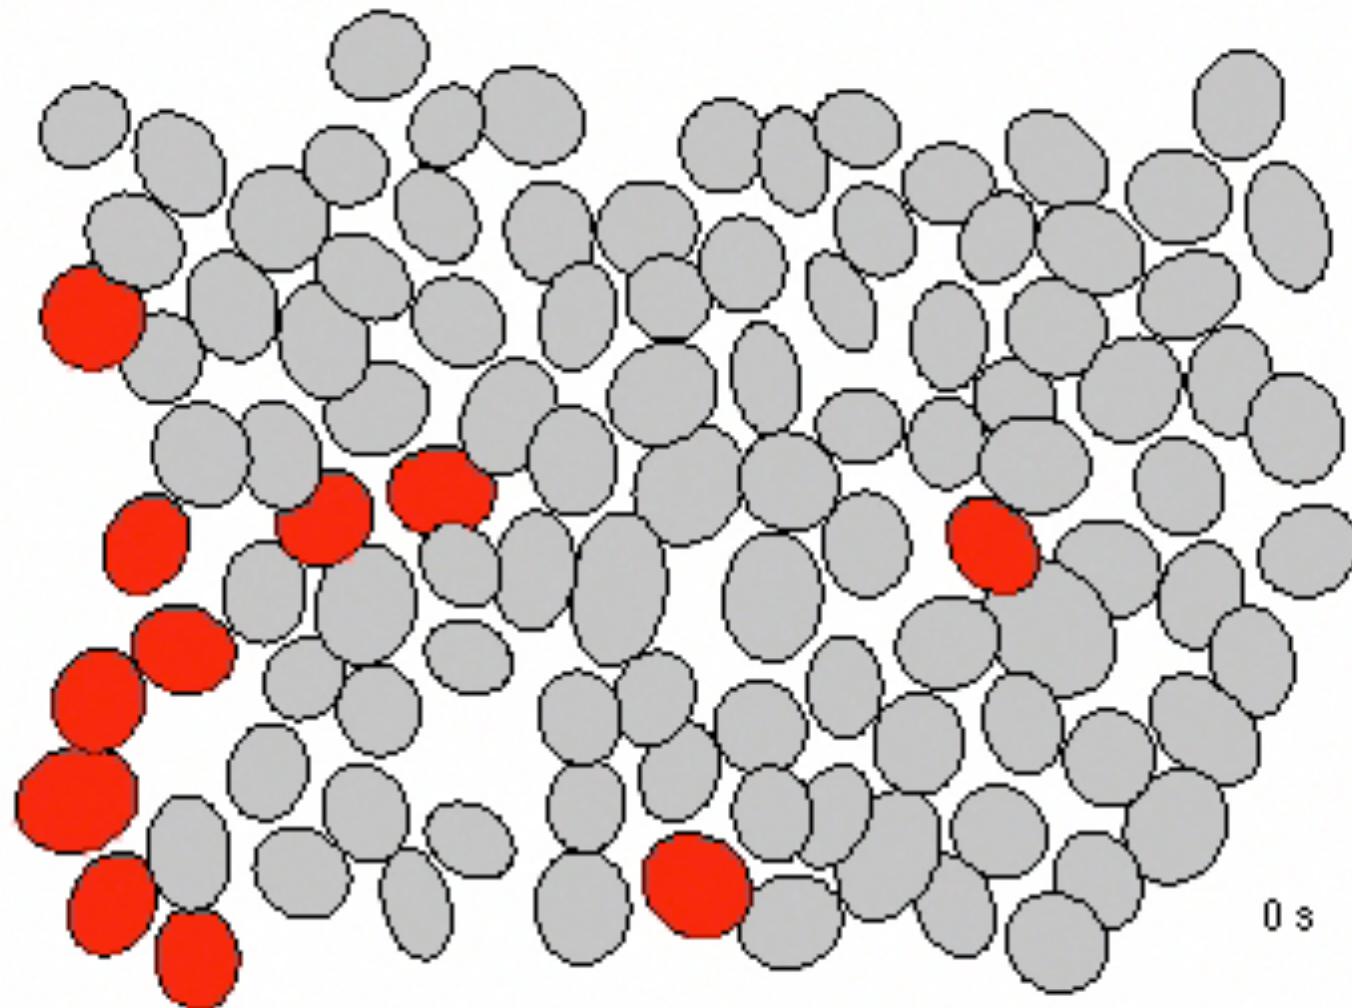
Princeton Neuroscience Institute



Tutorial
Cosyne 2018

Retinal responses to white noise

(ON parasol cells)



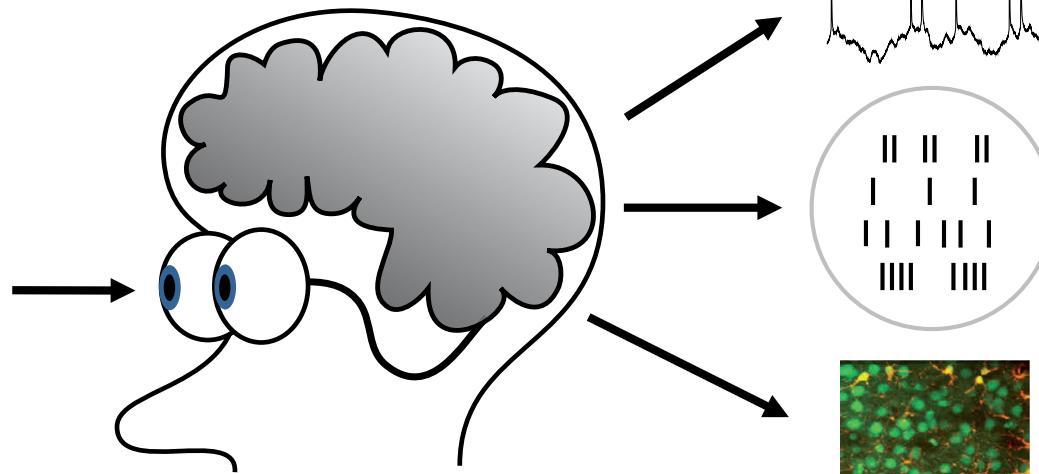
Shlens, Field, Gauthier, Greschner, Sher , Litke & Chichilnisky (2009).

neural coding problem



stimulus

x



membrane
potential

spikes

imaging

neural activity

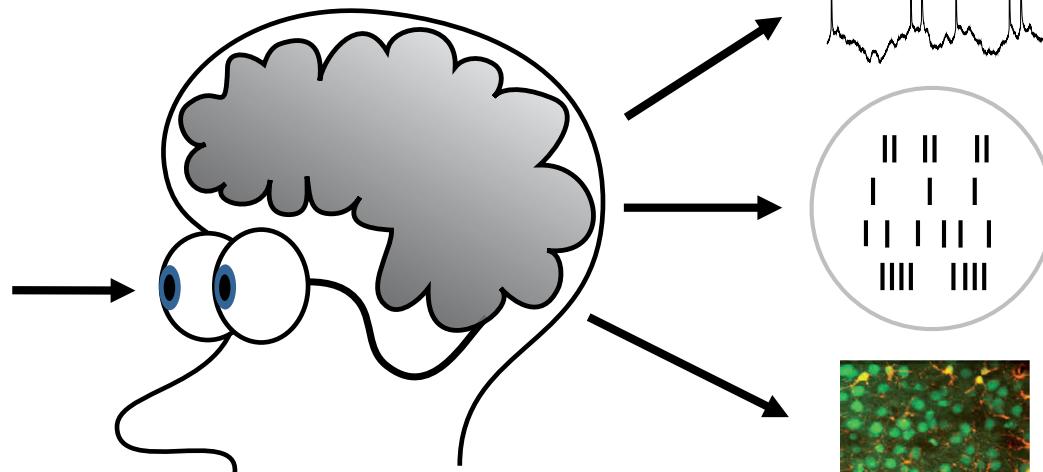
y

- How are stimuli and actions encoded in neural activity?
- What aspects of neural activity carry information?

neural coding problem



stimulus
 x



$P(y|x)$
encoding models

membrane
potential

spikes

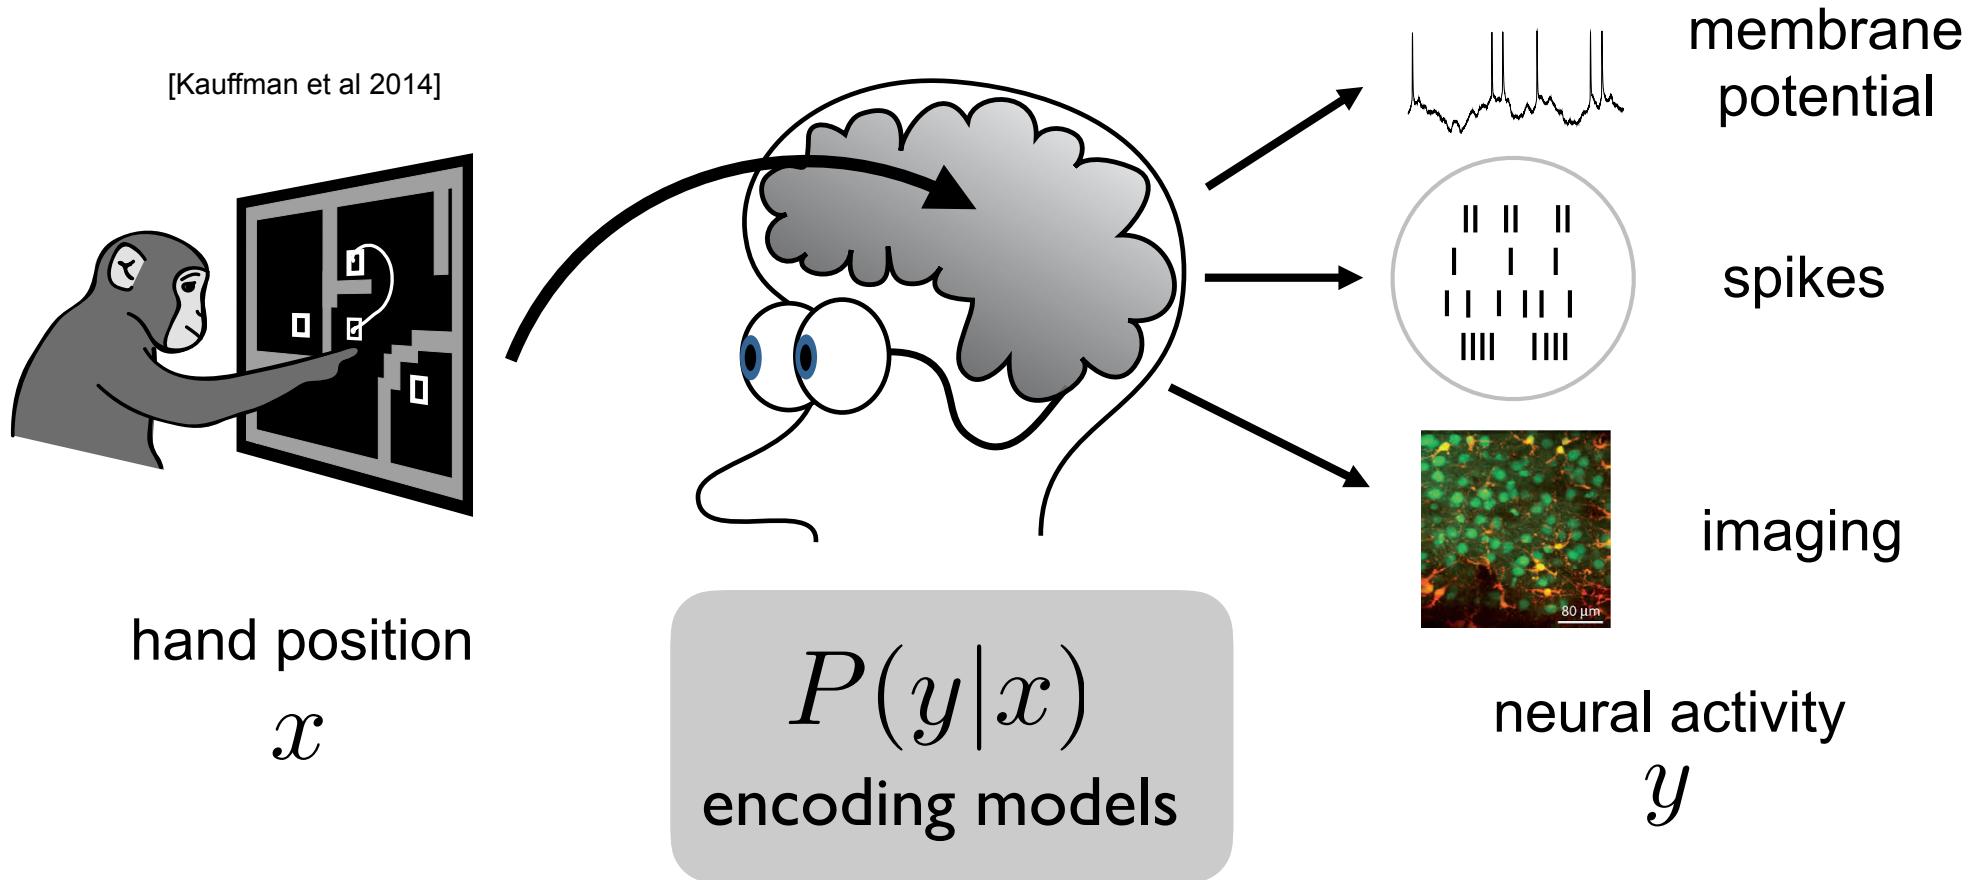
imaging

neural activity
 y

Approach:

- develop flexible statistical models of $P(y|x)$
- quantify information carried in neural responses

neural coding problem



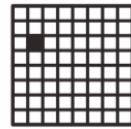
“regression models”

- not restricted to sensory variables

neural coding problem

[Hardcastle et al 2015]

Position (P)



Head direction



Speed (S)



Theta phase

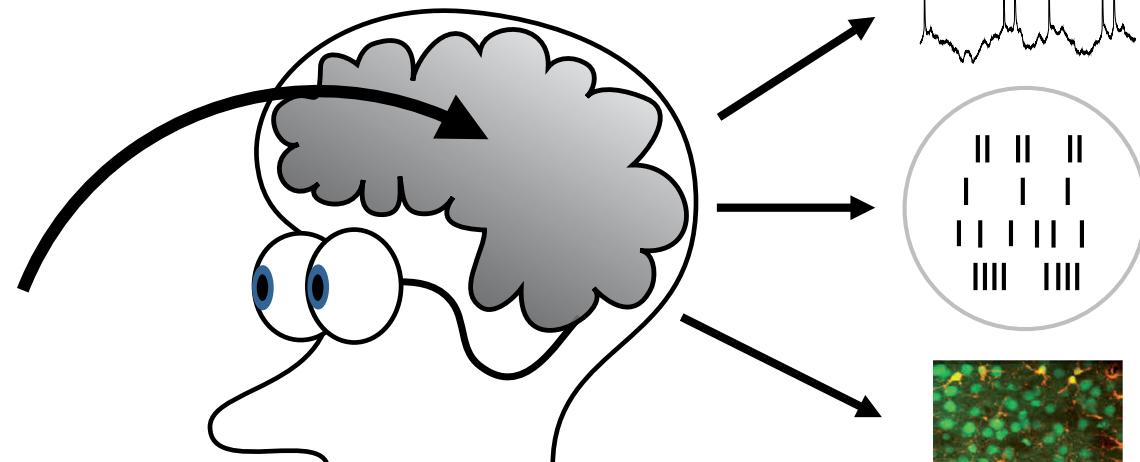


x

“external covariates”

“regression models”

- not restricted to sensory variables



$P(y|x)$
encoding models

membrane
potential

spikes

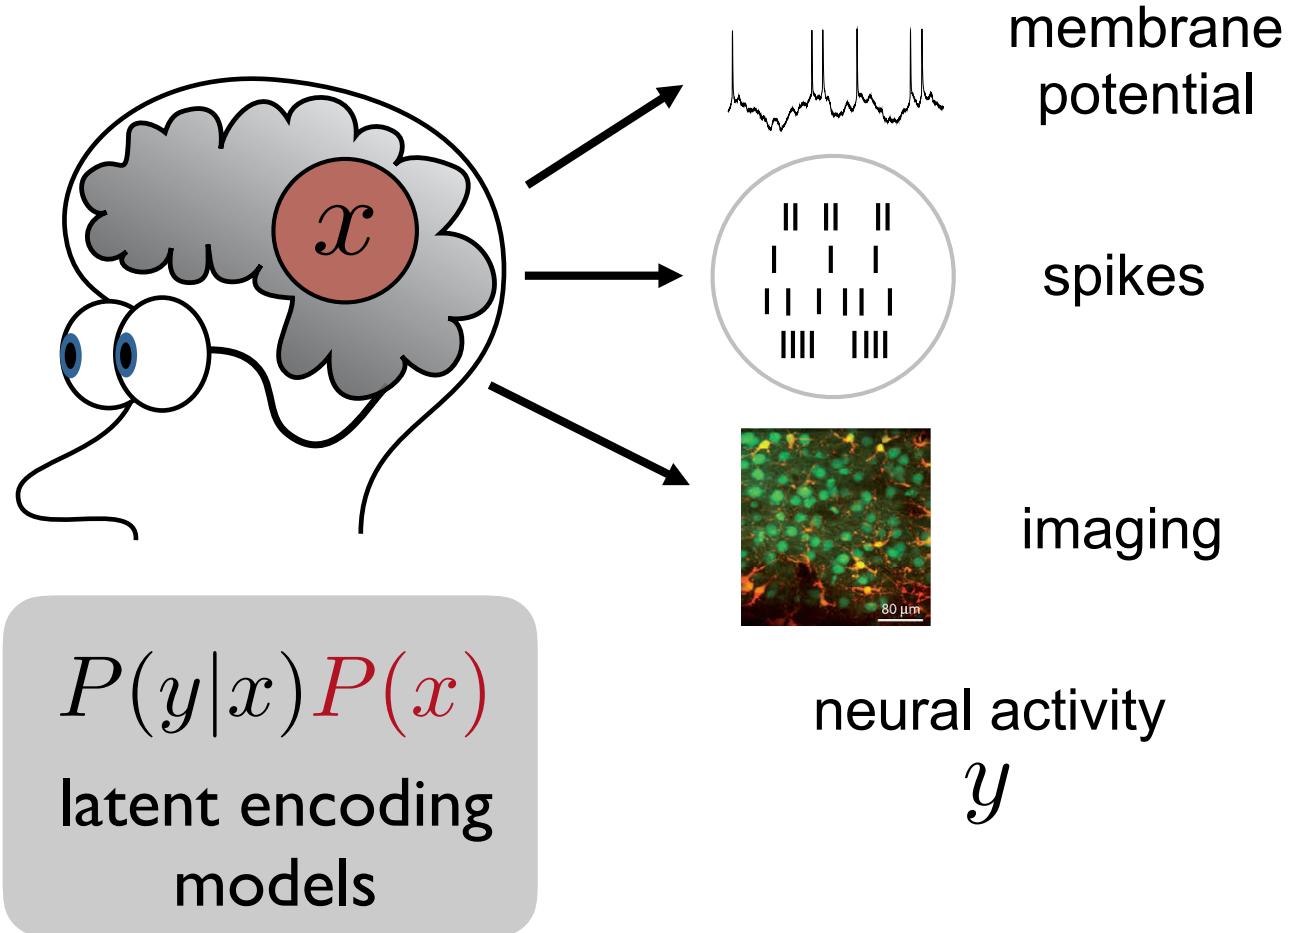
imaging

neural activity

y

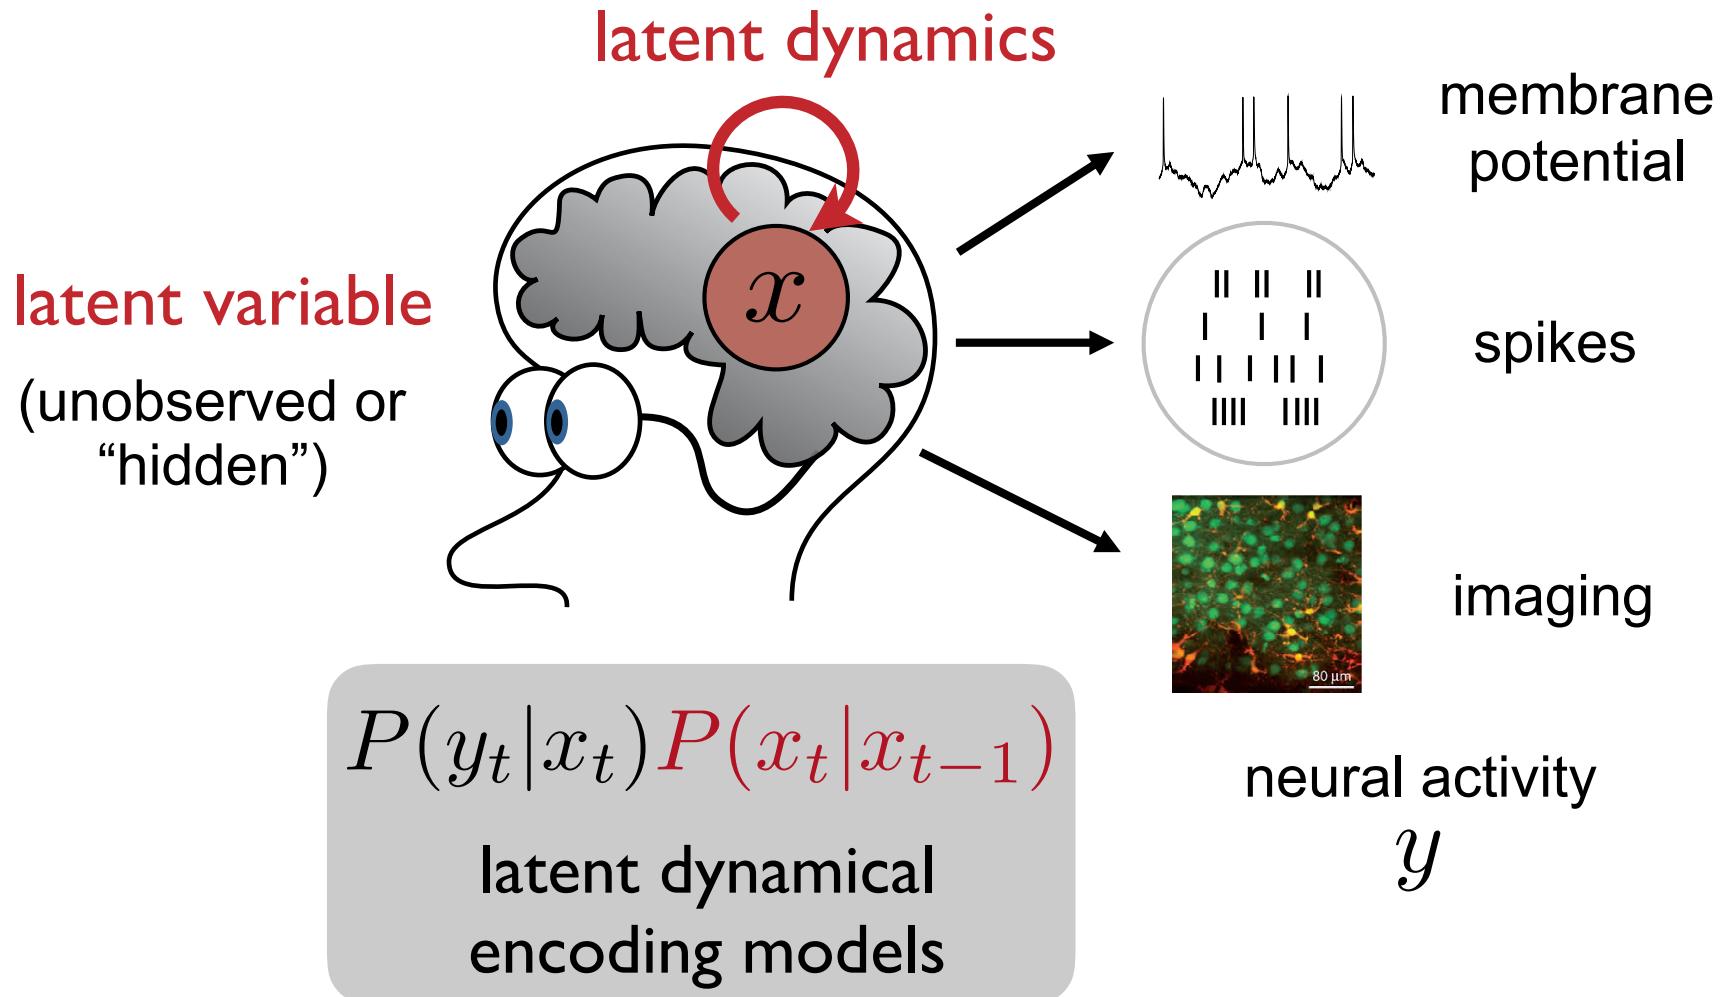
latent variable models

latent variable
(unobserved or
“hidden”)



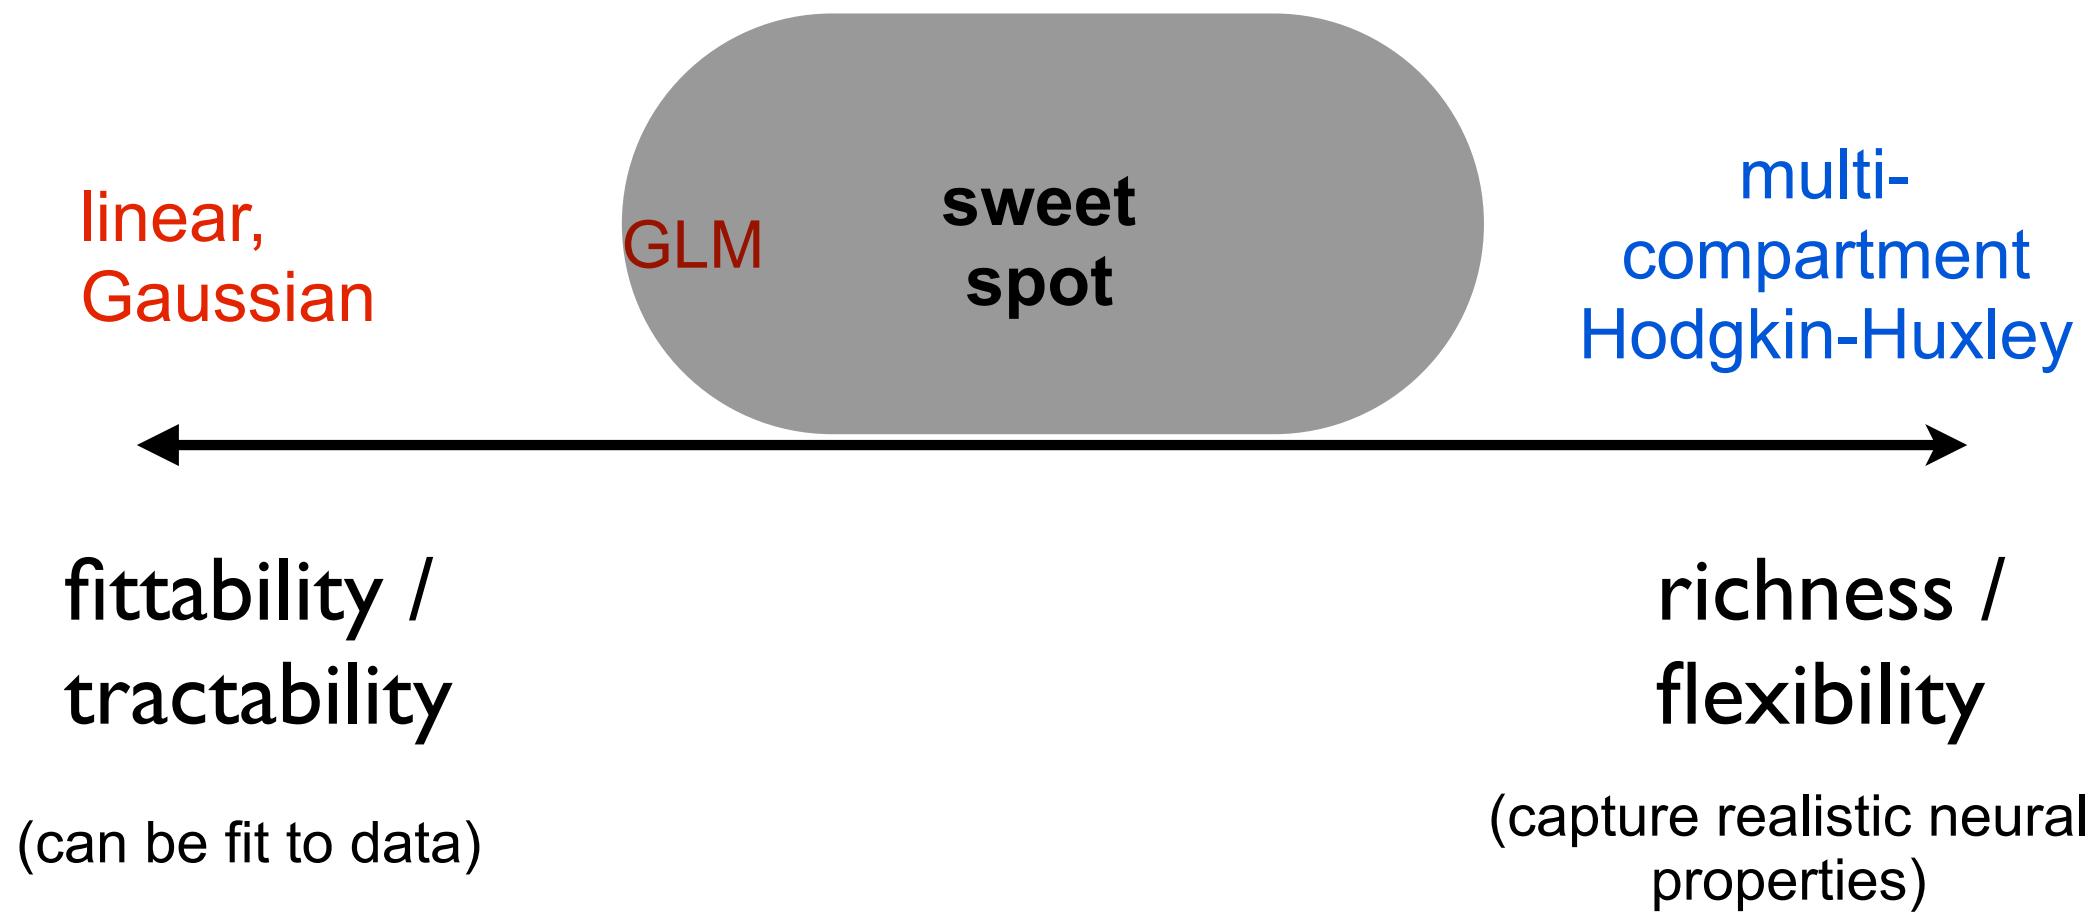
- capture hidden structure underlying neural activity
(eg. low-dimensional or discrete states)

latent variable models



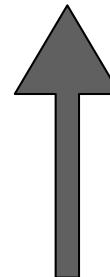
- capture hidden dynamics underlying neural activity

model desiderata



normative theories
(e.g. “efficient coding”)

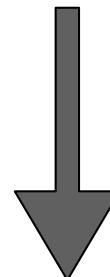
*Why does the code
take this form?*



**descriptive
statistical models**

anatomy,
biophysics

$P(y|x)$
What is the code?



How is it implemented?

Outline

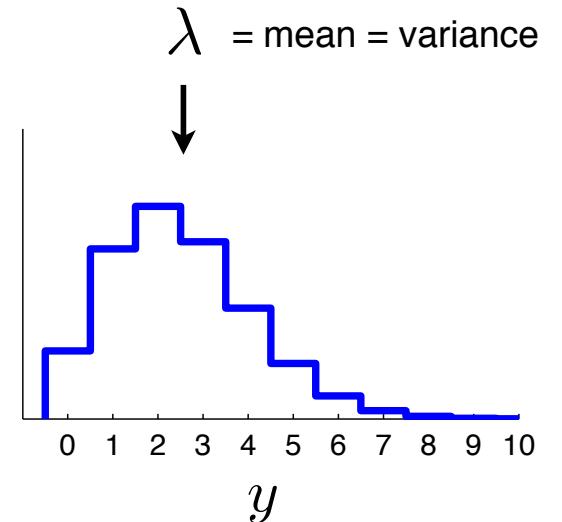
1. Spike count models & Maximum Likelihood
2. Spike train models (GLMs with spike history)
3. Multiple Spike Train Models (GLMs with coupling)
4. Regularization
5. Beyond GLM
6. Latent variable models

simple example #1: linear Poisson neuron

parameter stimulus

spike rate $\lambda = \theta x$

spike count $y \sim Poiss(\lambda)$

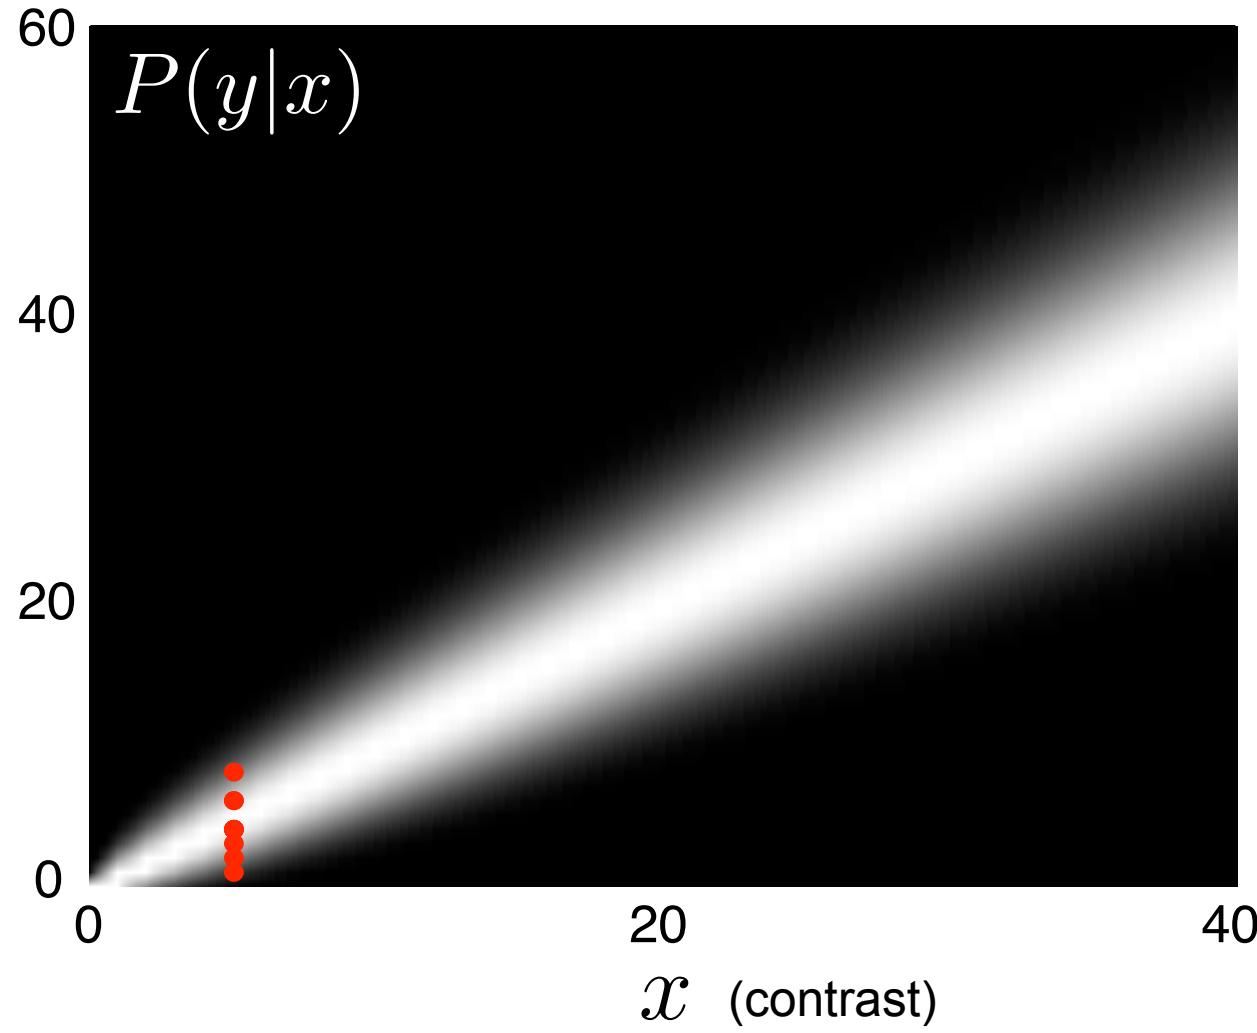


encoding model:

$$\begin{aligned} P(y|x, \theta) &= \frac{1}{y!} \lambda^y e^{-\lambda} \\ &= \frac{1}{y!} (\theta x)^y e^{-(\theta x)} \end{aligned}$$

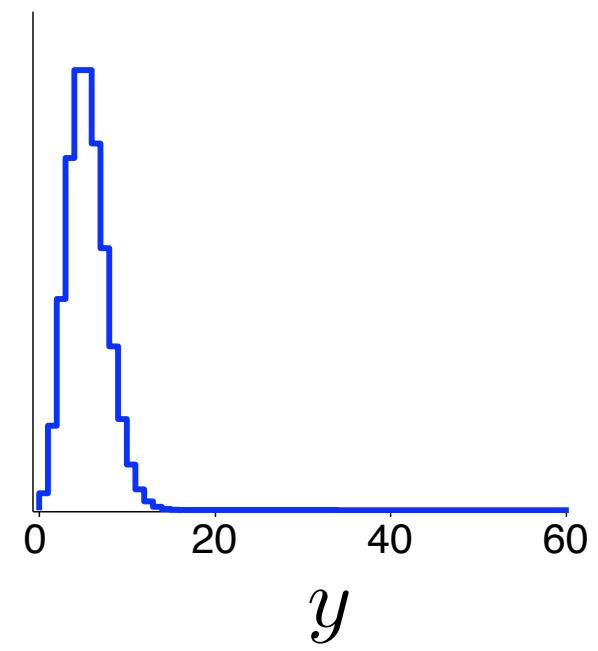
$$\text{mean}(y) = \theta x$$

$$\text{var}(y) = \theta x$$



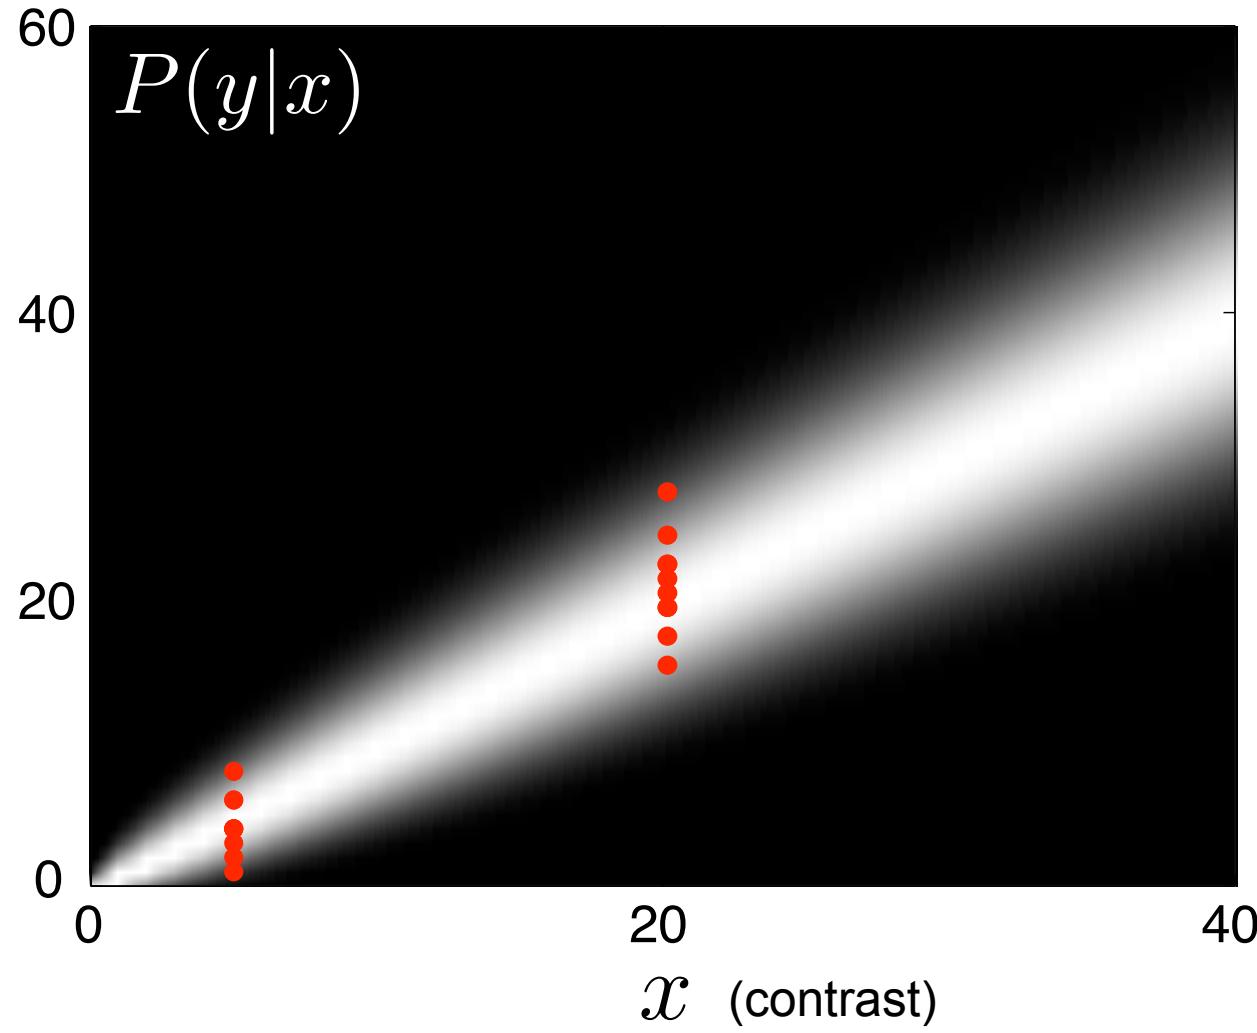
conditional distribution

$$p(y|x = 5)$$



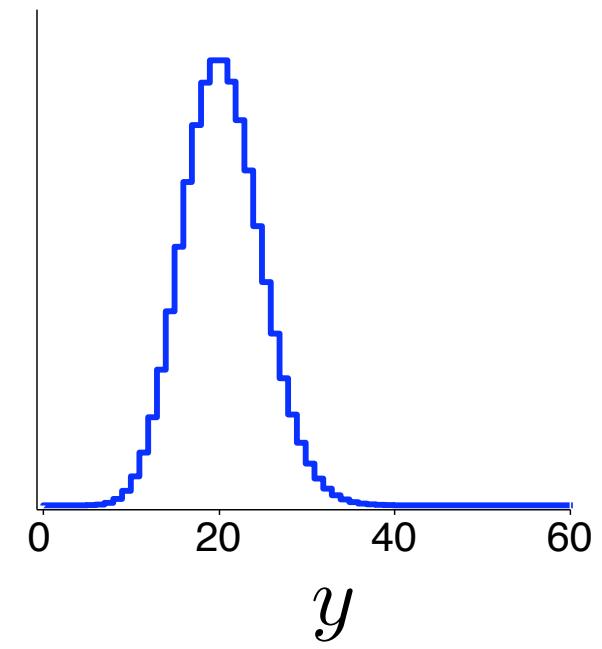
$$\text{mean}(y) = \theta x$$

$$\text{var}(y) = \theta x$$



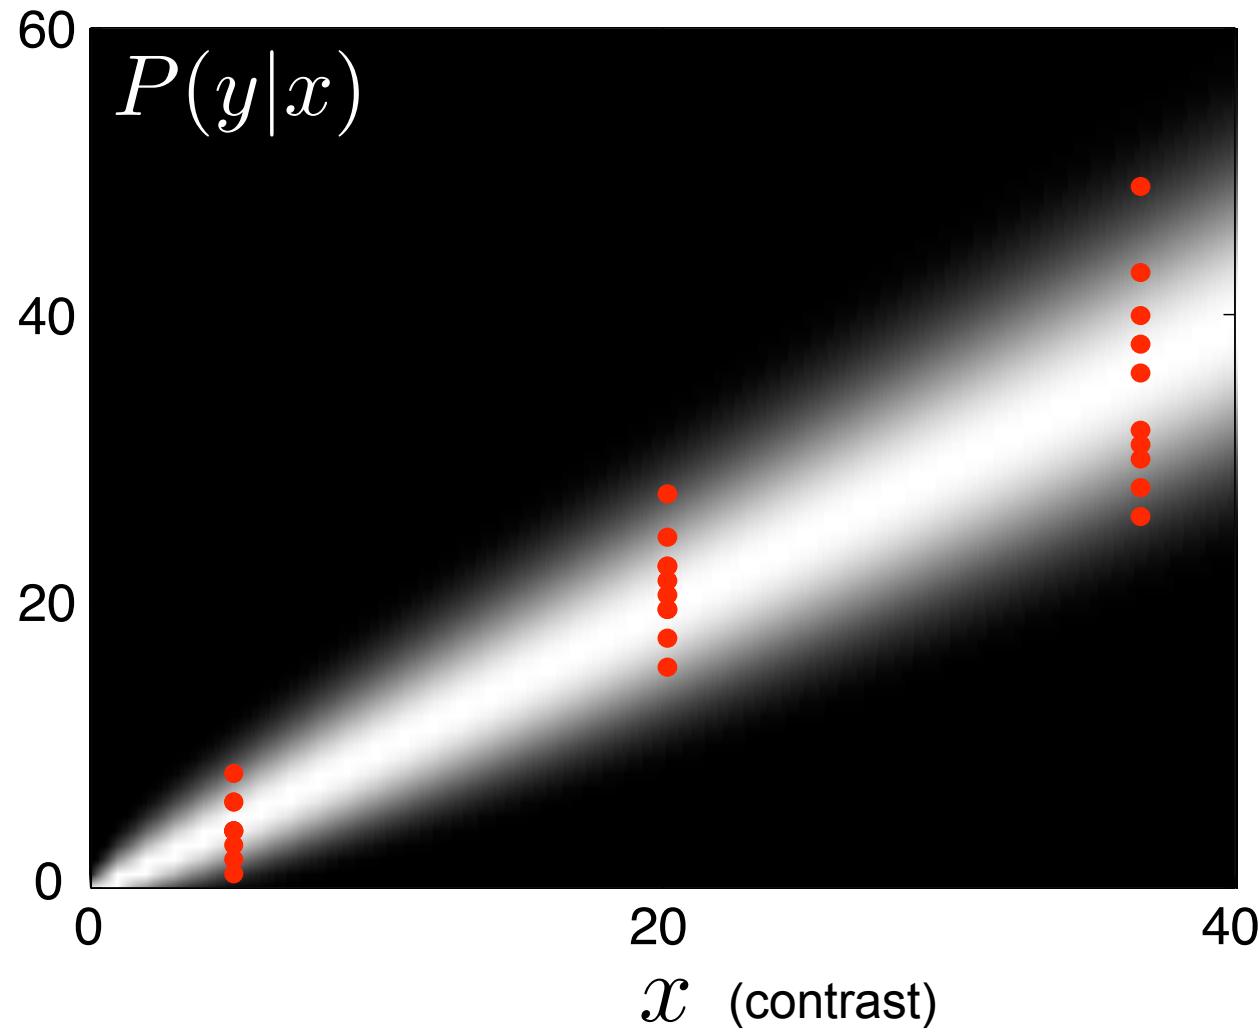
conditional distribution

$$p(y|x = 20)$$



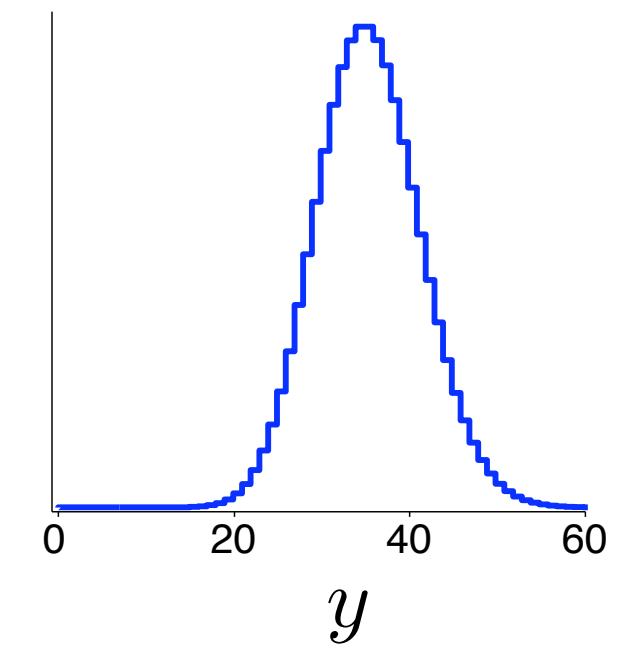
$$\text{mean}(y) = \theta x$$

$$\text{var}(y) = \theta x$$



conditional distribution

$$p(y|x = 35)$$



Maximum Likelihood Estimation:

- given observed data (Y, X) , find θ that maximizes $P(Y|X, \theta)$

↙ ↓ ↓
all spike all parameters
counts stimuli

$$P(Y|X, \theta) = \prod_{i=1}^N \underbrace{P(y_i|x_i, \theta)}_{\text{single-trial probability}}$$

Q: what assumption are we making about the responses?

A: conditional independence across trials!

Maximum Likelihood Estimation:

- given observed data (Y, X) , find θ that maximizes $P(Y|X, \theta)$


all spike counts all stimuli parameters

$$P(Y|X, \theta) = \prod_{i=1}^N \underbrace{P(y_i|x_i, \theta)}_{\text{single-trial probability}}$$

Q: what assumption are we making about the responses?

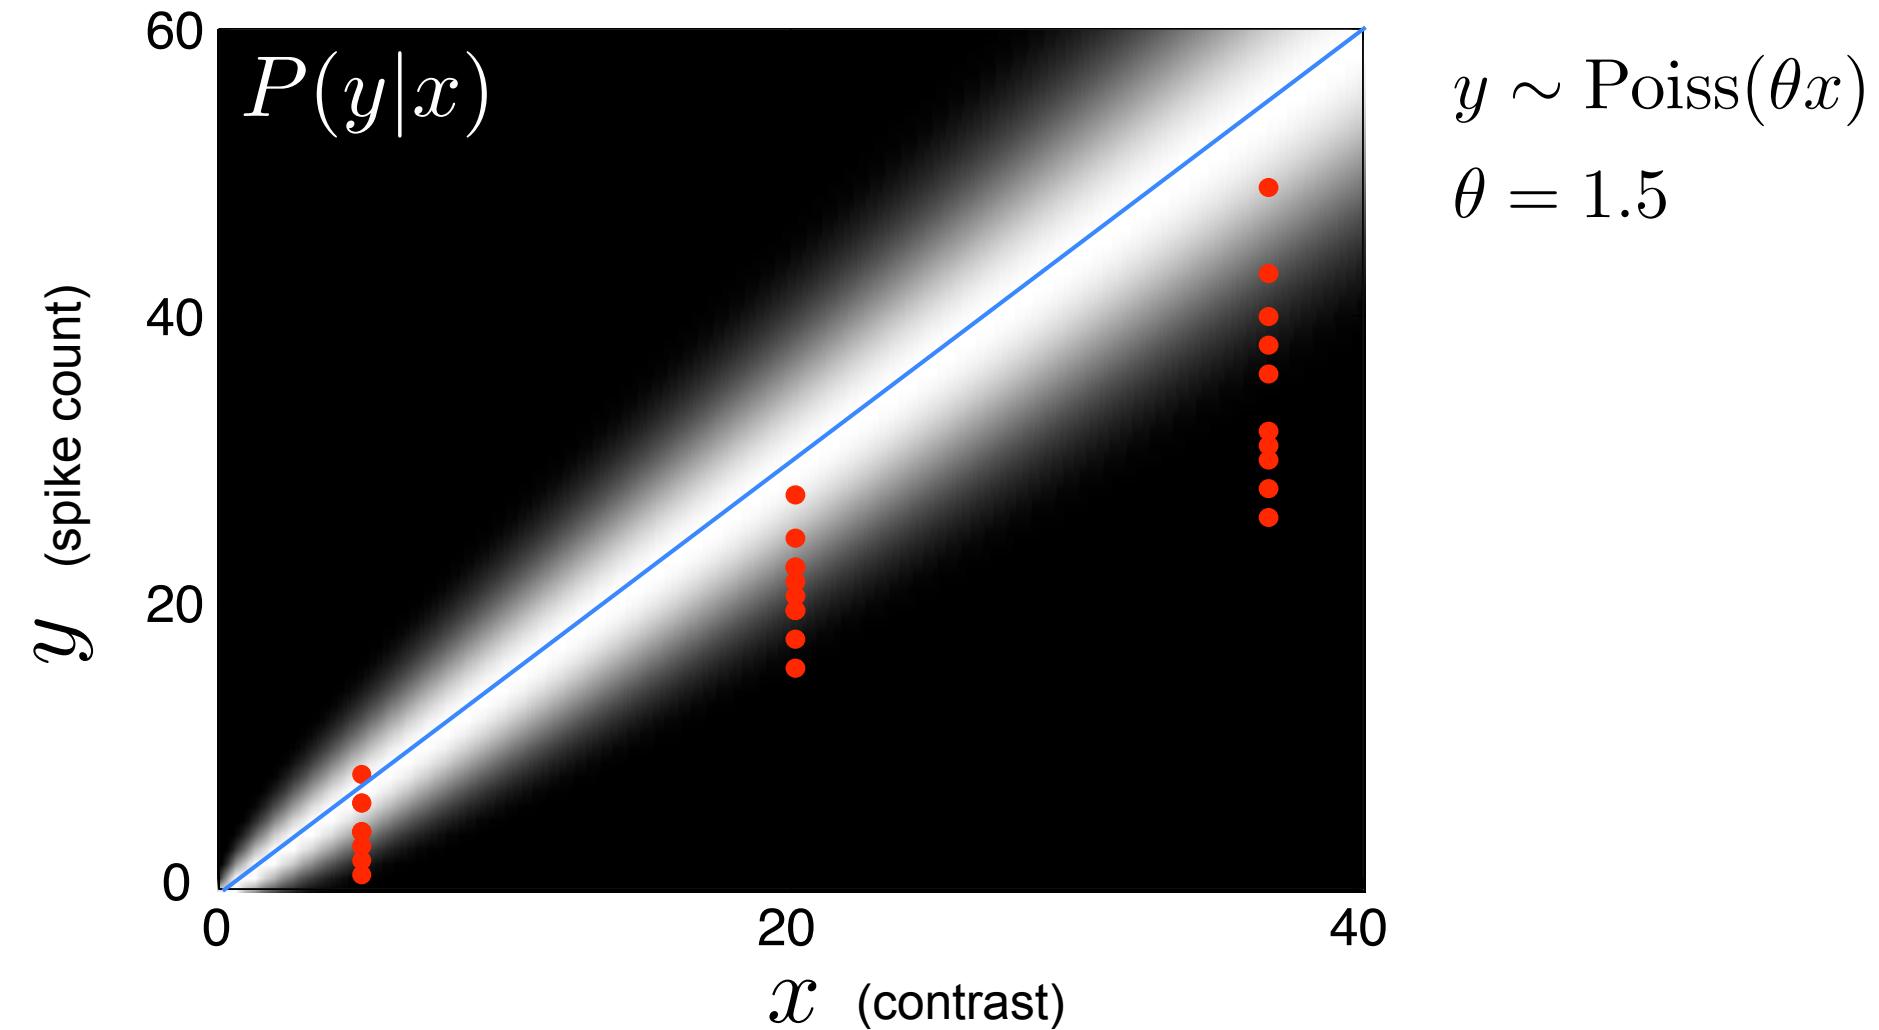
A: conditional independence across trials!

Q: when do we call $P(Y|X, \theta)$ a *likelihood*?

A: when considering it as a function of θ !

Maximum Likelihood Estimation:

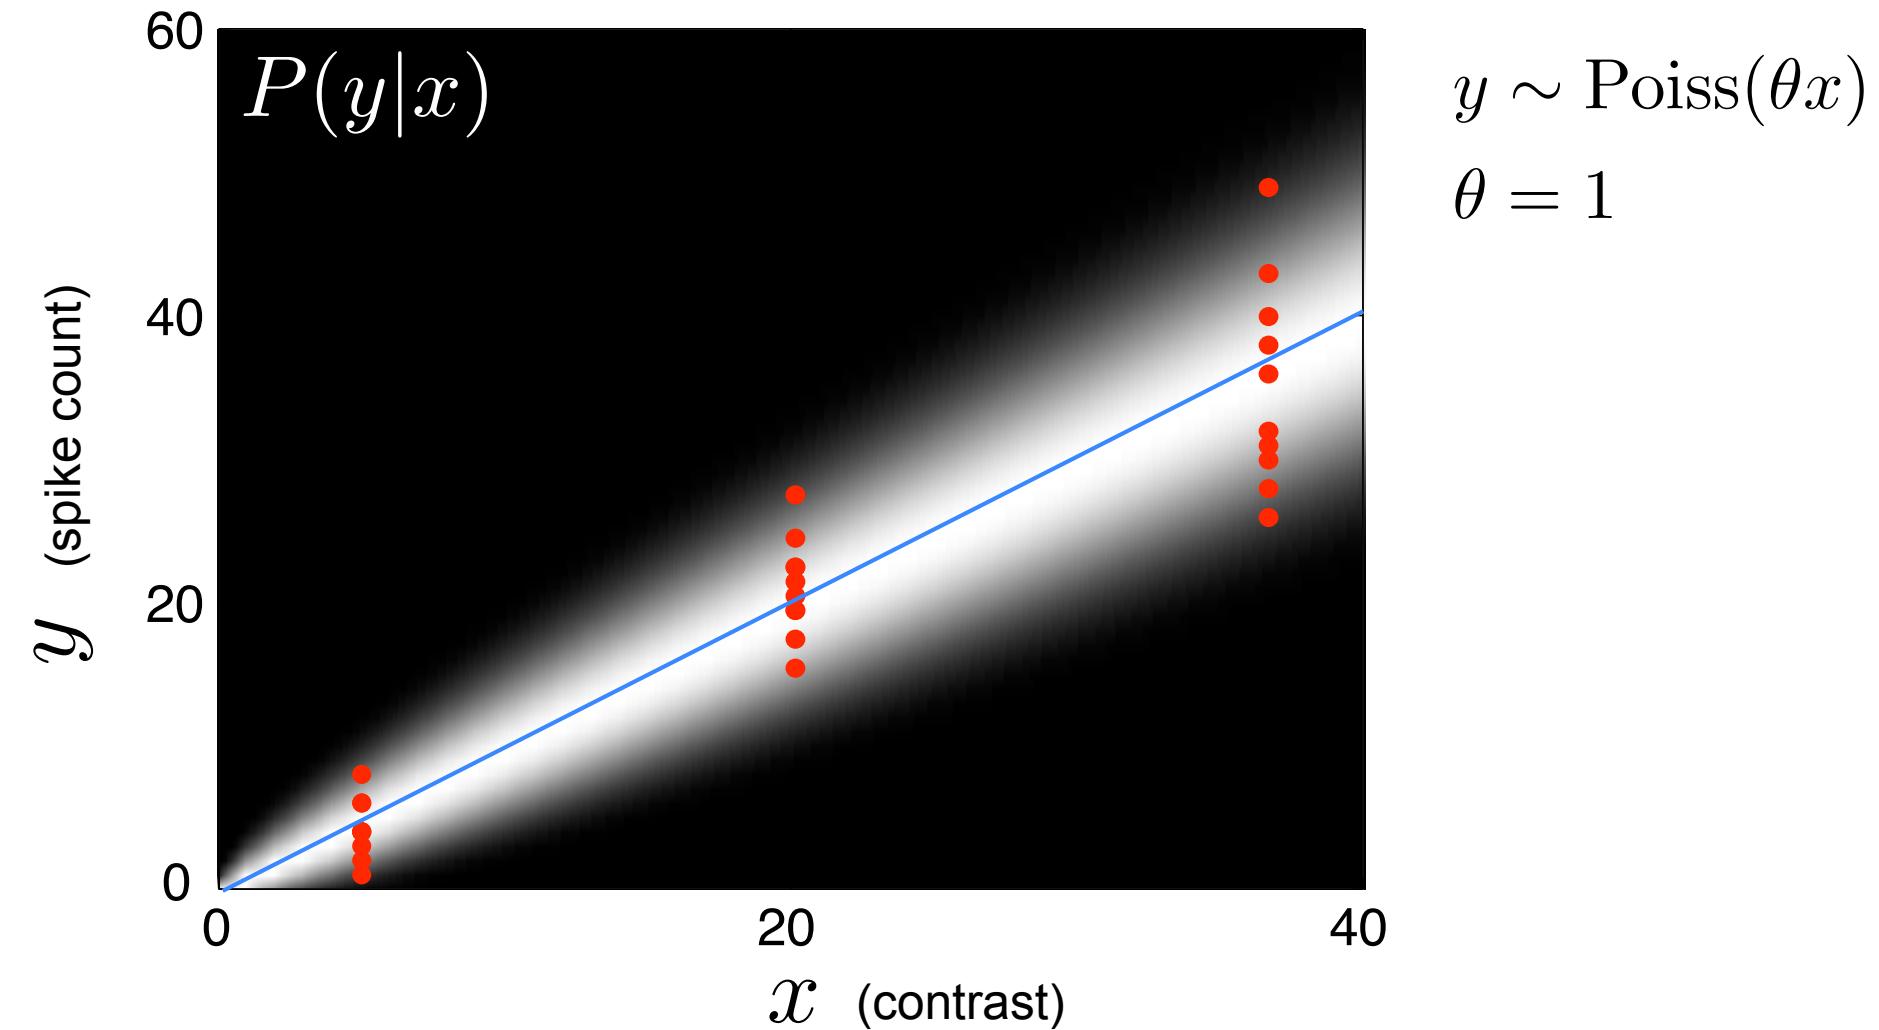
- given observed data (Y, X) , find θ that maximizes $P(Y|X, \theta)$



- could in theory do this by turning a knob

Maximum Likelihood Estimation:

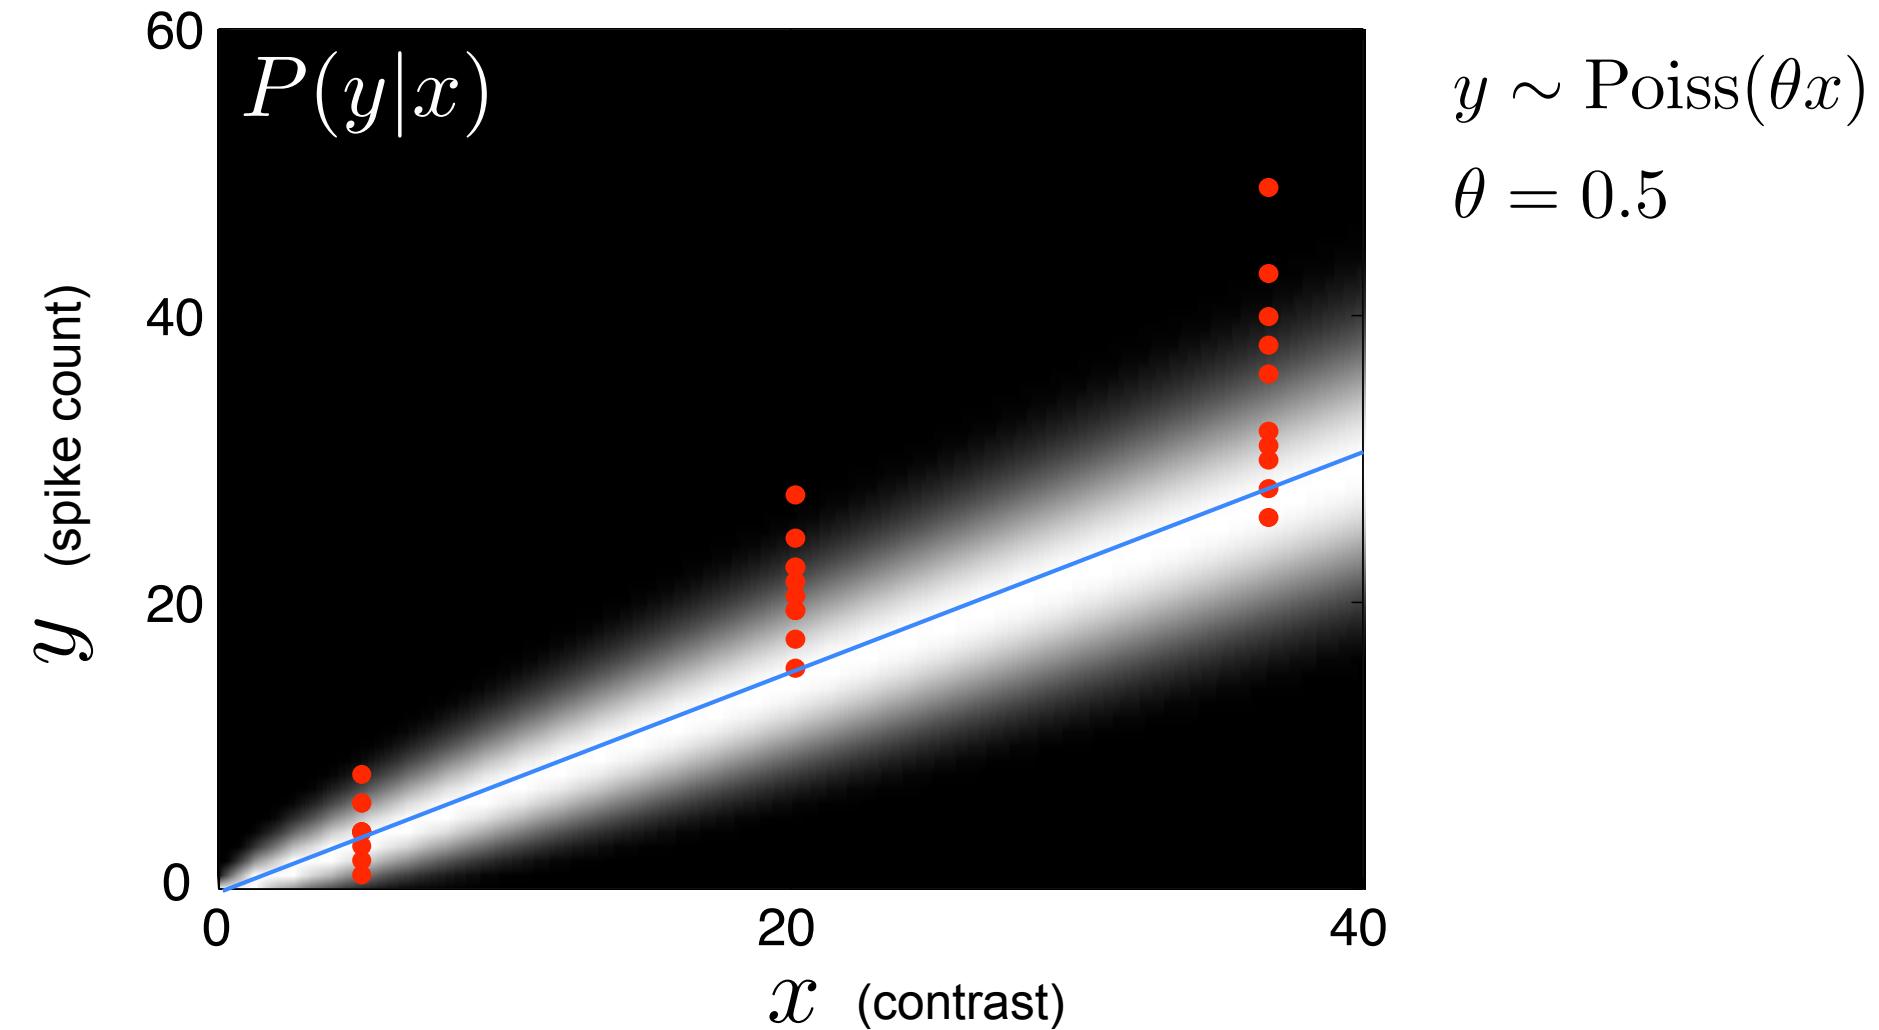
- given observed data (Y, X) , find θ that maximizes $P(Y|X, \theta)$



- could in theory do this by turning a knob

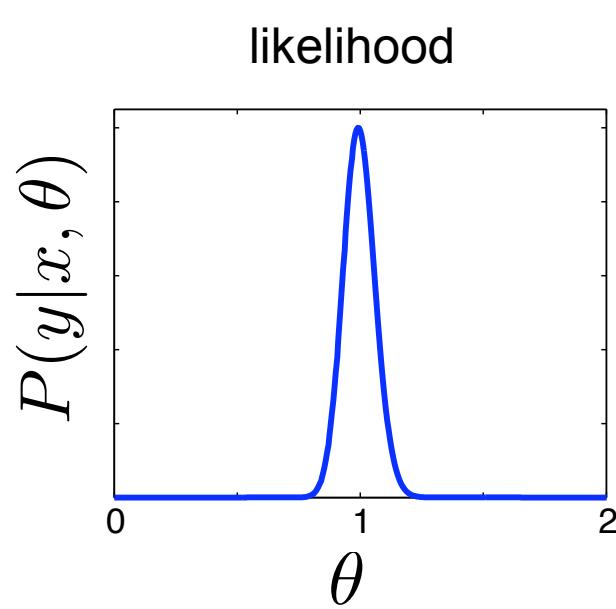
Maximum Likelihood Estimation:

- given observed data (Y, X) , find θ that maximizes $P(Y|X, \theta)$



- could in theory do this by turning a knob

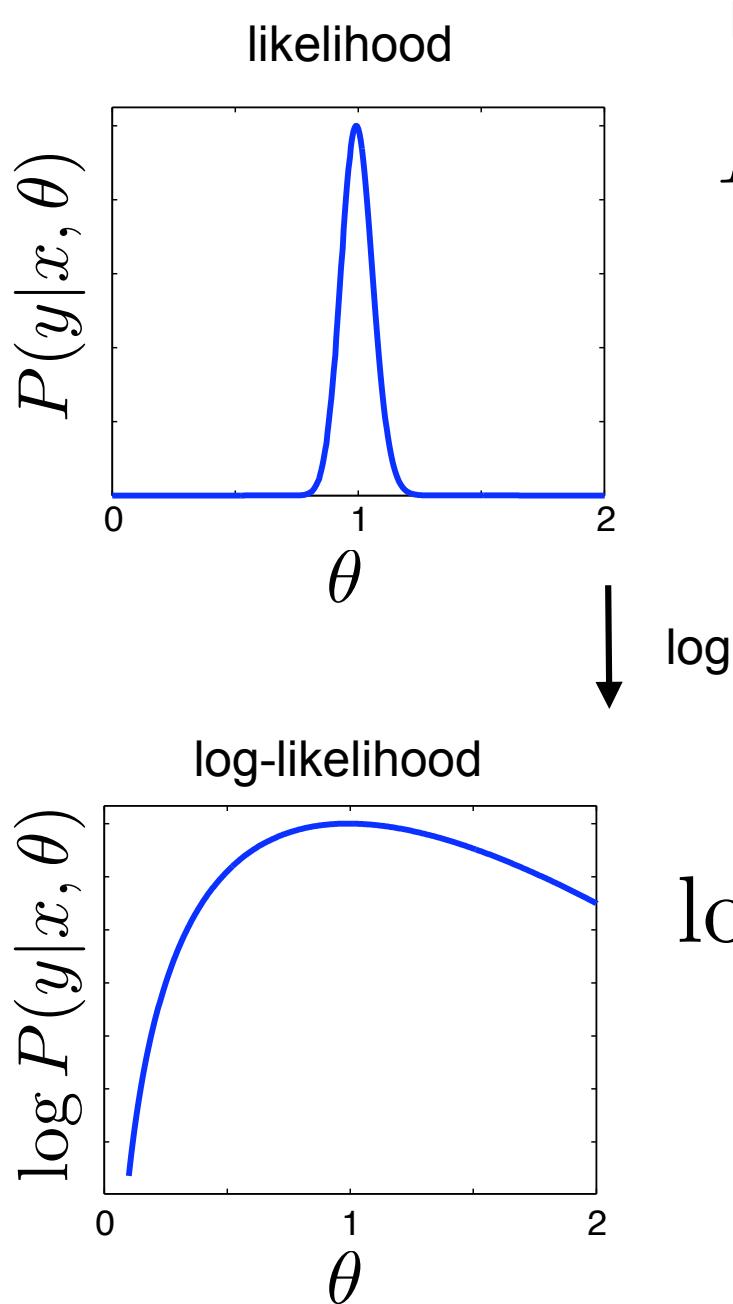
Likelihood function: $P(Y|X, \theta)$ as a function of θ



Because data are independent:

$$\begin{aligned}P(Y|X, \theta) &= \prod_i P(y_i|x_i, \theta) \\&= \prod_i \frac{1}{y_i!} (\theta x_i)^{y_i} e^{-(\theta x_i)}\end{aligned}$$

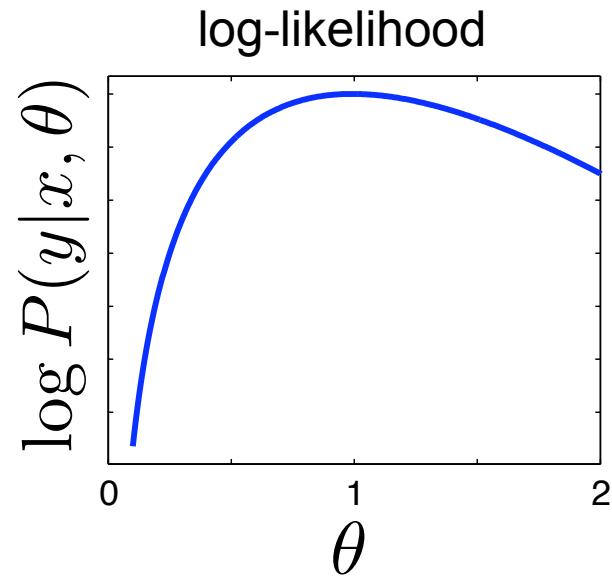
Likelihood function: $P(Y|X, \theta)$ as a function of θ



Because data are independent:

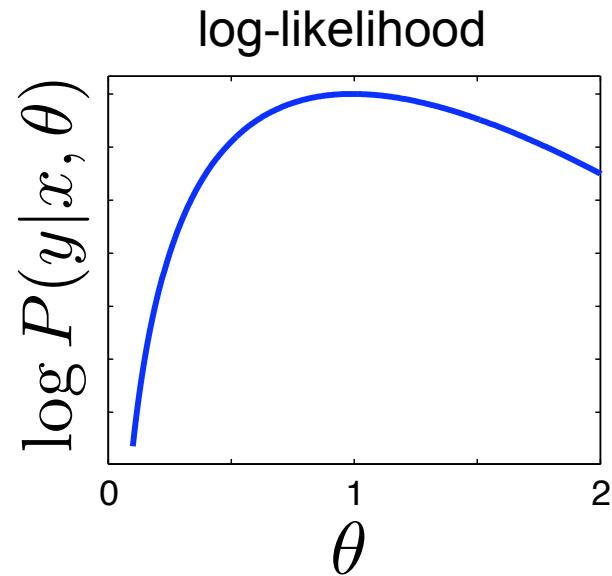
$$\begin{aligned}P(Y|X, \theta) &= \prod_i P(y_i|x_i, \theta) \\&= \prod_i \frac{1}{y_i!} (\theta x_i)^{y_i} e^{-(\theta x_i)}\end{aligned}$$

$$\begin{aligned}\log P(Y|X, \theta) &= \sum_i \log P(y_i|x_i, \theta) \\&= \sum y_i \log \theta - \theta x_i + c\end{aligned}$$



$$\begin{aligned}\log P(Y|X, \theta) &= \sum_i \log P(y_i|x_i, \theta) \\ &= \sum y_i \log \theta - \theta x_i + c \\ &= \log \theta(\sum y_i) - \theta(\sum x_i)\end{aligned}$$

Do it: solve for θ



$$\begin{aligned}
 \log P(Y|X, \theta) &= \sum_i \log P(y_i|x_i, \theta) \\
 &= \sum y_i \log \theta - \theta \sum x_i + c \\
 &= \log \theta (\sum y_i) - \theta (\sum x_i)
 \end{aligned}$$

- Closed-form solution when model in “exponential family”

$$\frac{d}{d\theta} \log P(Y|X, \theta) = \frac{1}{\theta} \sum y_i - \sum x_i = 0$$

$$\implies \hat{\theta}_{ML} = \frac{\sum y_i}{\sum x_i}$$

Properties of the MLE

(maximum likelihood estimator)

- **consistent**
(converges to true θ in limit of infinite data)
- **efficient**
(converges as quickly as possible,
i.e., achieves minimum possible asymptotic error)

simple example #2: linear Gaussian neuron

$$\text{spike rate} \quad \mu = \theta x$$

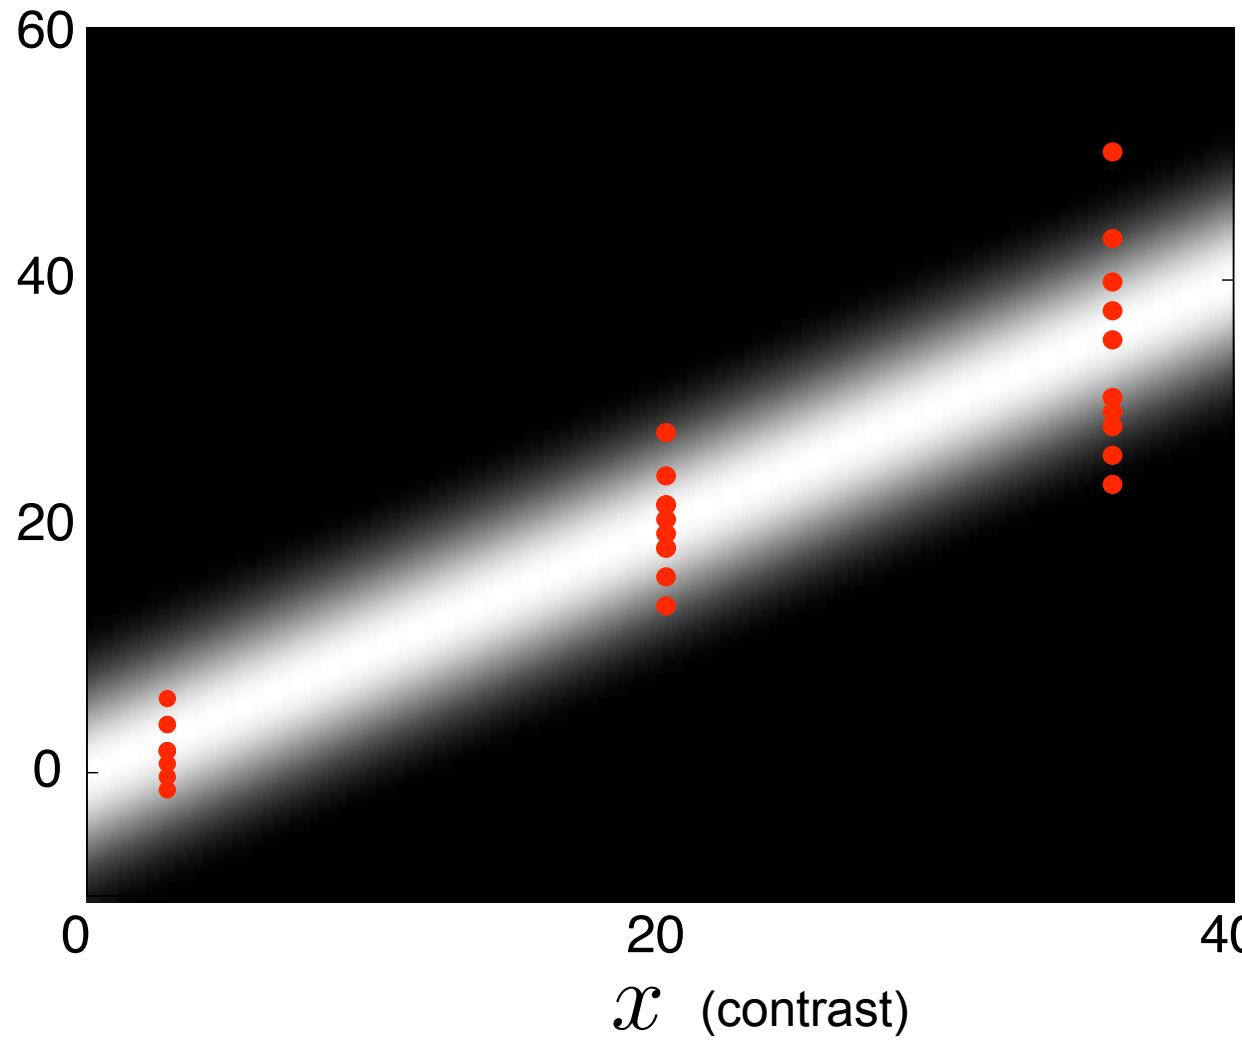
parameter stimulus

spike count $y \sim \mathcal{N}(\mu, \sigma^2)$

encoding model:

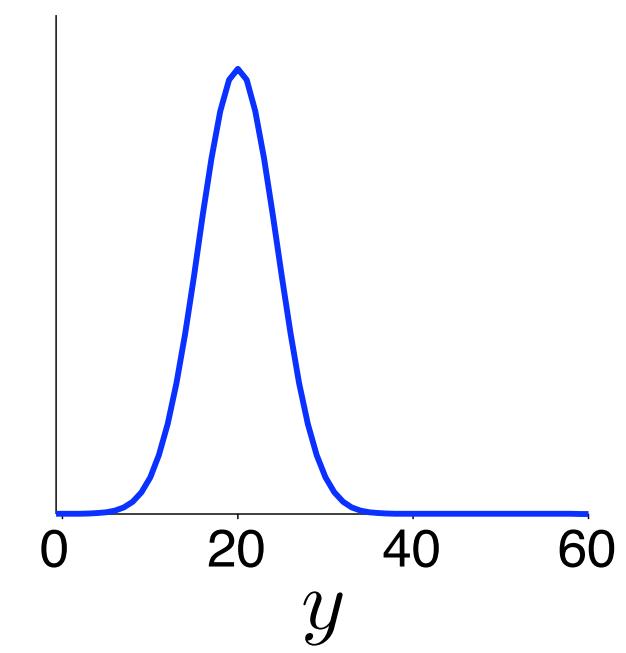
$$P(y|x, \theta) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{\frac{-(y-\theta x)^2}{2\sigma^2}}$$

$$\text{mean}(y) = \theta x$$
$$\text{var}(y) = \sigma^2$$



encoding distribution

$$p(y|x = 20)$$



All slices have same width

$$P(y|x, \theta) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{\frac{-(y-\theta x)^2}{2\sigma^2}}$$

Log-Likelihood $\log P(Y|X, \theta) = -\sum \frac{(y_i - \theta x_i)^2}{2\sigma^2} + c$

Differentiate, set to zero, and solve for θ

$$P(y|x, \theta) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{\frac{-(y-\theta x)^2}{2\sigma^2}}$$

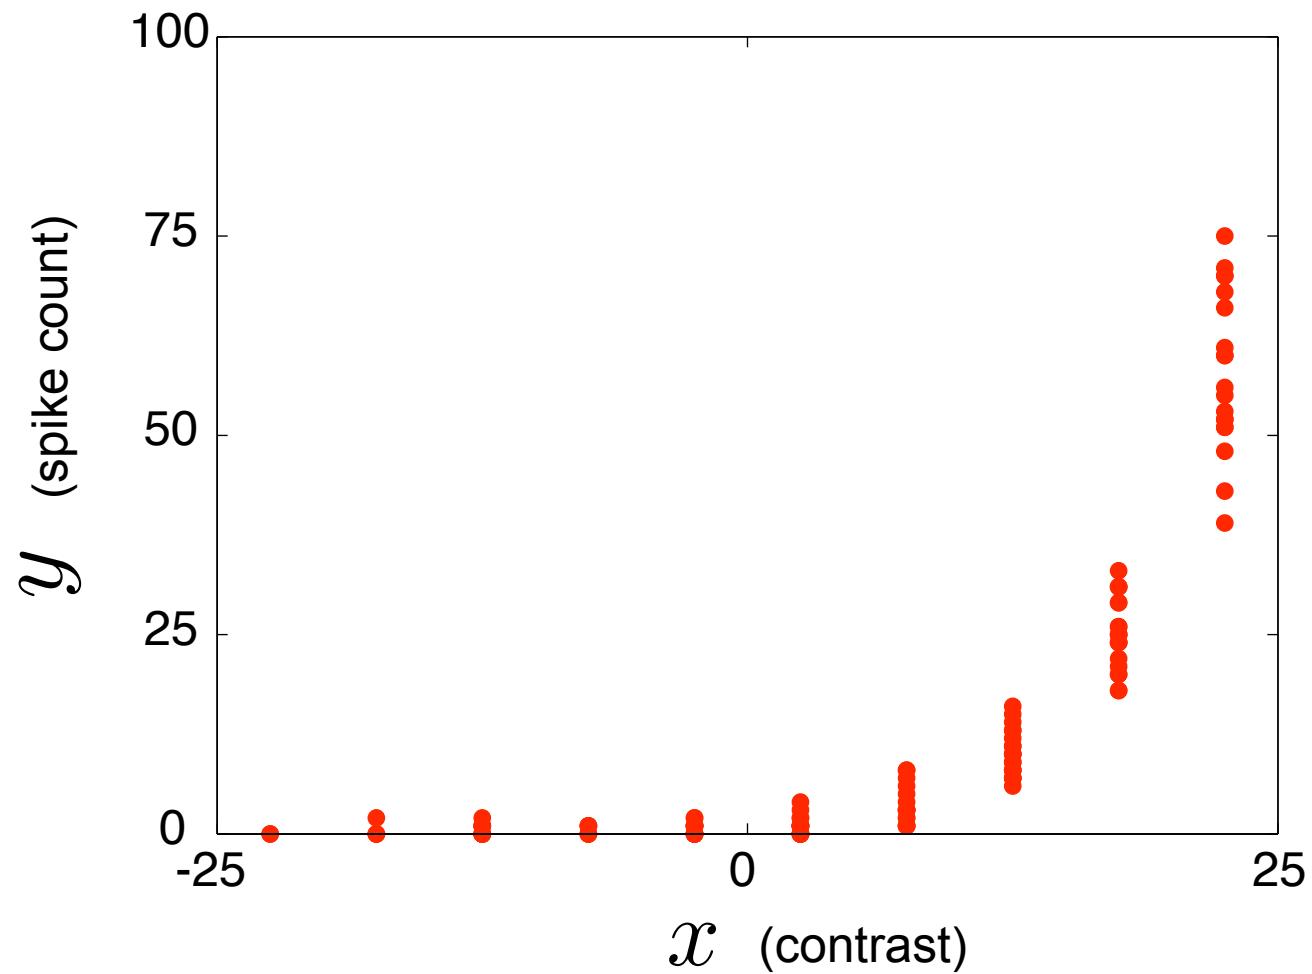
Log-Likelihood $\log P(Y|X, \theta) = -\sum \frac{(y_i - \theta x_i)^2}{2\sigma^2} + c$

$$\frac{d}{d\theta} \log P(Y|X, \theta) = -\sum \frac{(y_i - \theta x_i)x_i}{\sigma^2} = 0$$

Maximum-Likelihood Estimator: $\hat{\theta}_{ML} = \frac{\sum y_i x_i}{\sum x_i^2}$
 (“Least squares regression” solution)

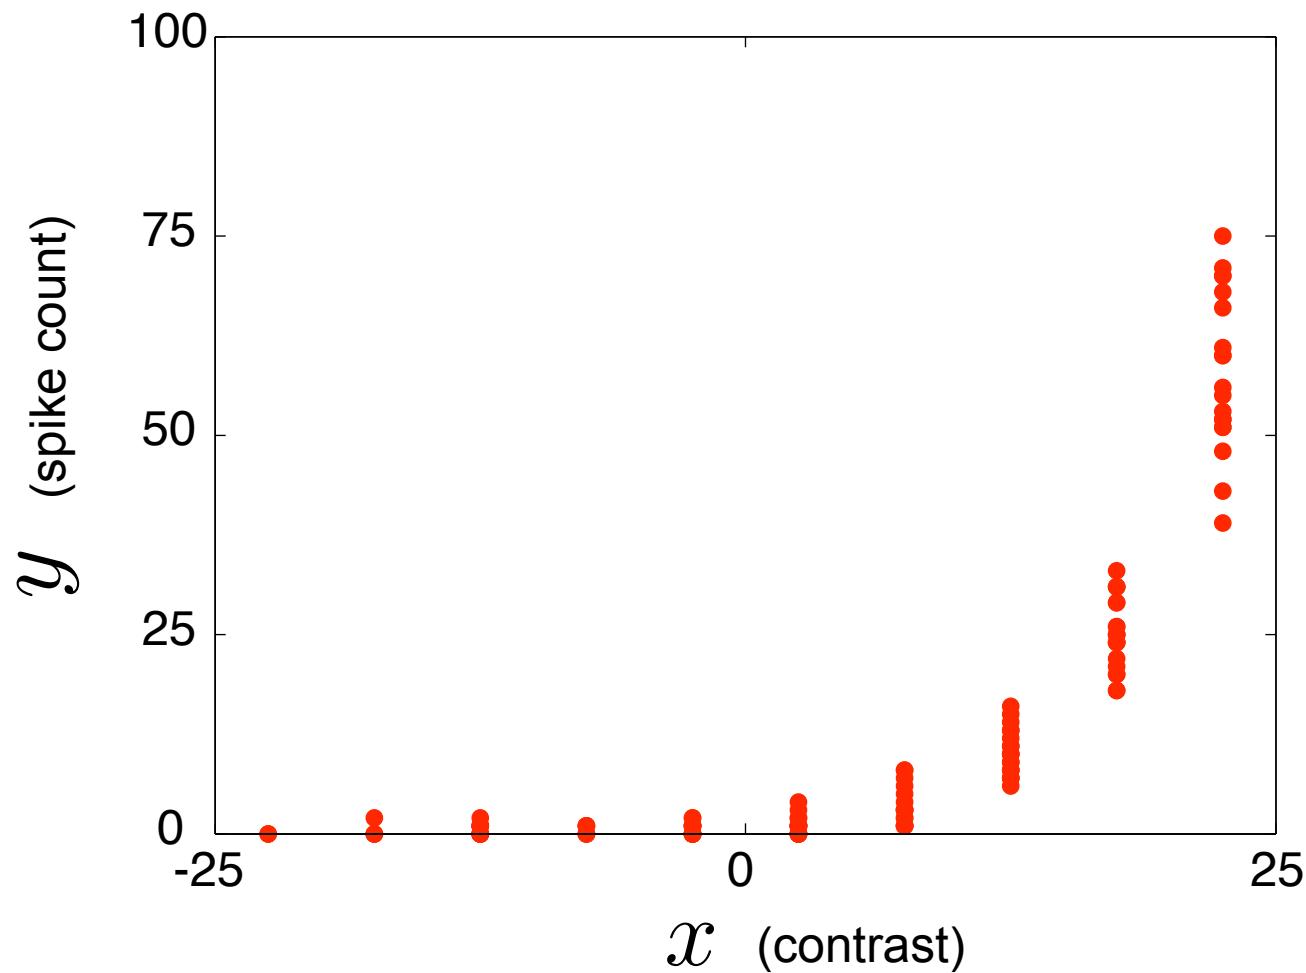
(Recall that for Poisson, $\hat{\theta}_{ML} = \frac{\sum y_i}{\sum x_i}$)

example #3: unknown neuron



Be the computational neuroscientist: what model would you use?

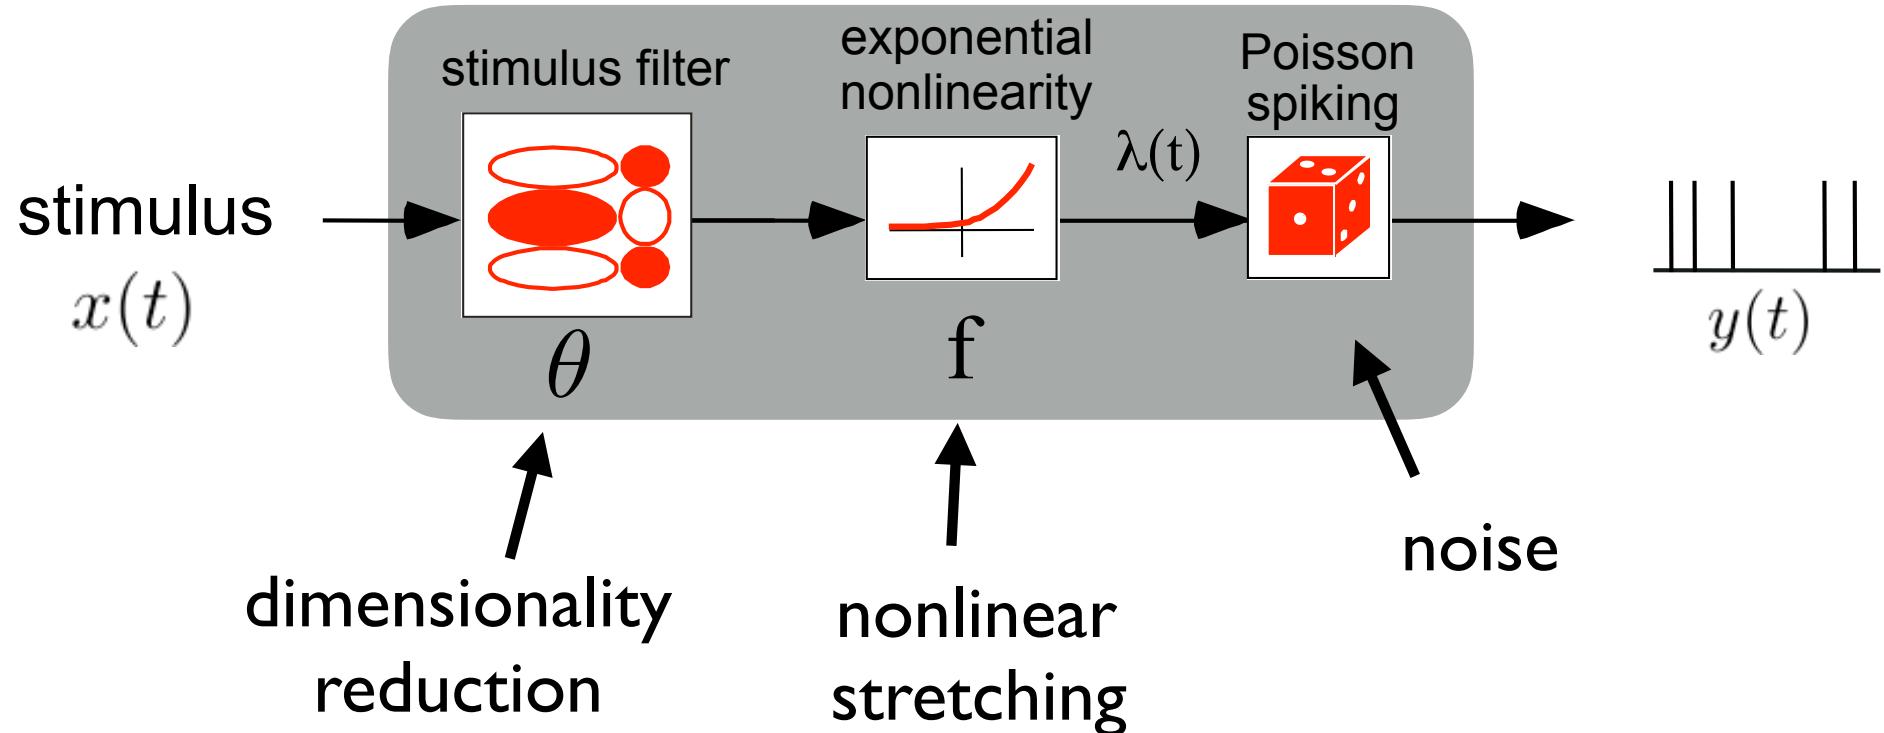
Example 3: unknown neuron



More general setup: $\lambda = f(\theta x)$ firing rate is nonlinear
This is a GLM! $y \sim Poiss(\lambda)$ Poisson firing

“basic” Poisson generalized linear model (GLM)

Linear-Nonlinear-Poisson (LNP) model



$$\text{spike rate } \lambda = f(\vec{k} \cdot \vec{x})$$
$$\text{spike count } y \sim \text{Poiss}(\lambda)$$

- also known as a “cascade” model

What is a GLM?

Be careful about terminology:

GLM

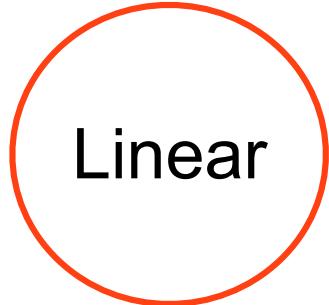
\neq

GLM

General Linear Model

Generalized Linear Model

(Nelder 1972)



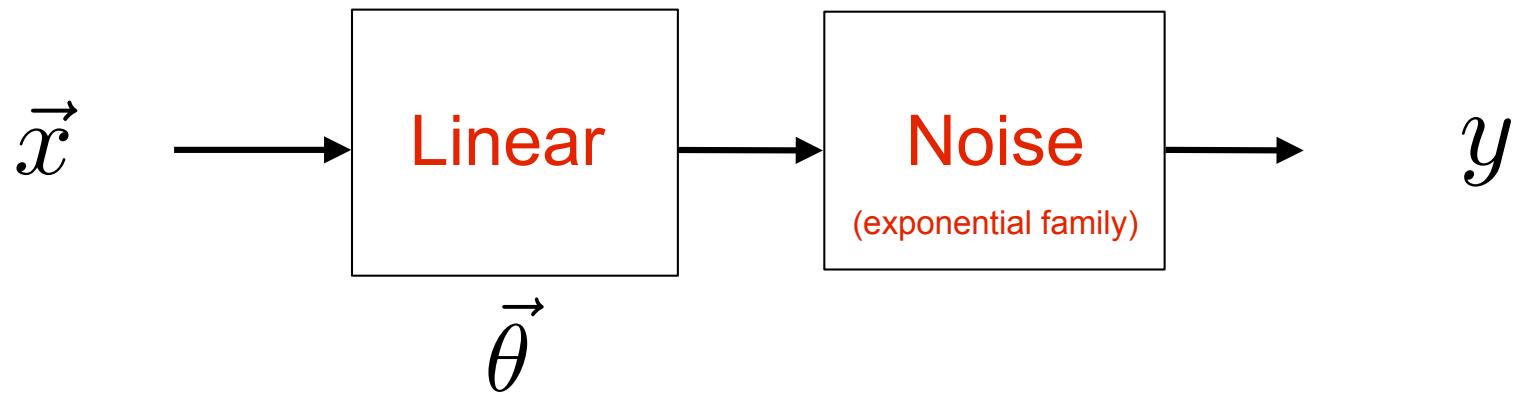
Linear



Linear

Moral:
Be careful when naming your model!

1. General Linear Model



“Dimensionality
Reduction”

Examples:

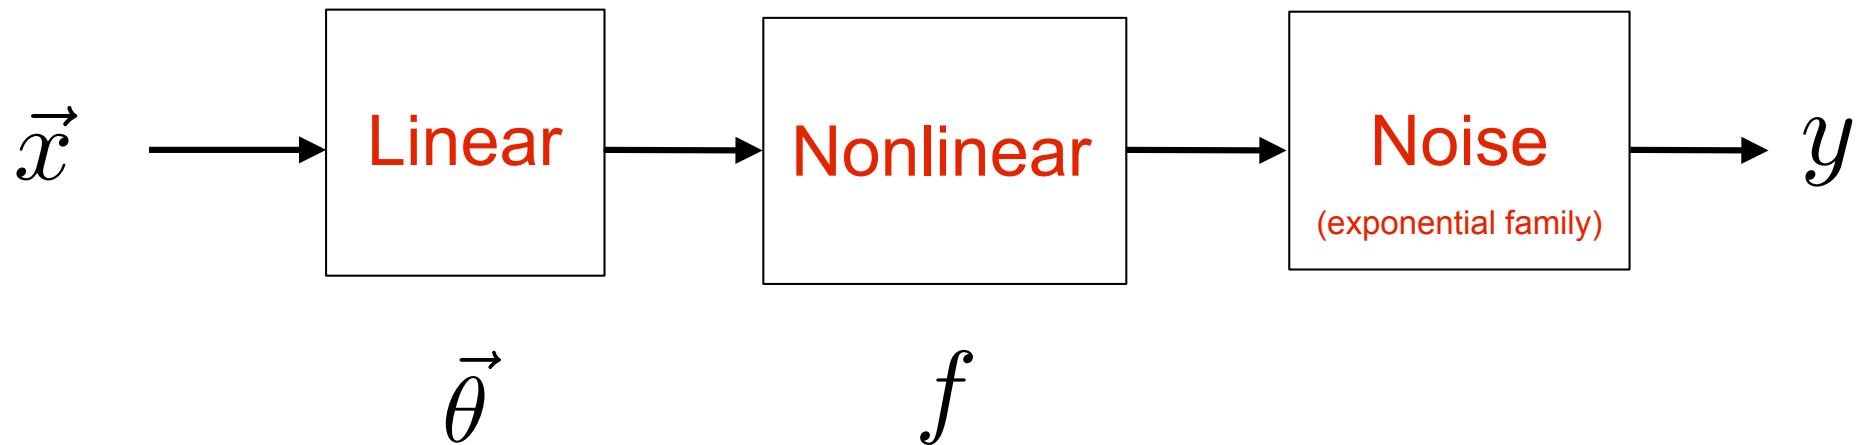
1. Gaussian

$$y = \vec{\theta} \cdot \vec{x} + \epsilon$$

2. Poisson

$$y \sim \text{Poiss}(\vec{\theta} \cdot \vec{x})$$

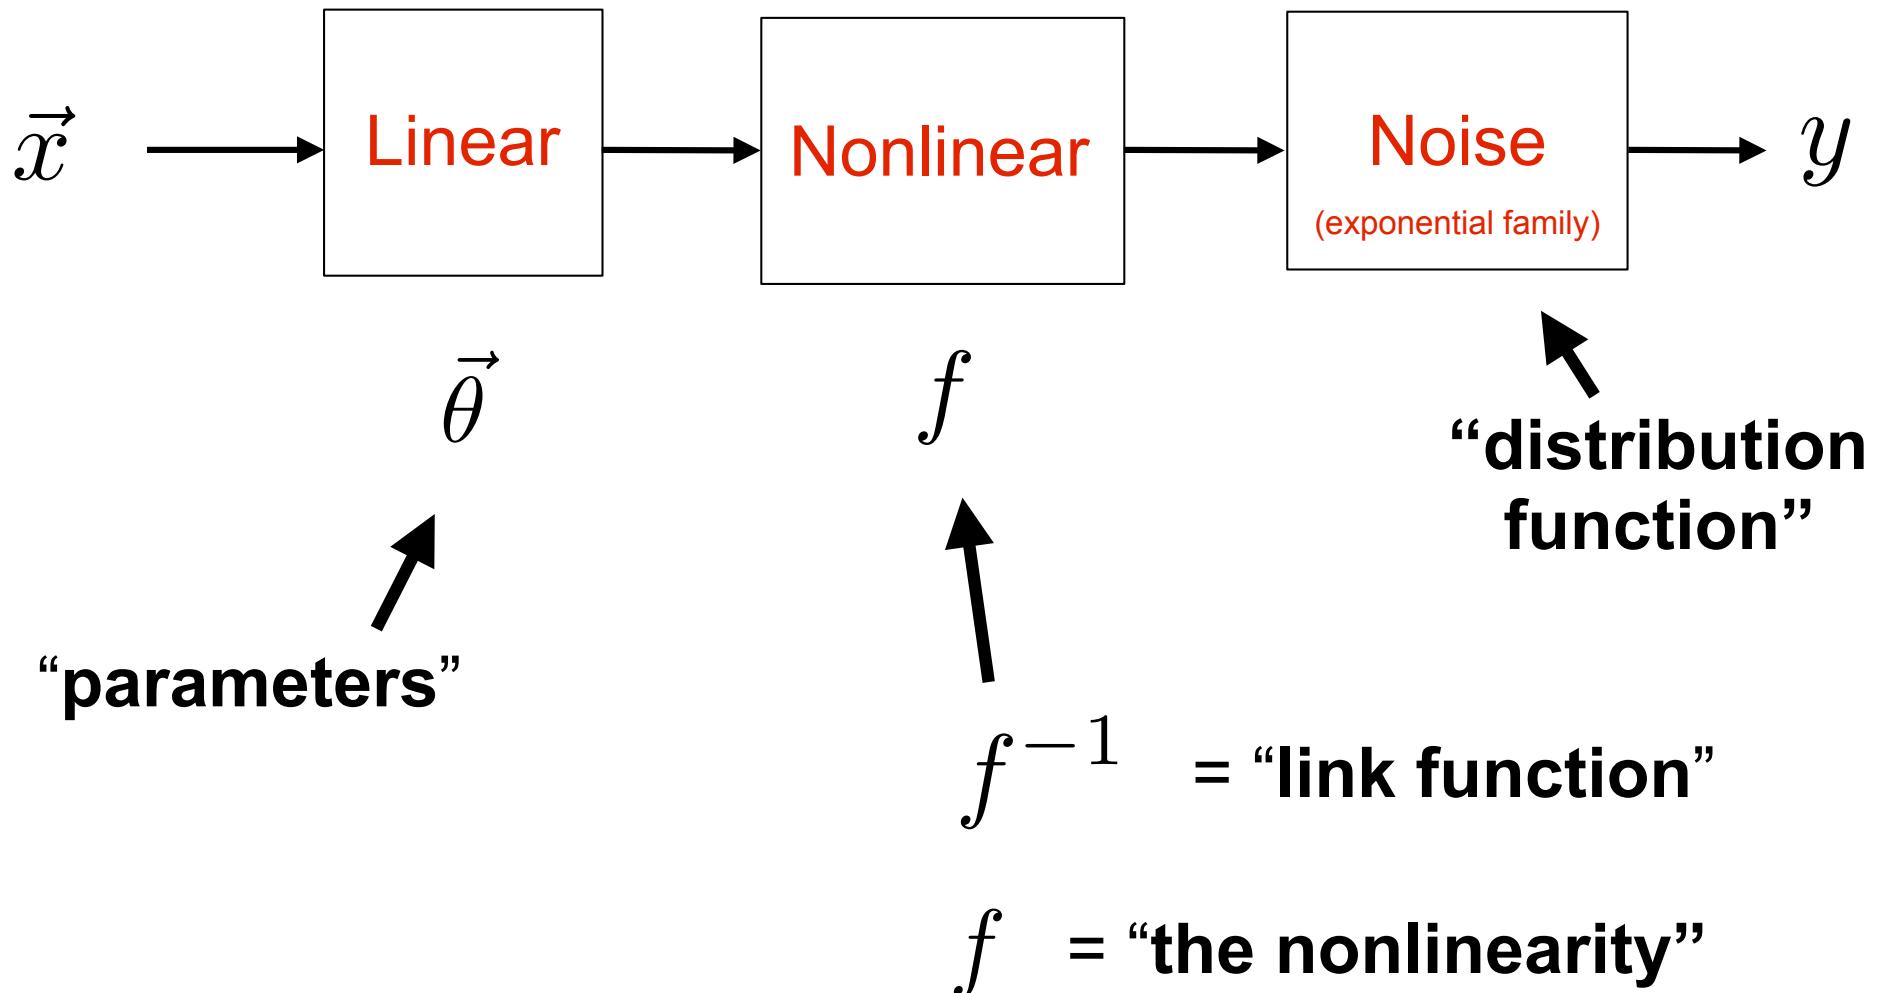
2. Generalized Linear Model



- Examples:
- 1. Gaussian $y = f(\vec{\theta} \cdot \vec{x}) + \epsilon$
 - 2. Poisson $y \sim \text{Poiss}(f(\vec{\theta} \cdot \vec{x}))$

2. Generalized Linear Model

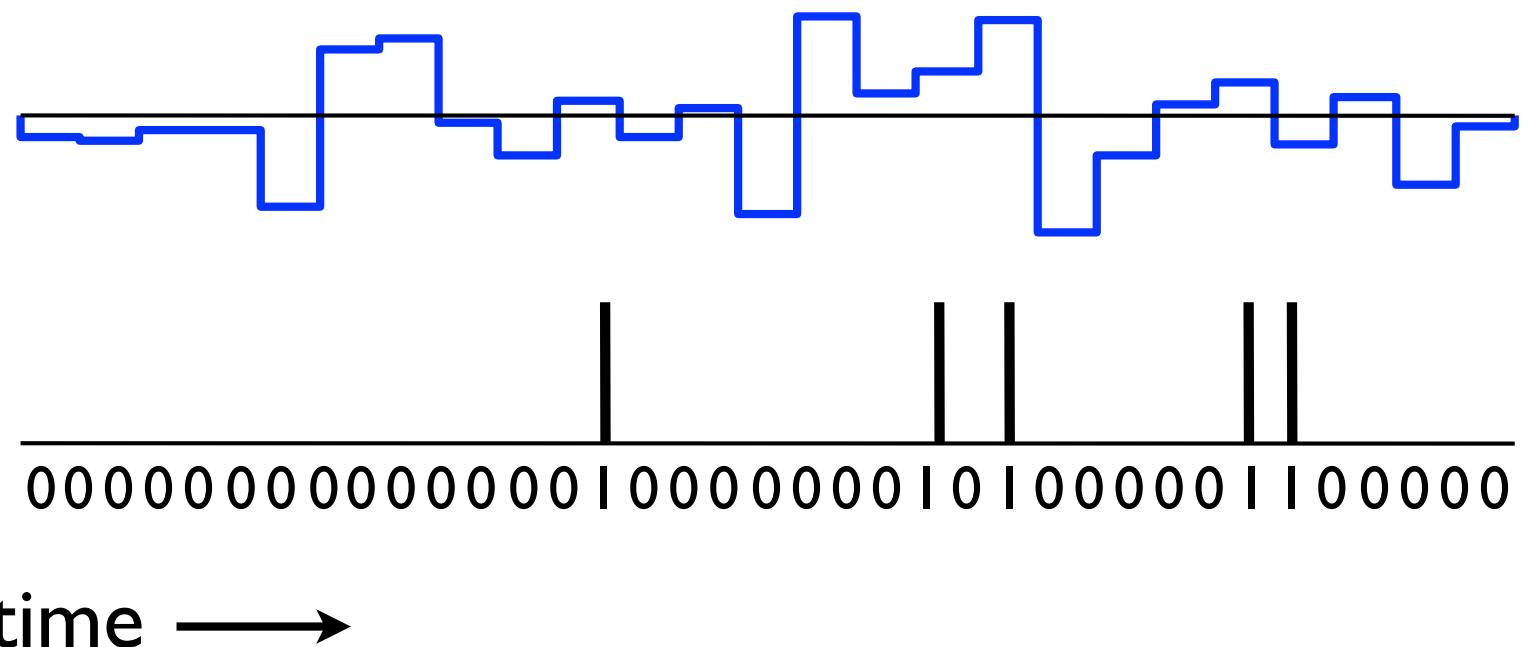
Terminology:



Applying it to data

response at time t

stimulus



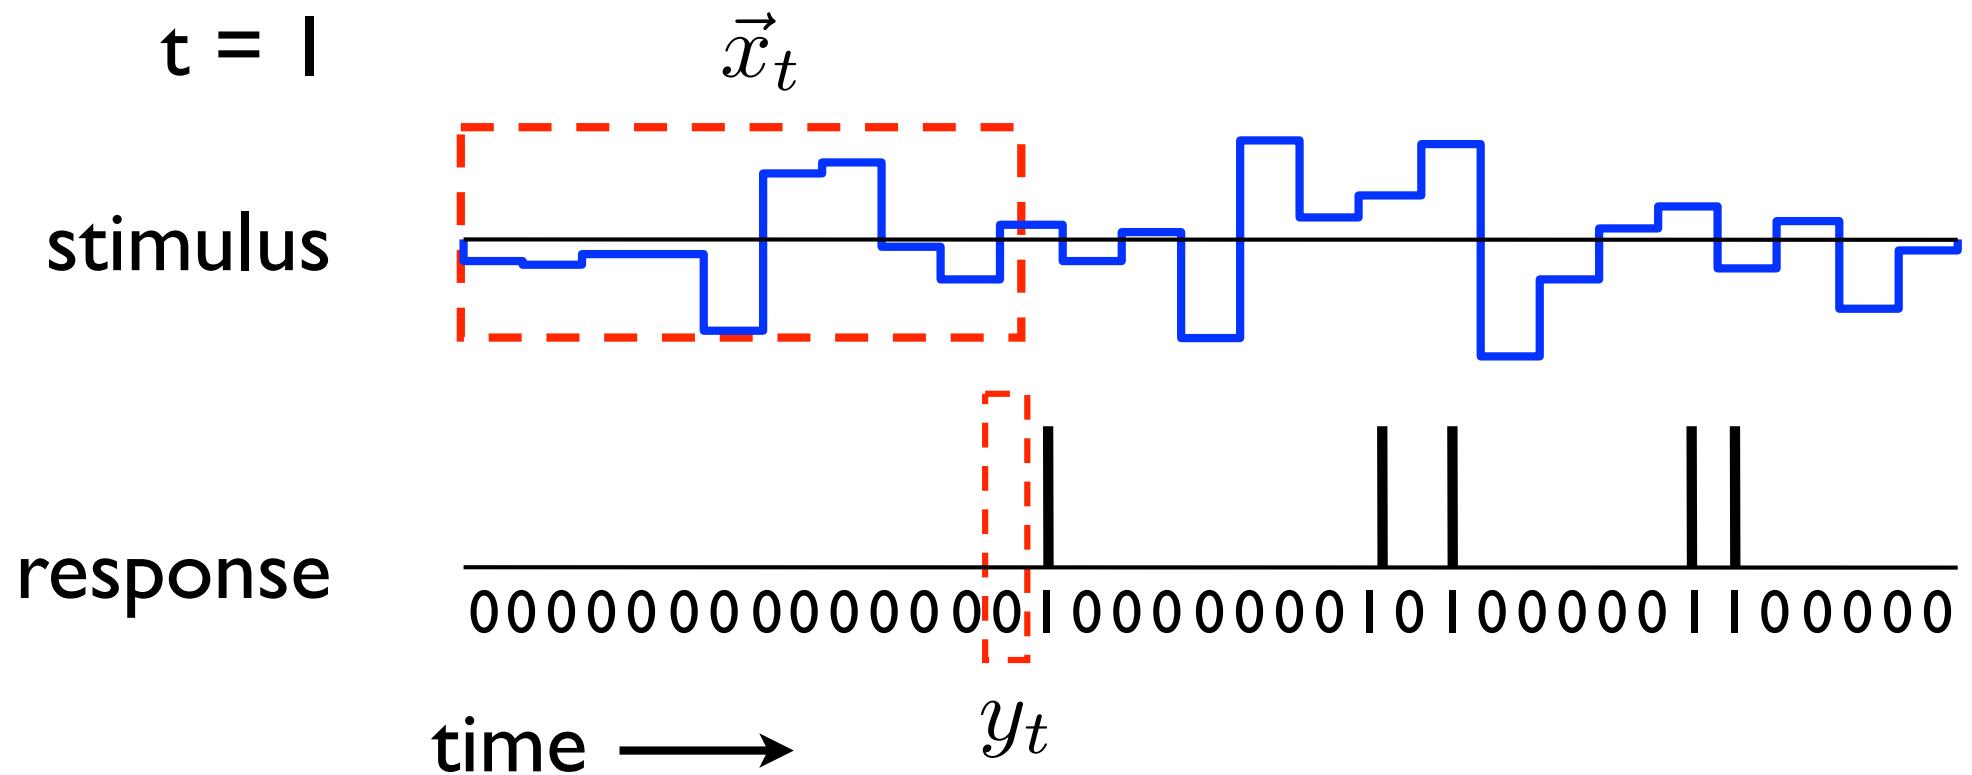
response
at time t

$$y_t = \vec{k} \cdot \vec{x}_t + \text{noise}$$

↑ ↑
linear filter vector stimulus
 at time t

walk through the data
one time bin at a time

$$t = 1$$



response
at time t

$$y_t = \vec{k} \cdot \vec{x}_t + \text{noise}$$

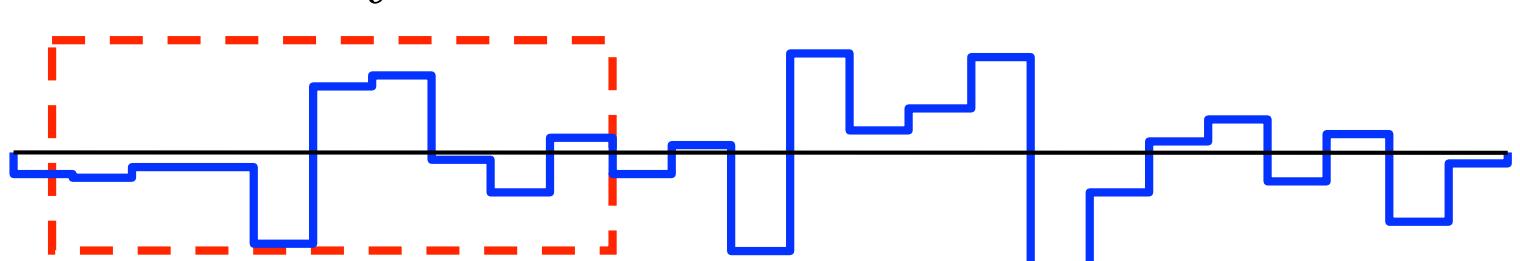
↑ ↑
linear filter vector stimulus
 at time t

walk through the data
one time bin at a time

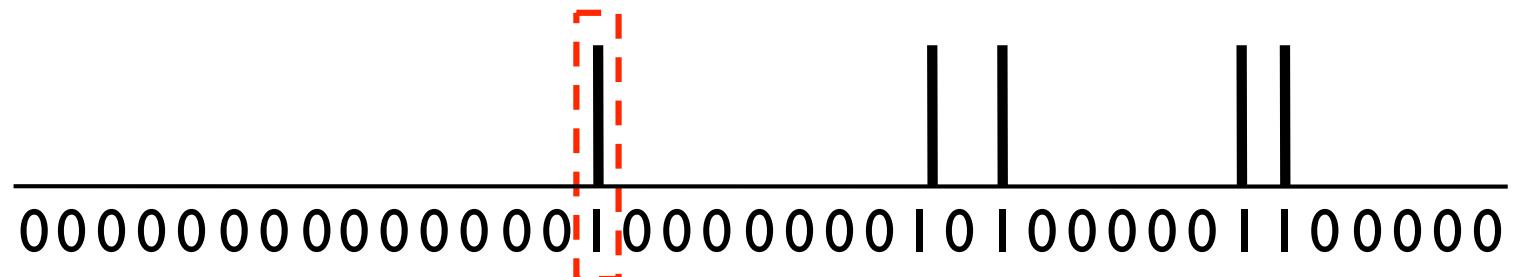
$t = 2$

\vec{x}_t

stimulus



response



time →

y_t

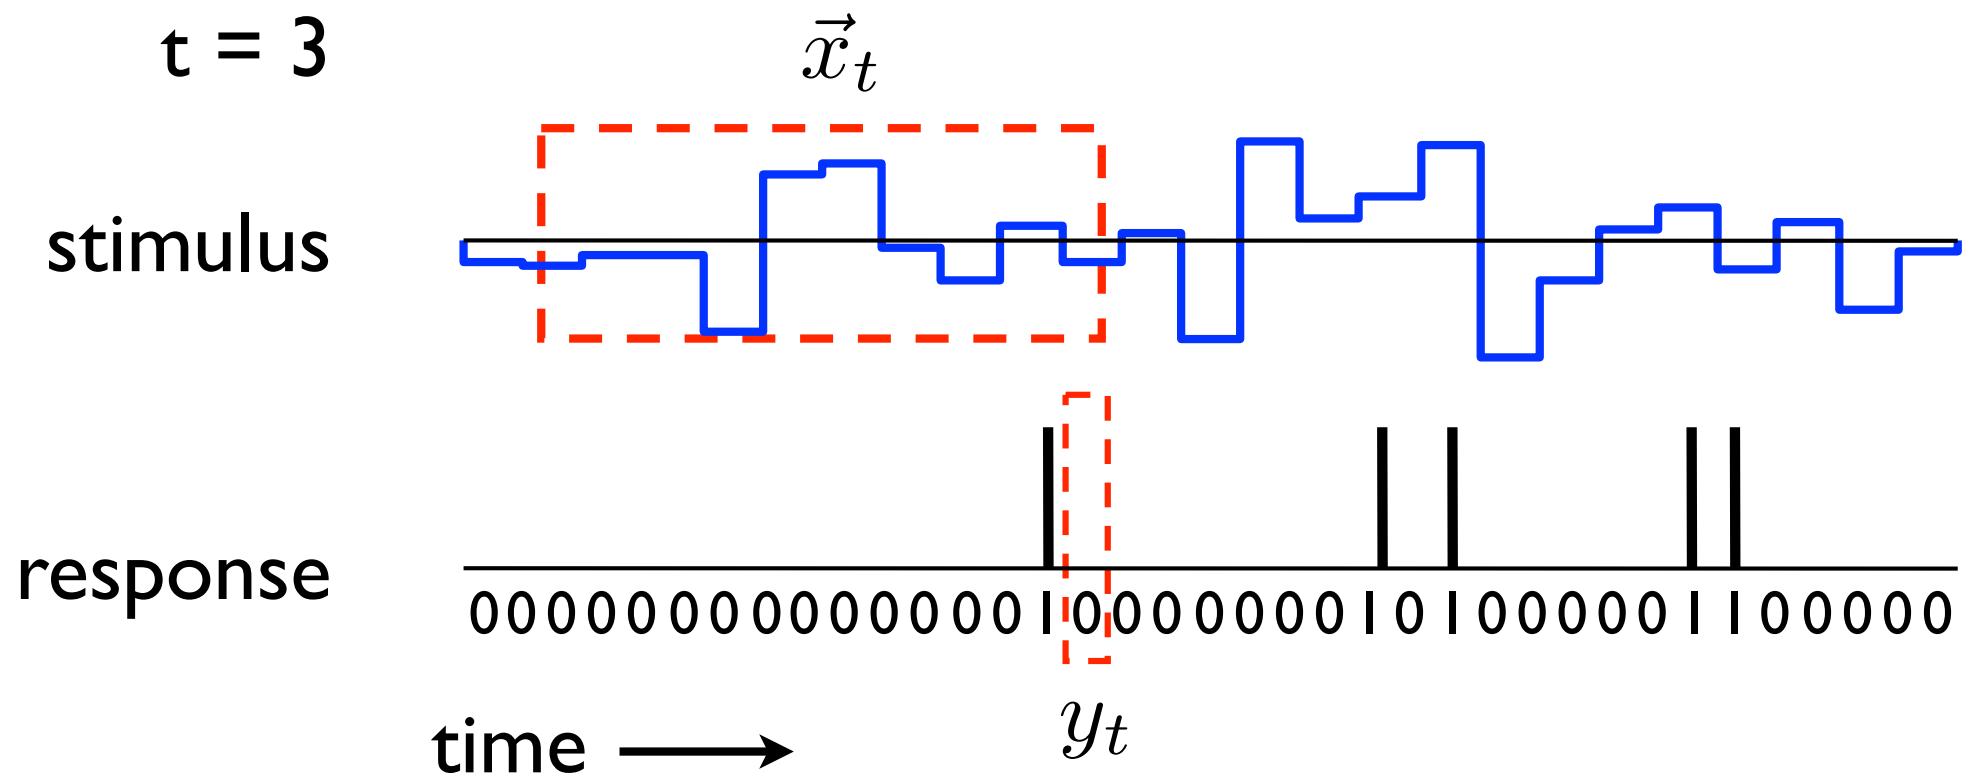
response
at time t

$$y_t = \vec{k} \cdot \vec{x}_t + \text{noise}$$

↑ ↑
linear filter vector stimulus
 at time t

walk through the data
one time bin at a time

$t = 3$



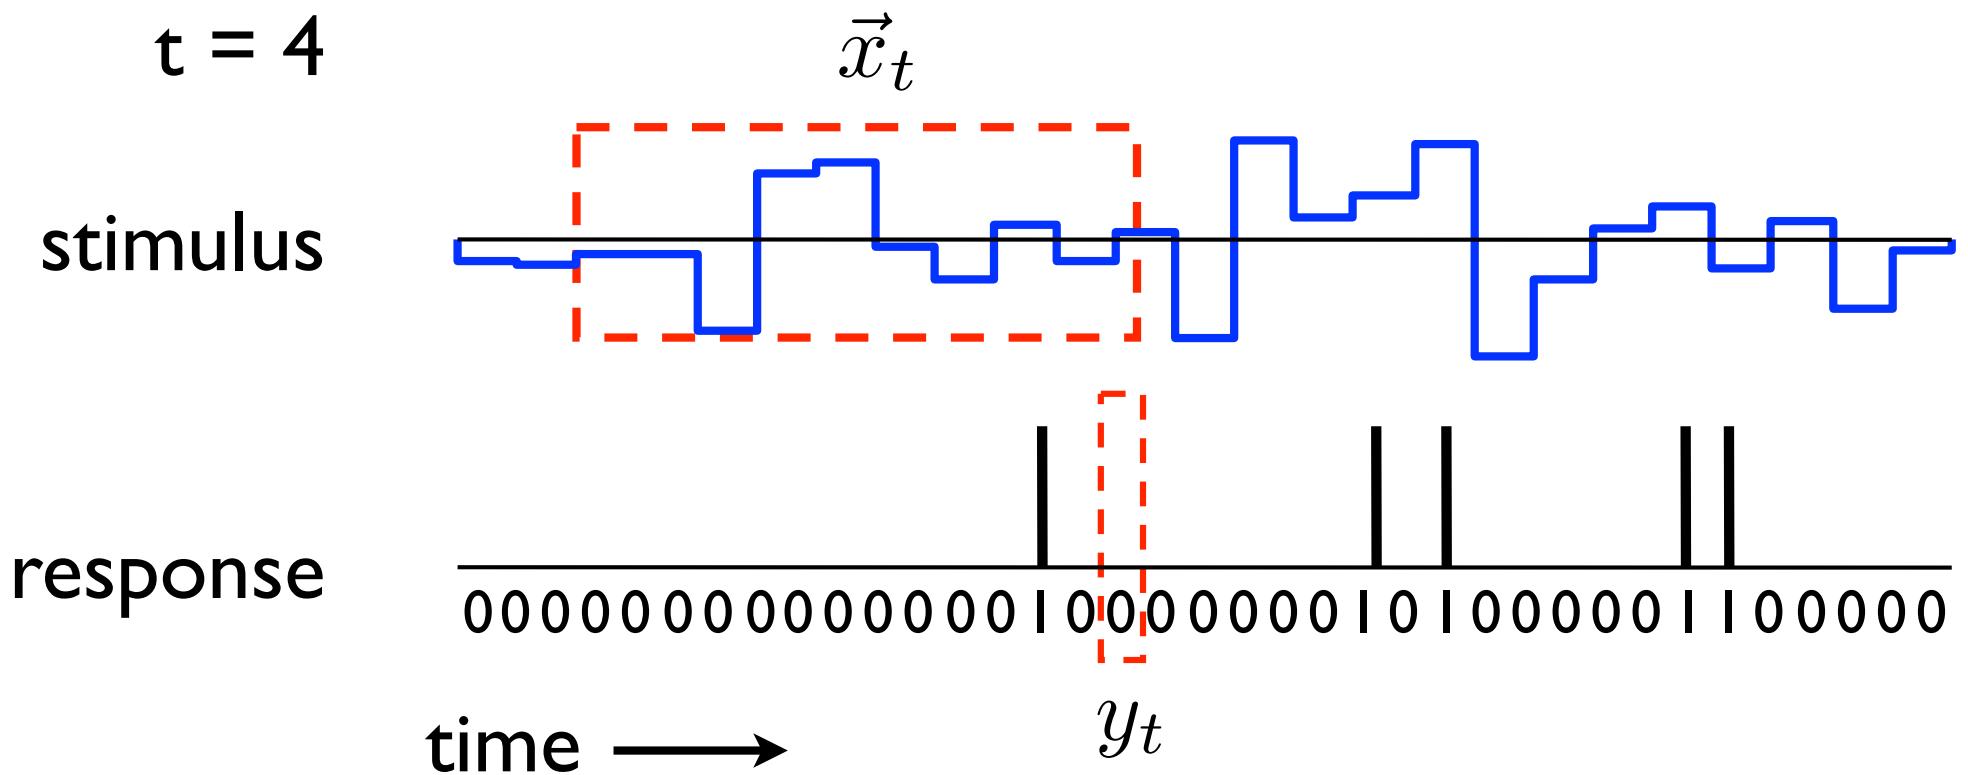
response
at time t

$$y_t = \vec{k} \cdot \vec{x}_t + \text{noise}$$

↑ ↑
linear filter vector stimulus
 at time t

walk through the data
one time bin at a time

$t = 4$



response
at time t

$$y_t = \vec{k} \cdot \vec{x}_t + \text{noise}$$

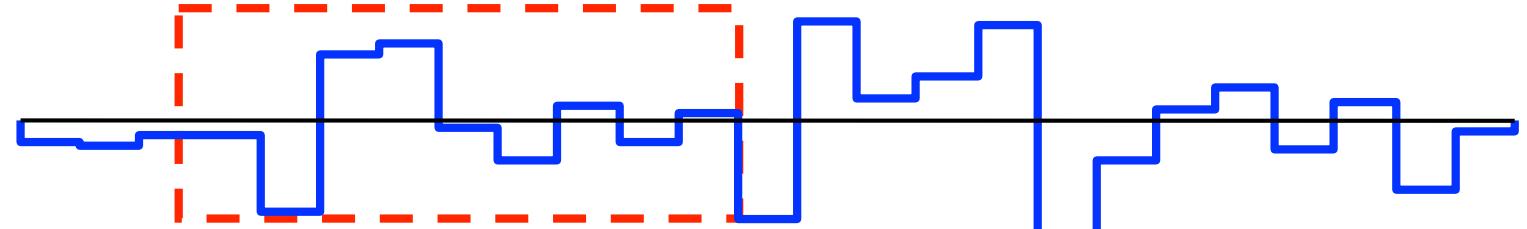
↑ ↑
linear filter vector stimulus
 at time t

walk through the data
one time bin at a time

$t = 5$

\vec{x}_t

stimulus



response

000000000000000010000000101000001100000

time →

y_t

response
at time t

$$y_t = \vec{k} \cdot \vec{x}_t + \text{noise}$$

↑ ↑
linear filter vector stimulus
 at time t

walk through the data
one time bin at a time

$t = 6$

stimulus

\vec{x}_t

response

000000000000000010000000101000001100000

time —→

y_t

Build up to following matrix version:

Y = $X\vec{k}$ + *noise*

↑
time ↓

=

$\begin{bmatrix} 0 \\ 0 \\ 1 \\ \vdots \end{bmatrix}$

$\begin{bmatrix} \text{blue step functions} \\ \text{blue step functions} \\ \text{blue step functions} \\ \vdots \end{bmatrix}$

$\begin{bmatrix} \vec{k} \end{bmatrix}$

design matrix

Computing maximum likelihood estimate

I. “Linear-Gaussian” GLM:

$$\hat{k} = (X^T X)^{-1} X^T Y$$

The diagram illustrates the decomposition of the inverse covariance matrix $(X^T X)^{-1}$ into two components: 'stimulus covariance' and 'spike-triggered average (STA)'. The matrix is shown as a large black bracket above two red brackets. The left red bracket covers the first term $(X^T X)^{-1}$ and is labeled 'stimulus covariance' below it. The right red bracket covers the second term $X^T Y$ and is labeled 'spike-triggered avg (STA)' below it.

Computing maximum likelihood estimate

$$Y = f(X \vec{k}) + noise$$

time ↓

2. Poisson GLM: `k = glmfit(X, Y, 'Poisson');`

maximum likelihood fit
(assumes exponential nonlinearity by default)

Computing maximum likelihood estimate

$$Y = f(X \vec{k}) + noise$$

time ↓

3. Bernoulli GLM: `k = glmfit(X, Y, 'binomial');`
outputs 0 and 1 (assumes **logistic** nonlinearity by default)

“logistic regression”

GLM summary

I. Linear-Gaussian GLM: $Y|X, \vec{k} \sim \mathcal{N}(X\vec{k}, \sigma^2 I)$ continuous

log-likelihood: $-\frac{1}{2\sigma^2} (Y - X\vec{k})^\top (Y - X\vec{k}) + const$

MLE: $\hat{\vec{k}} = (X^T X)^{-1} X^T Y$

2. Poisson GLM: $y|\vec{x}, \vec{k} \sim \text{Poiss}(f(\vec{x}_t \cdot \vec{k}))$ integer counts

log-likelihood: $\mathcal{L} = Y^\top \log f(X\vec{k}) - \mathbf{1}^\top f(X\vec{k})$

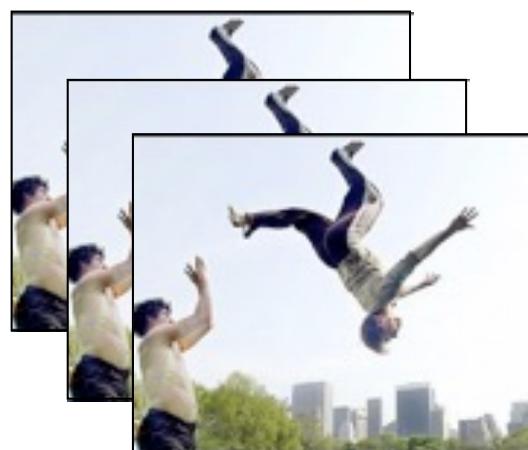
3. Bernoulli GLM: $y_t|\vec{x}_t, \vec{k} \sim \text{Ber}(f(\vec{x}_t \cdot \vec{k}))$ binary counts

log-likelihood: $\mathcal{L} = Y^\top \log f(X\vec{k}) - (1 - Y)^\top \log(1 - f(X\vec{k}))$

“logistic regression” if $f(x) = \frac{1}{1 + e^{-x}}$

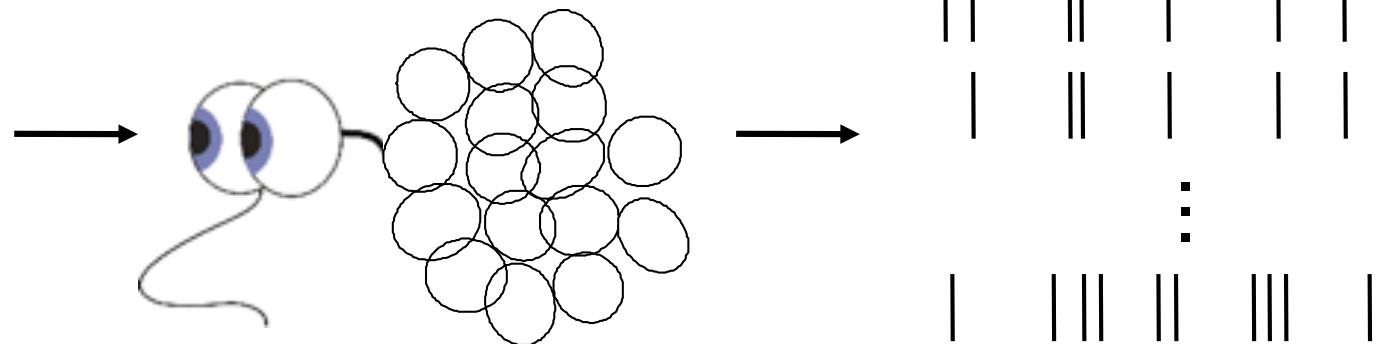
NEXT:

GLMs with spike-history and coupling



X

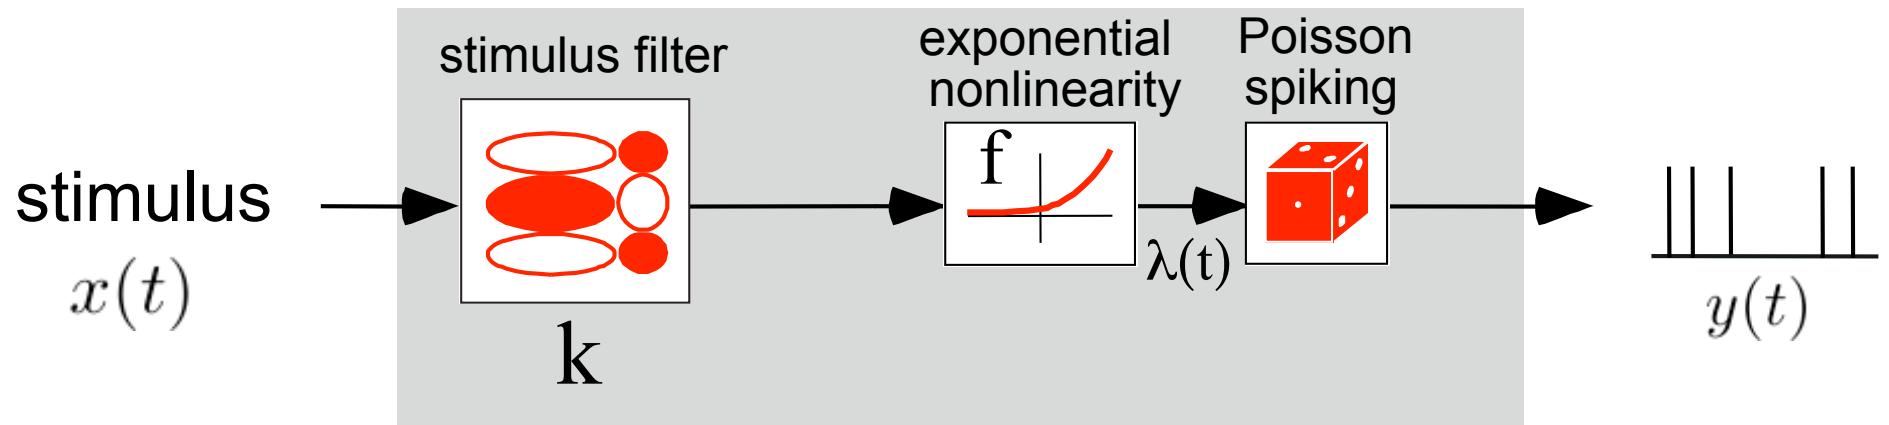
stimuli



y

spike trains

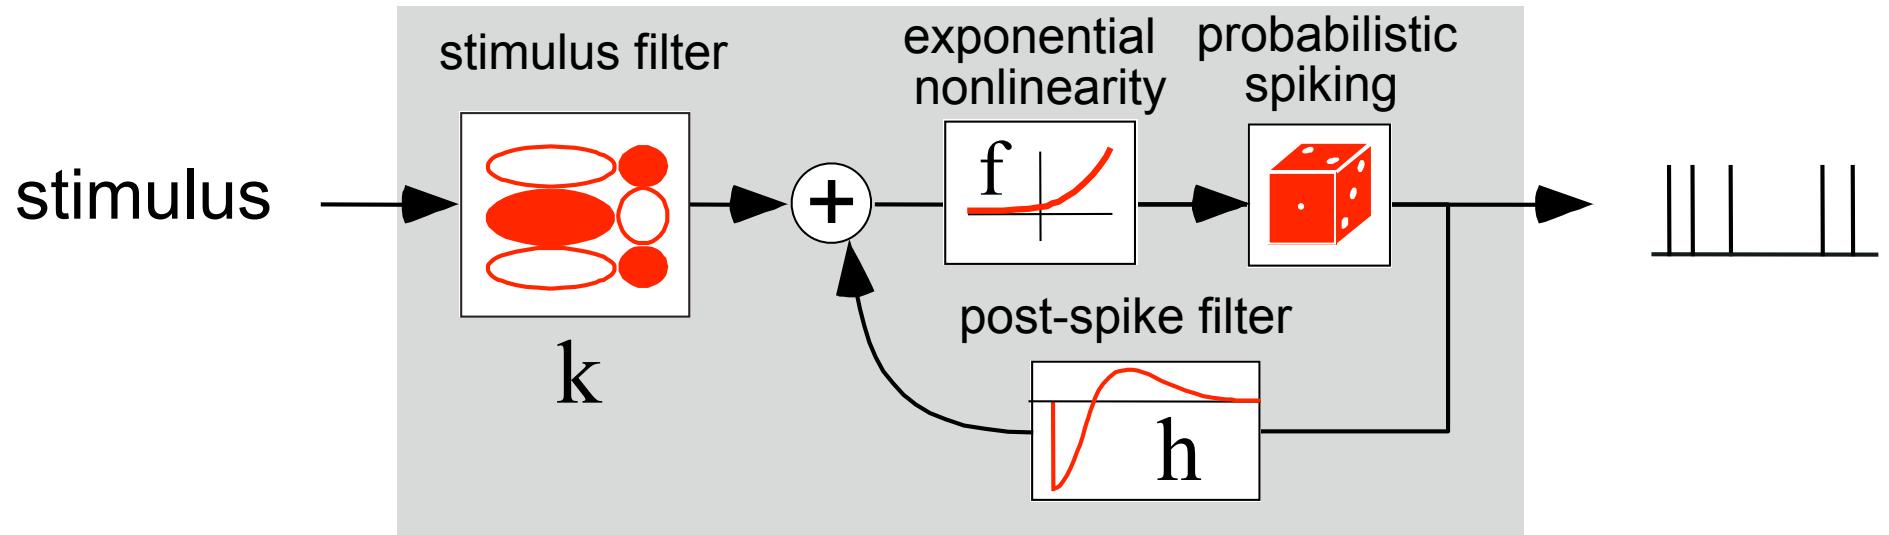
Poisson GLM



$$\text{spike rate } \lambda(t) = f(k \cdot x(t))$$

- problem: assumes spiking depends only on stimulus!

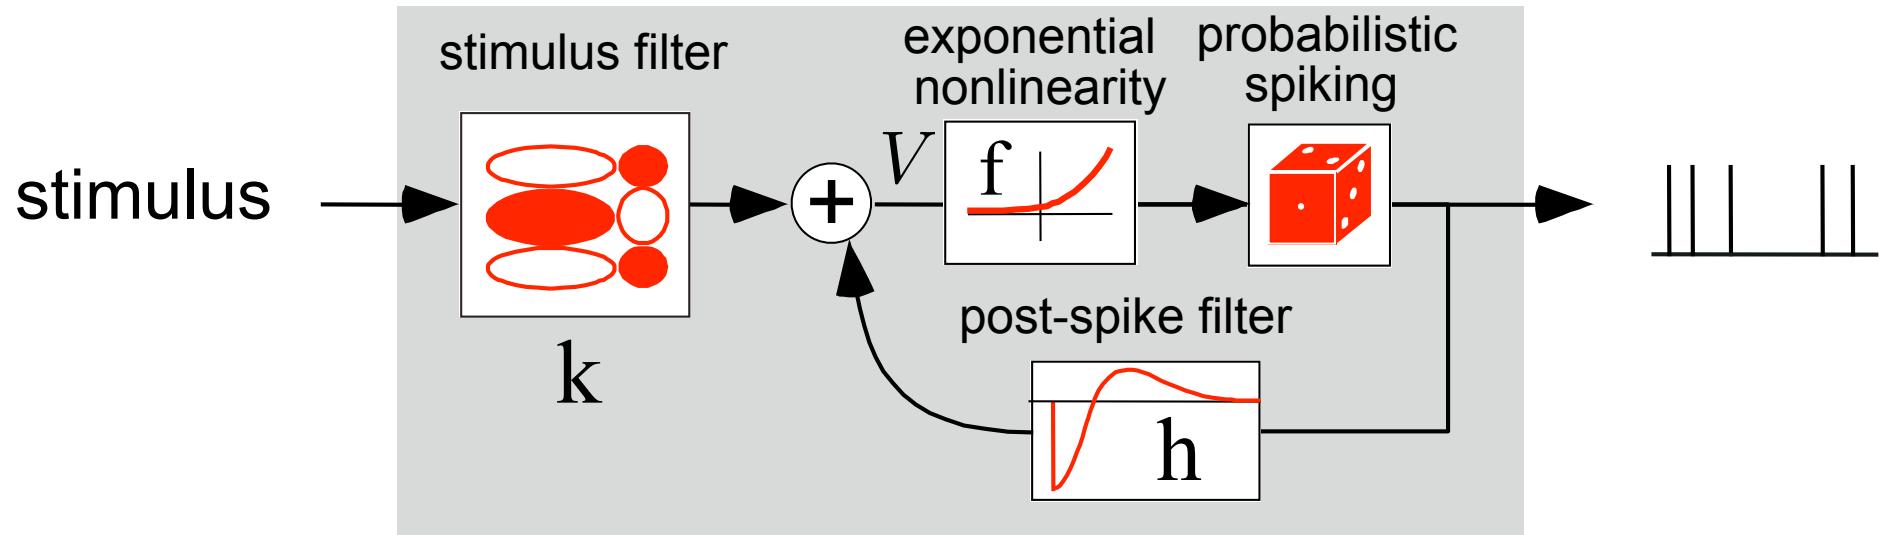
Poisson GLM with spike-history dependence



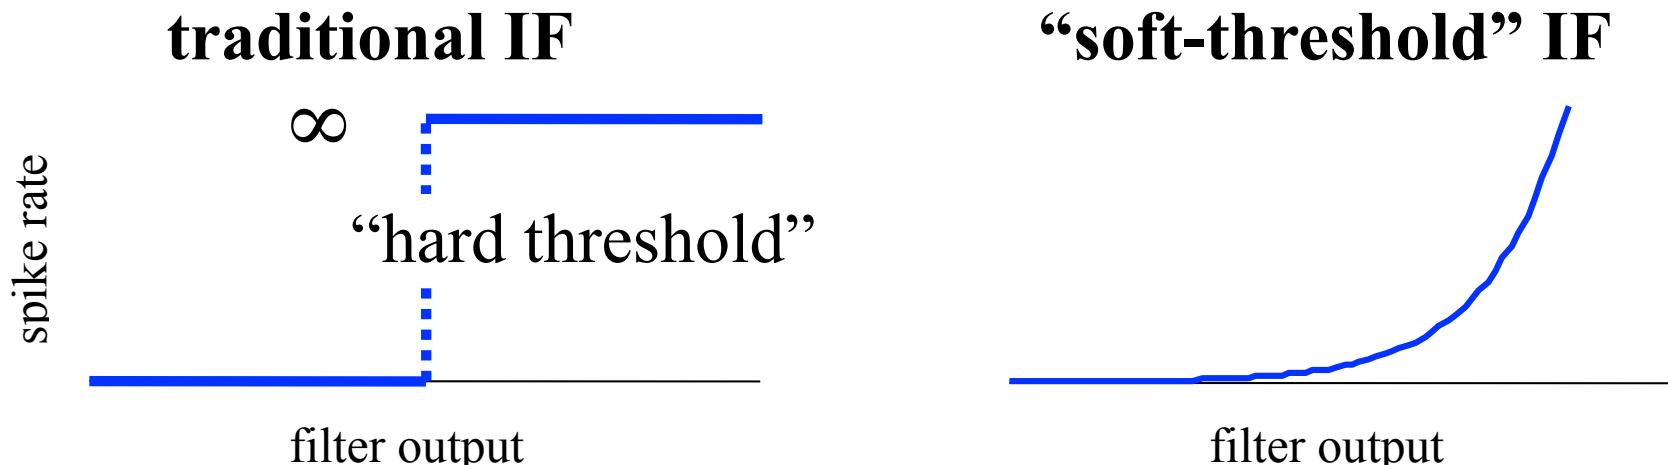
$$\begin{aligned} \text{spike rate: } \lambda(t) &= f(\vec{k} \cdot \vec{x}(t) + \vec{h} \cdot \vec{y}_{hst}(t)) \\ &= e^{\vec{k} \cdot \vec{x}(t)} \cdot e^{\vec{h} \cdot \vec{y}_{hst}(t)} \end{aligned}$$

- output: no longer a Poisson process
- interpretation: “soft-threshold” integrate-and-fire model

Poisson GLM with spike-history dependence

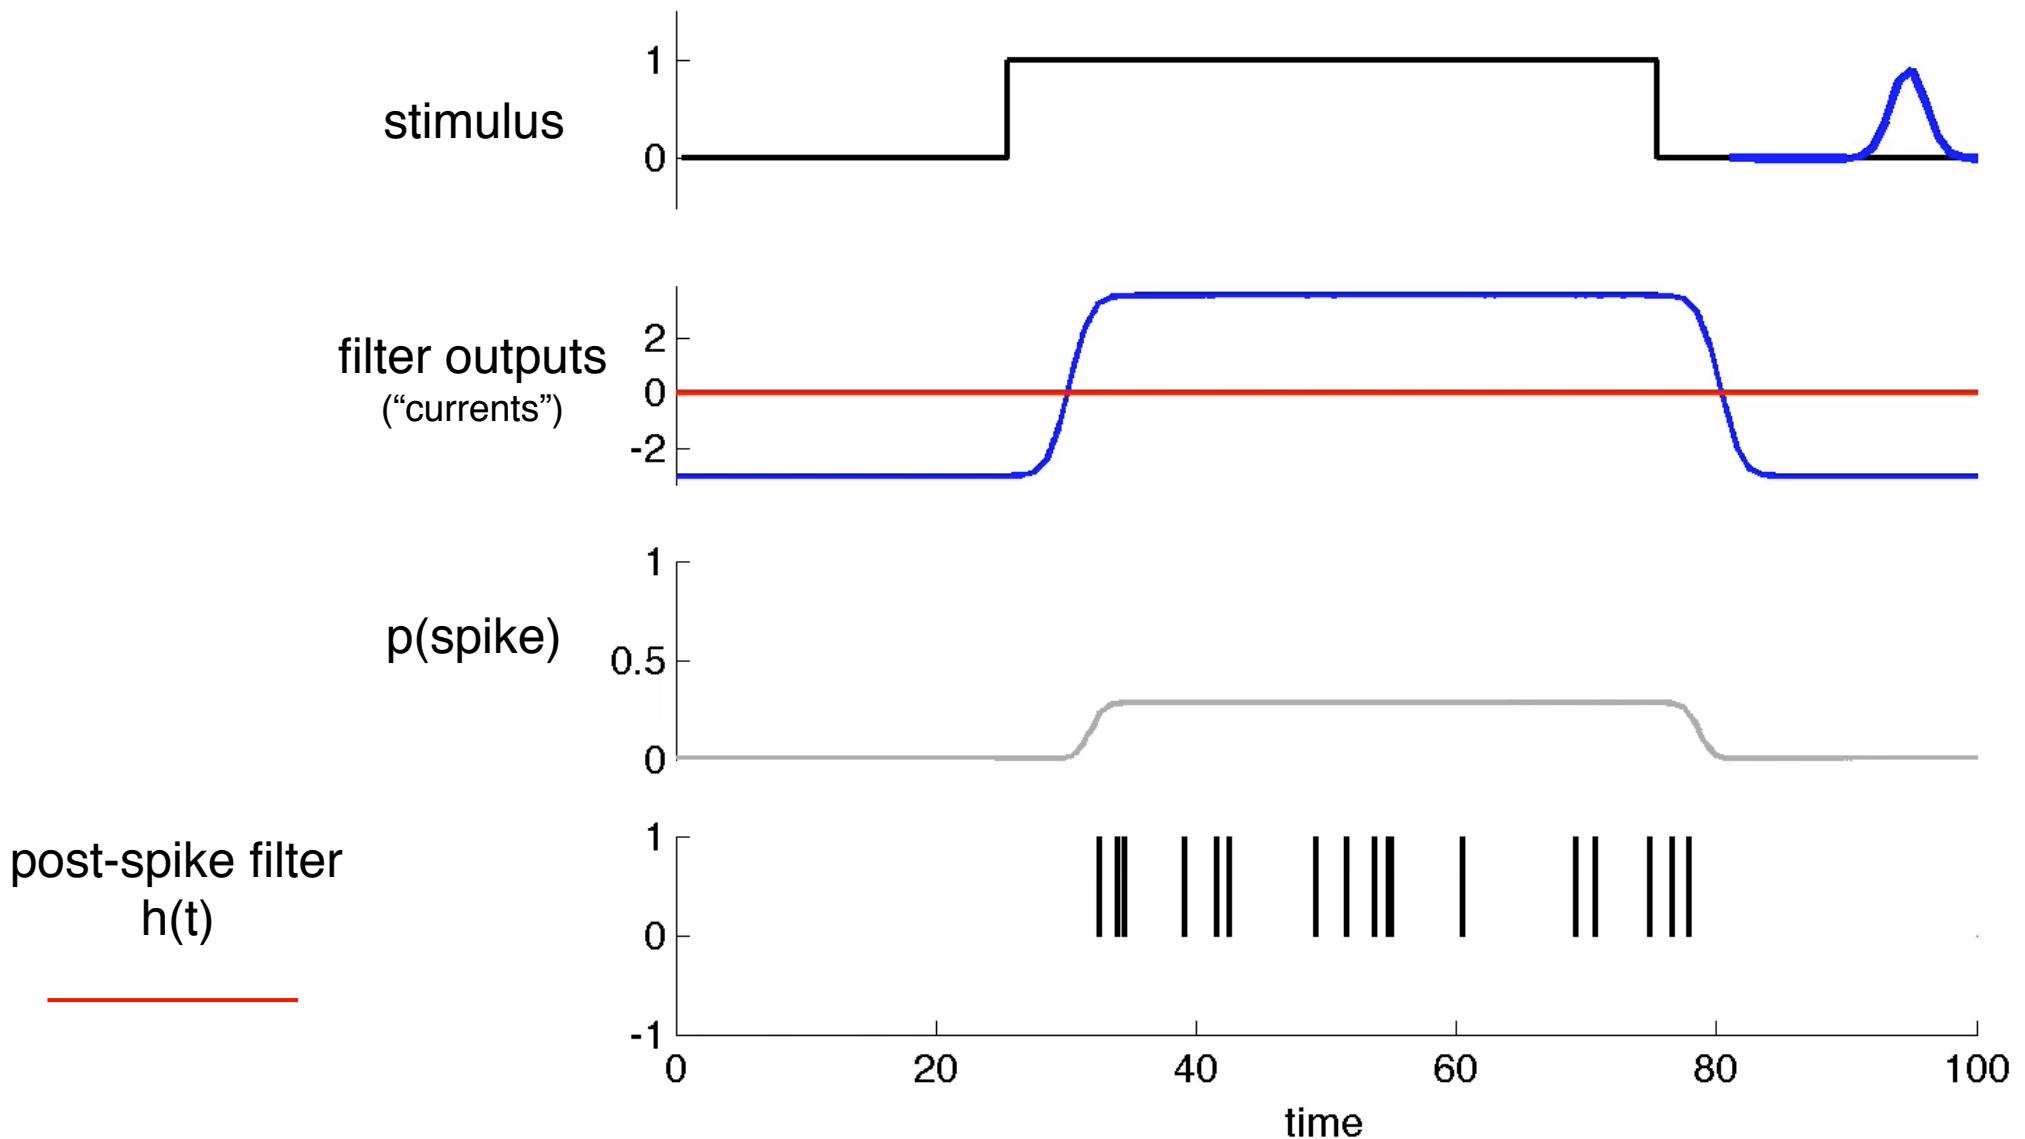


- interpretation: “soft-threshold” integrate-and-fire model



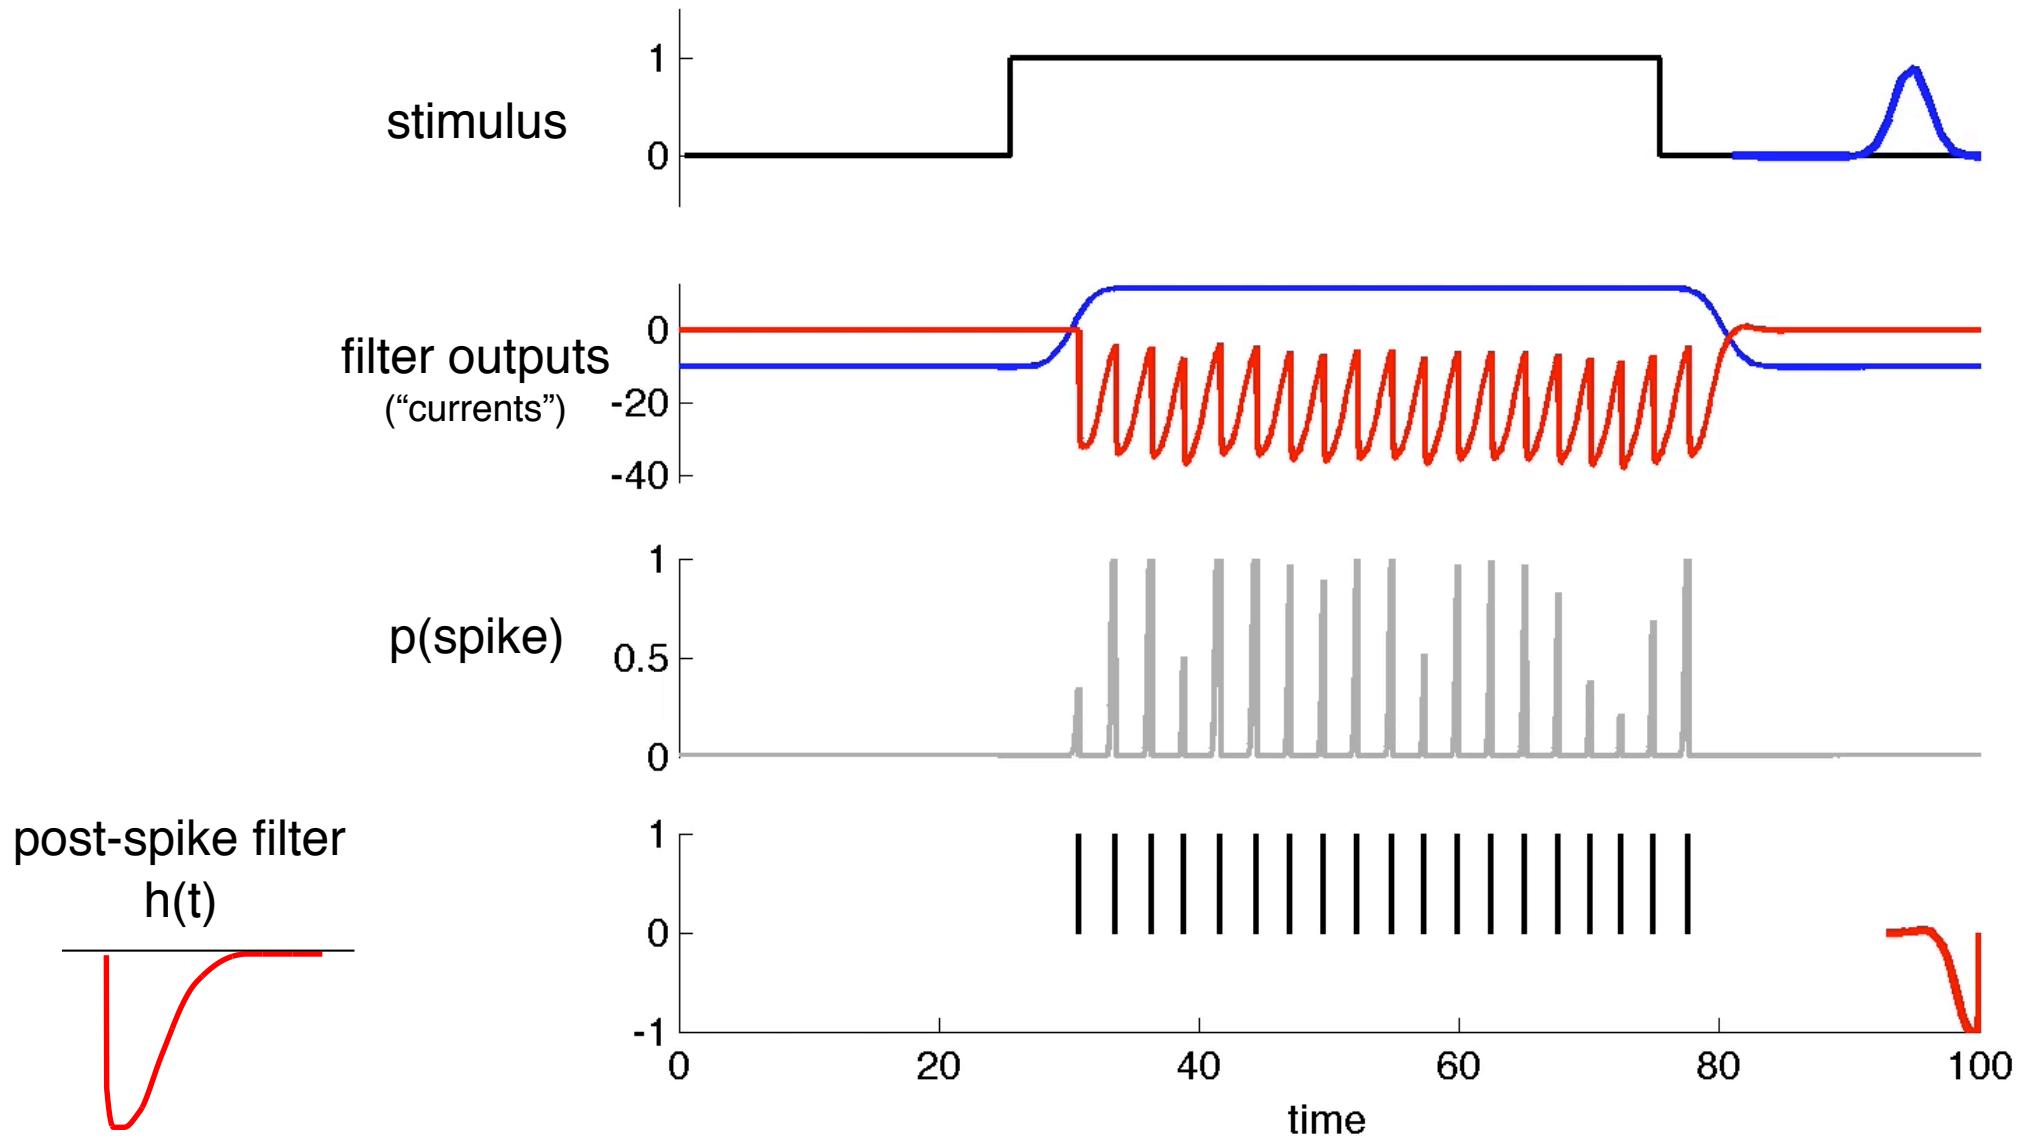
GLM dynamic behaviors

- irregular spiking



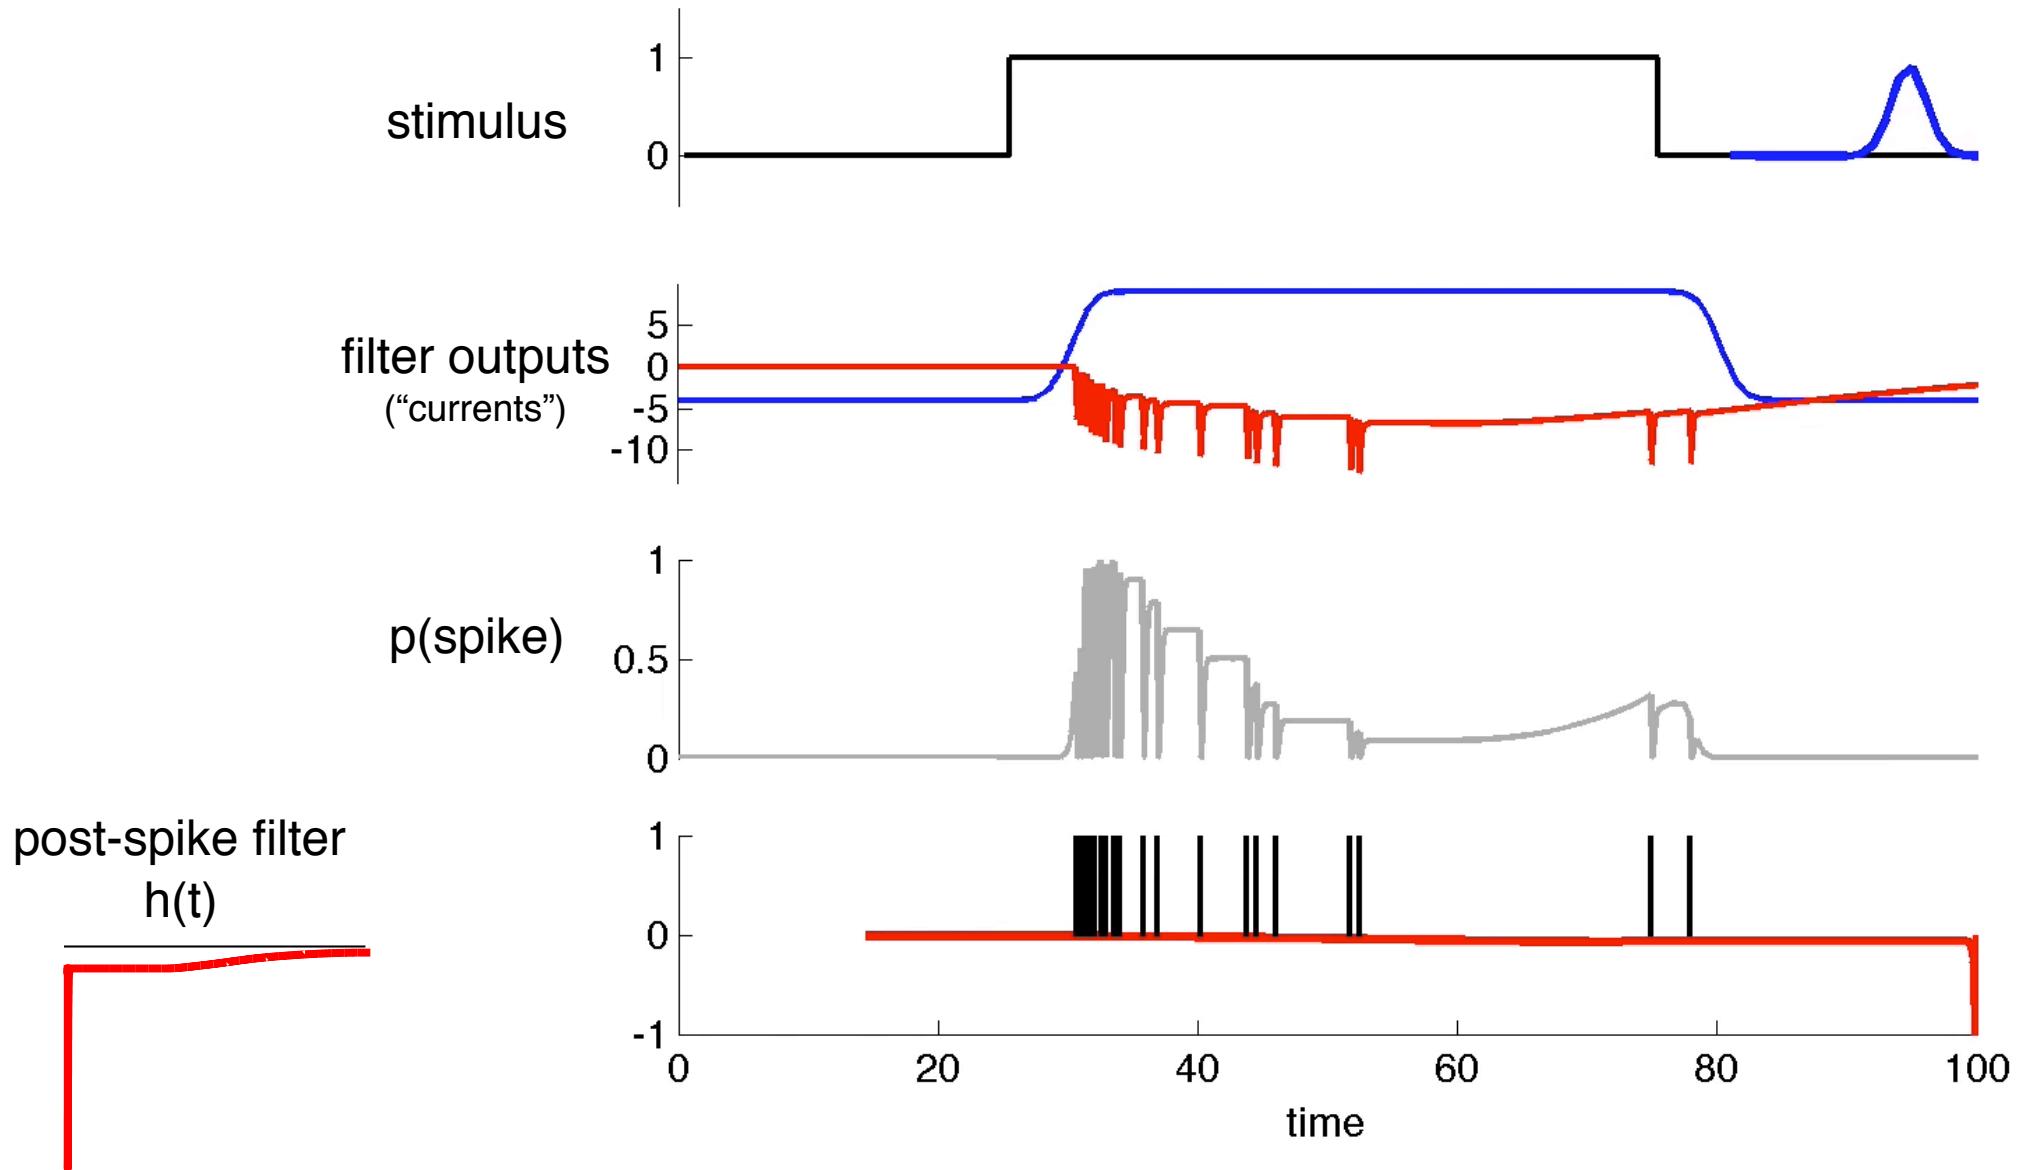
GLM dynamic behaviors

- regular spiking



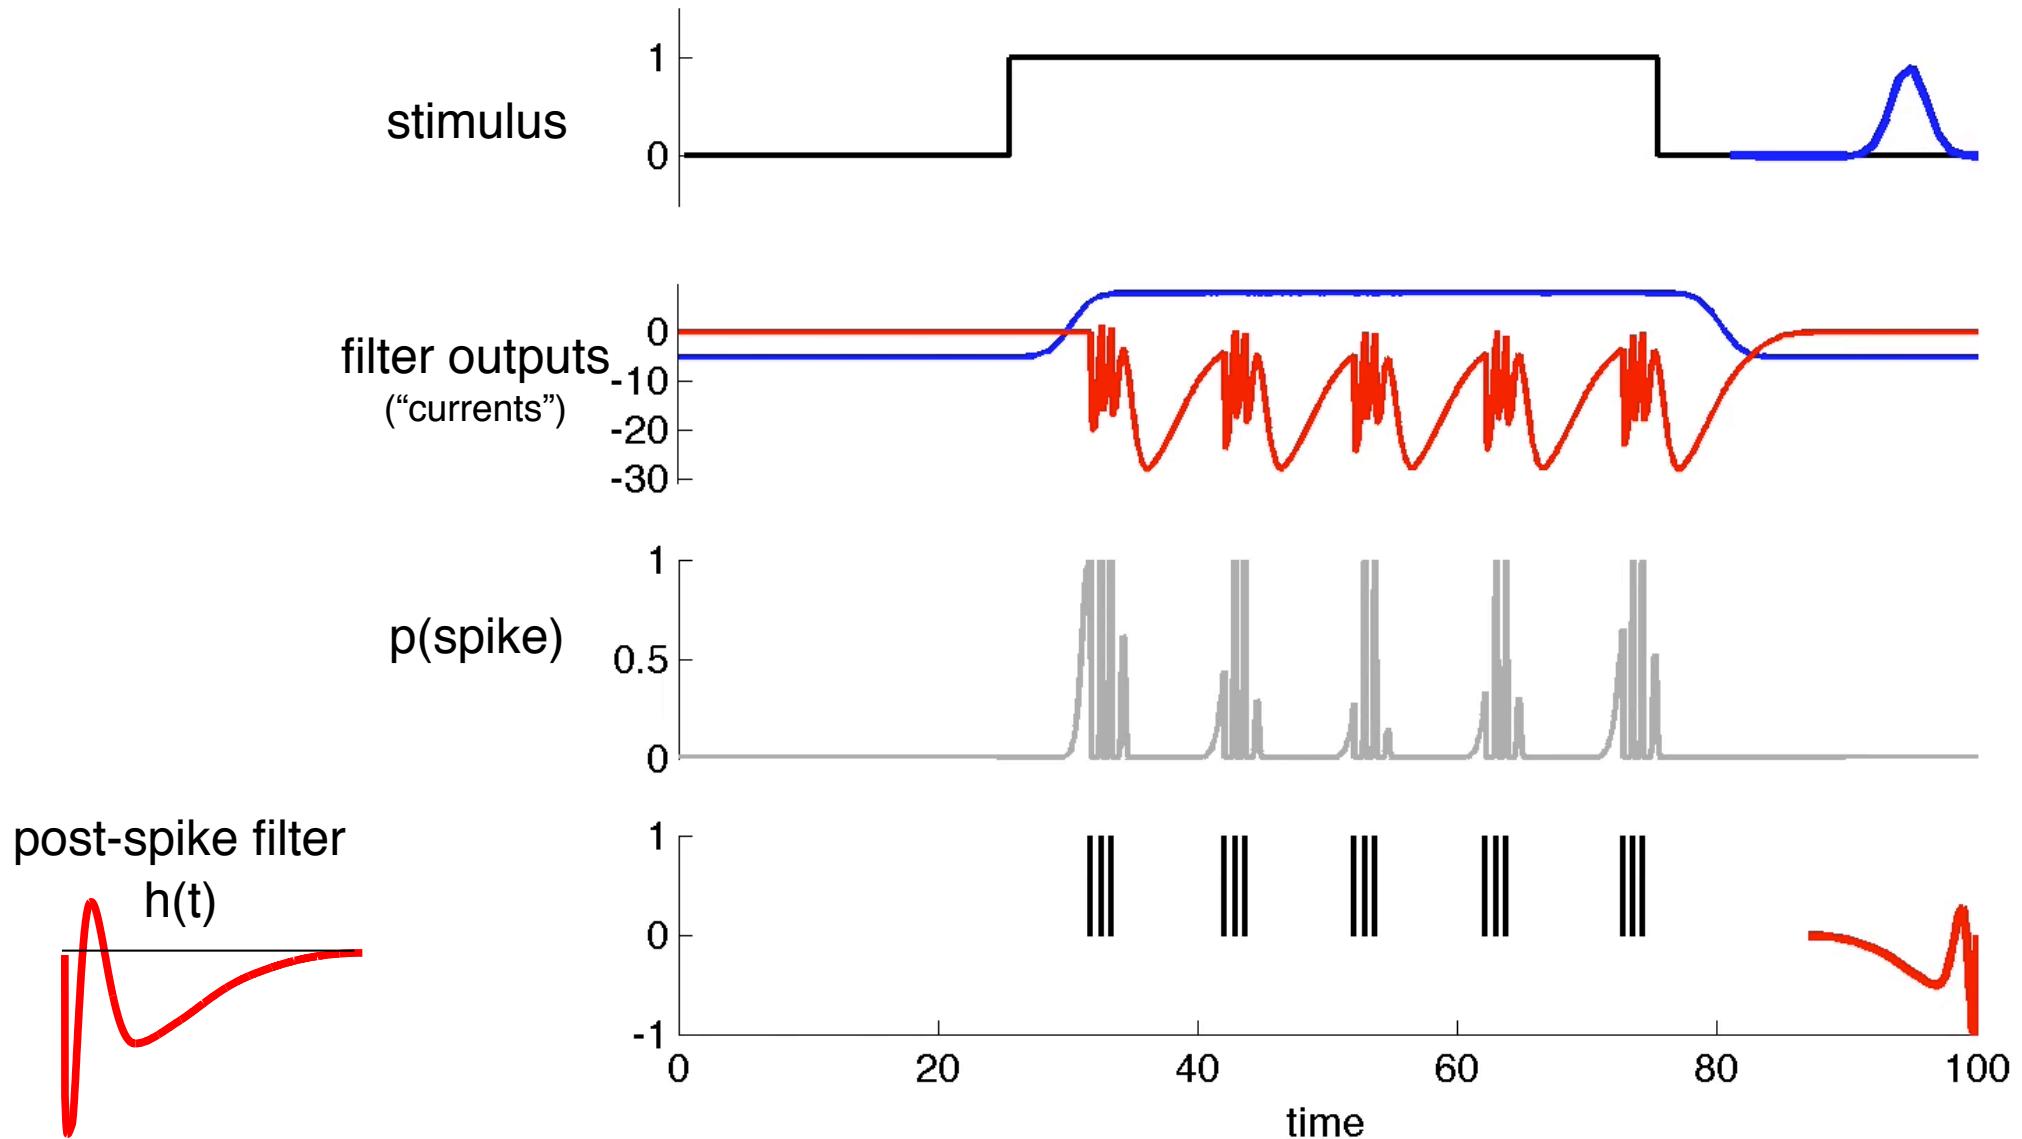
GLM dynamic behaviors

- adaptation



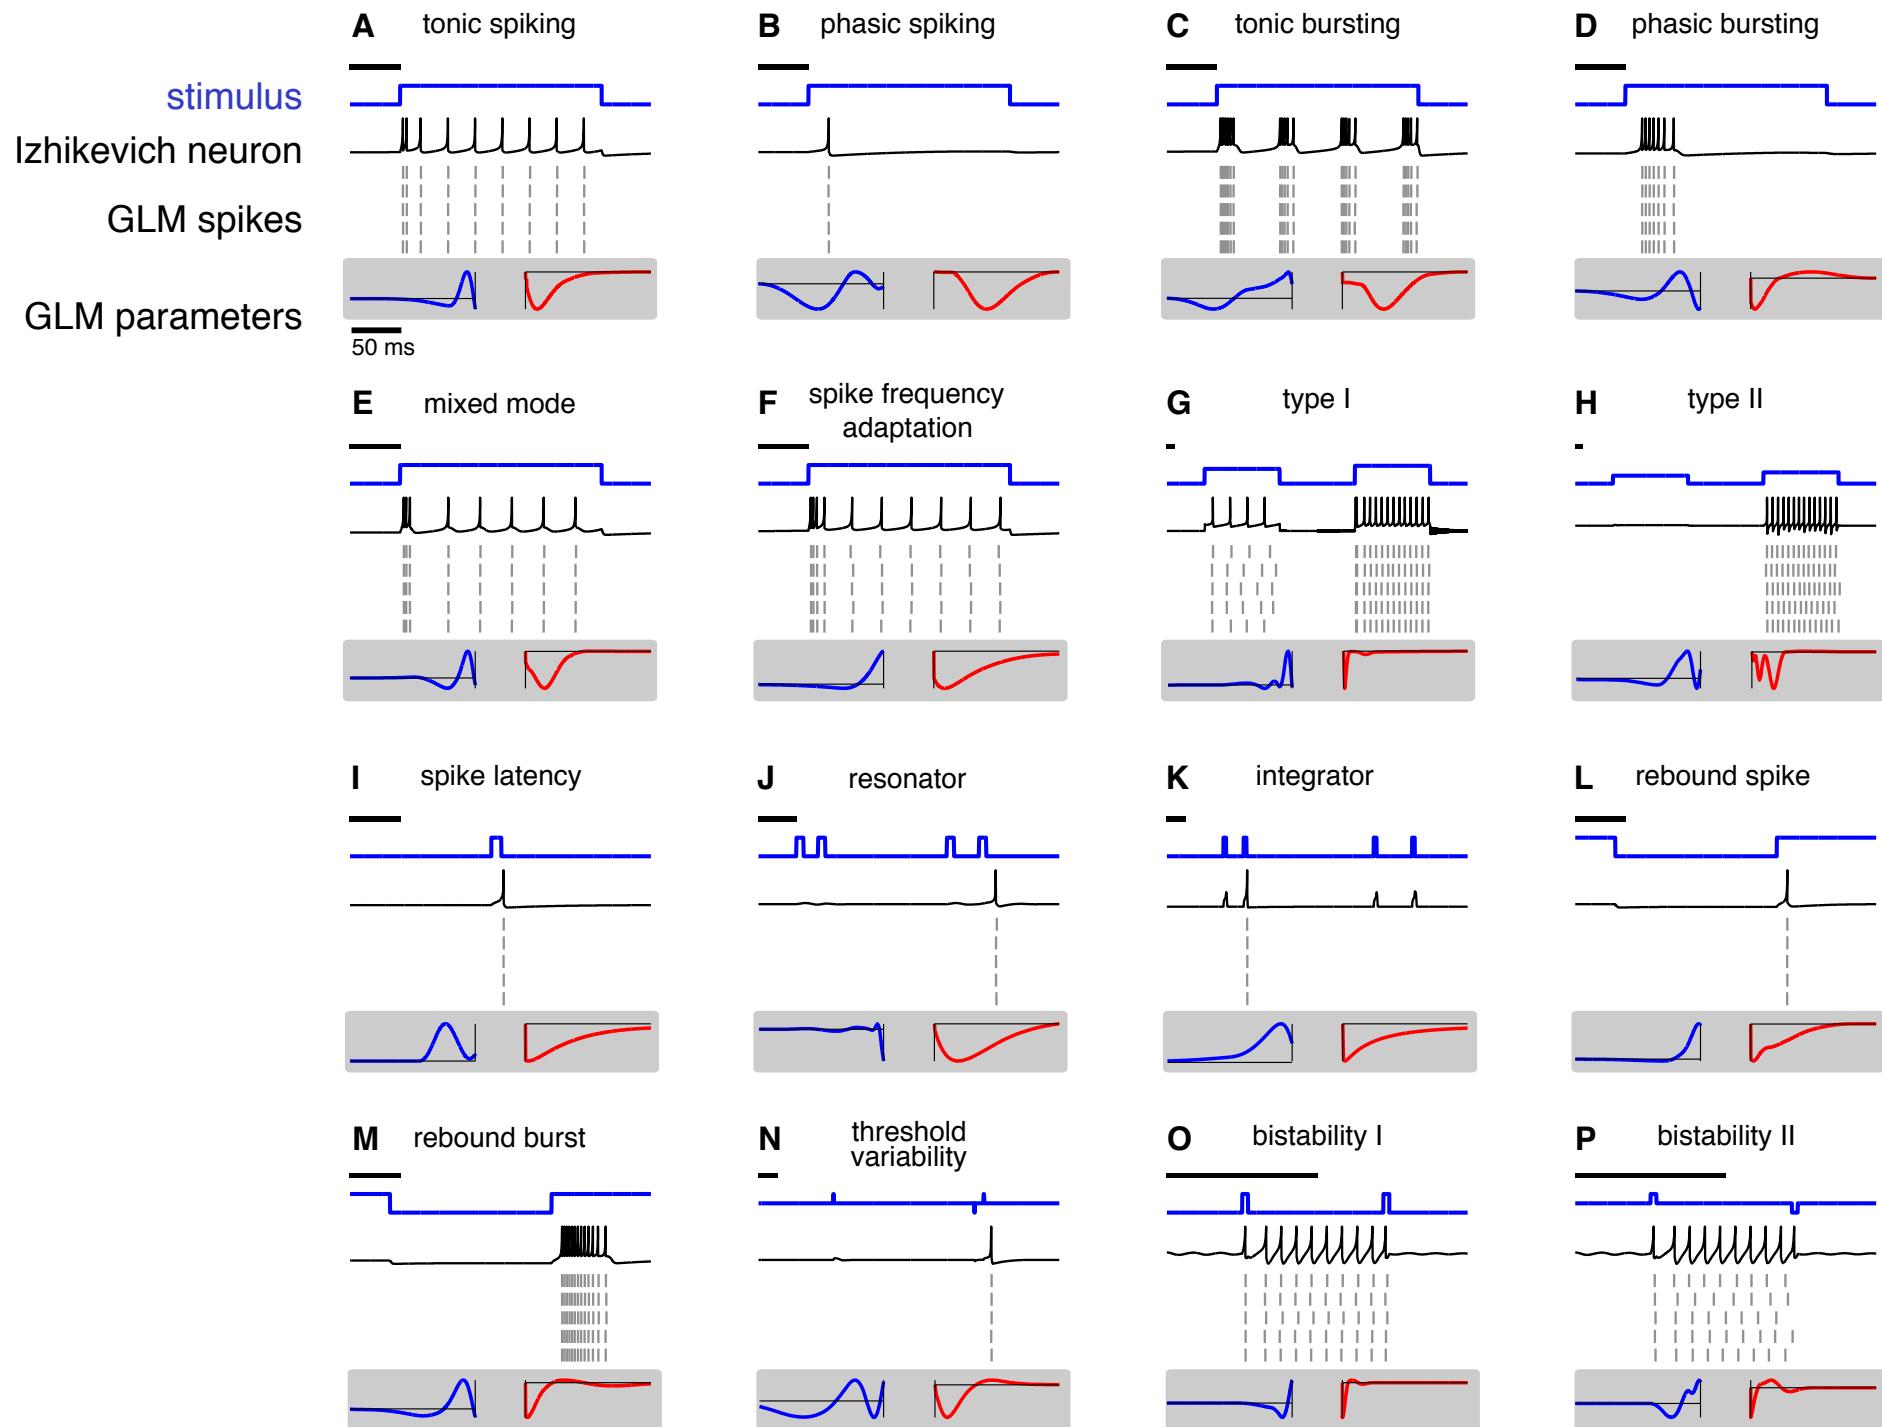
GLM dynamic behaviors

- bursting

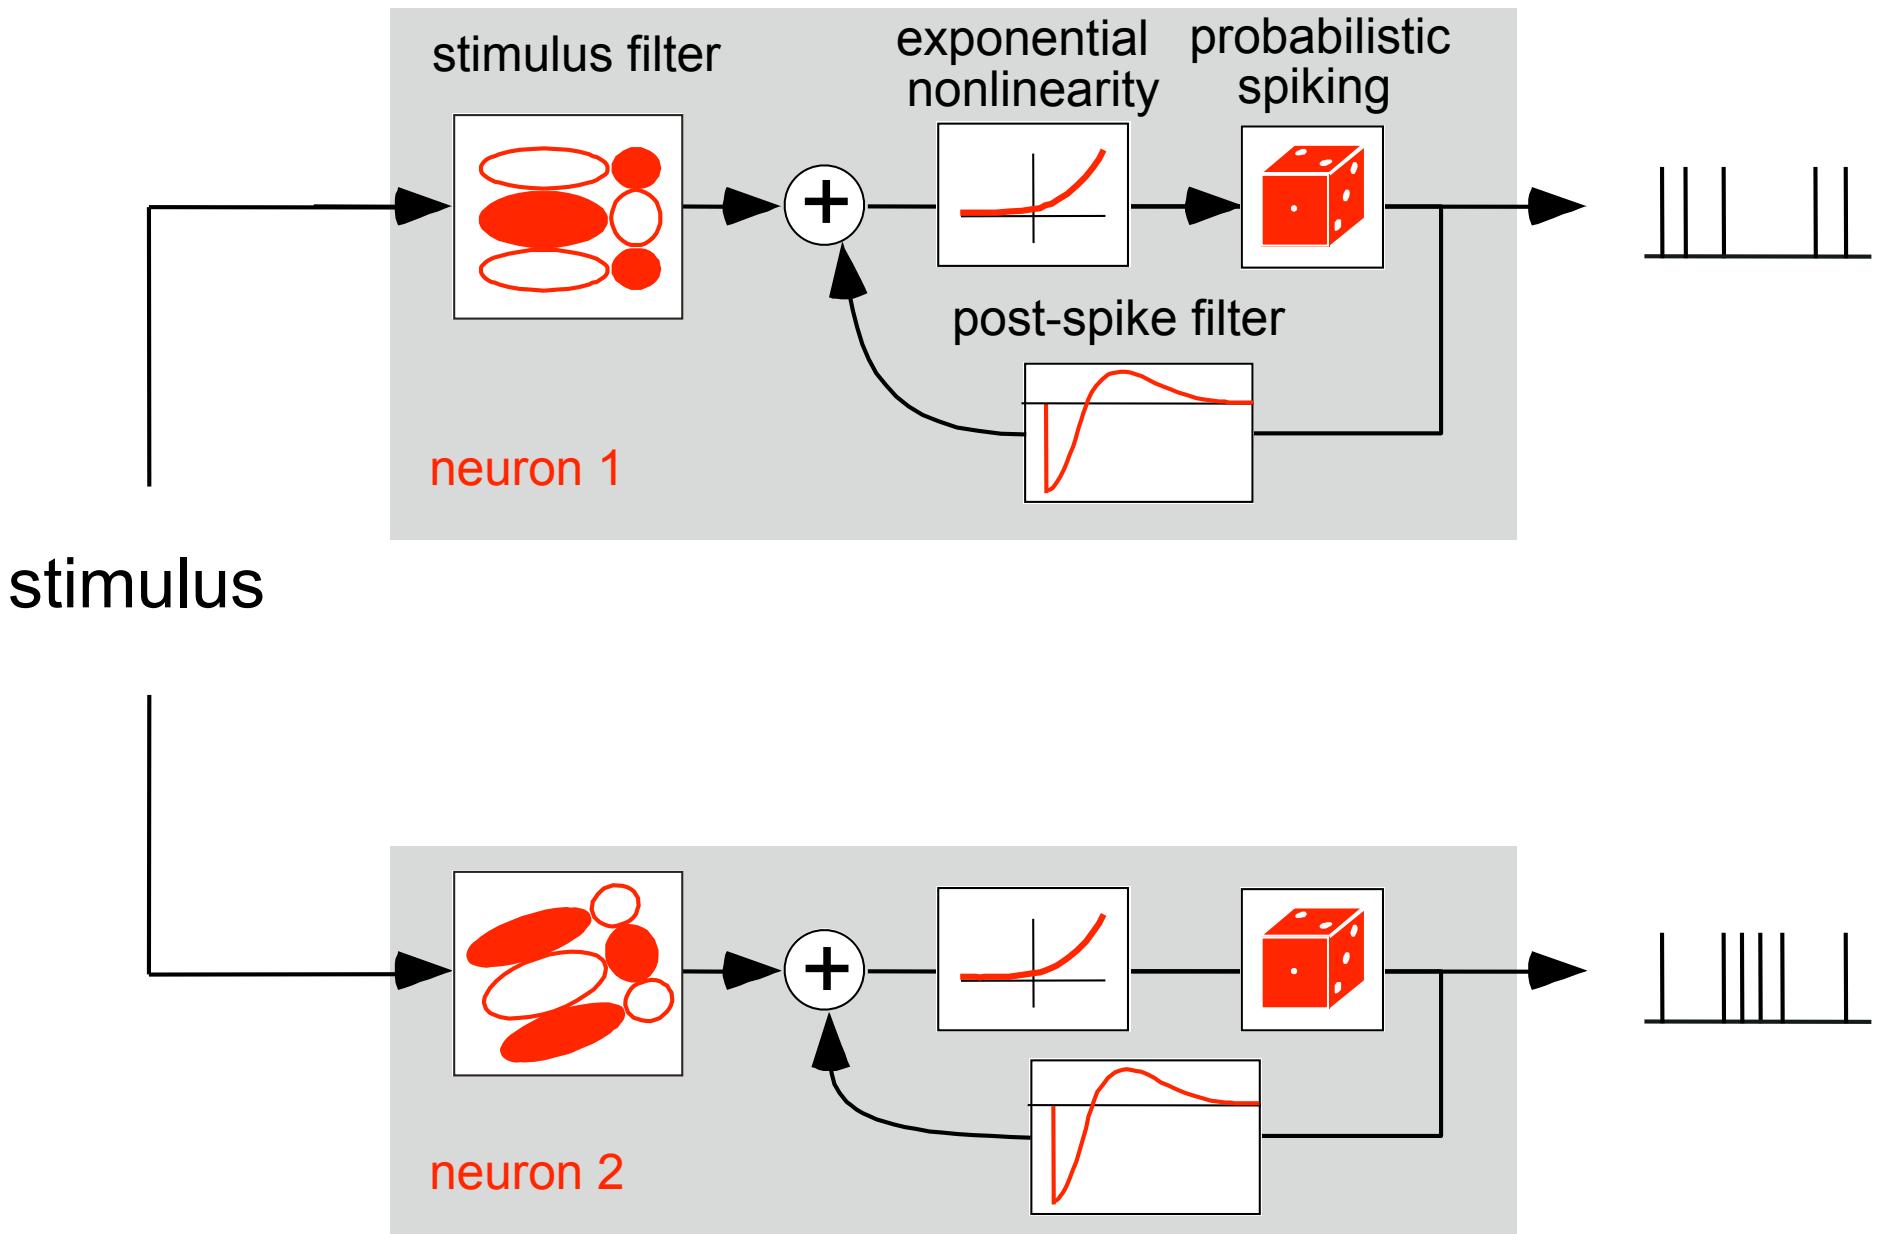


GLM dynamic behaviors (from Izhikevich)

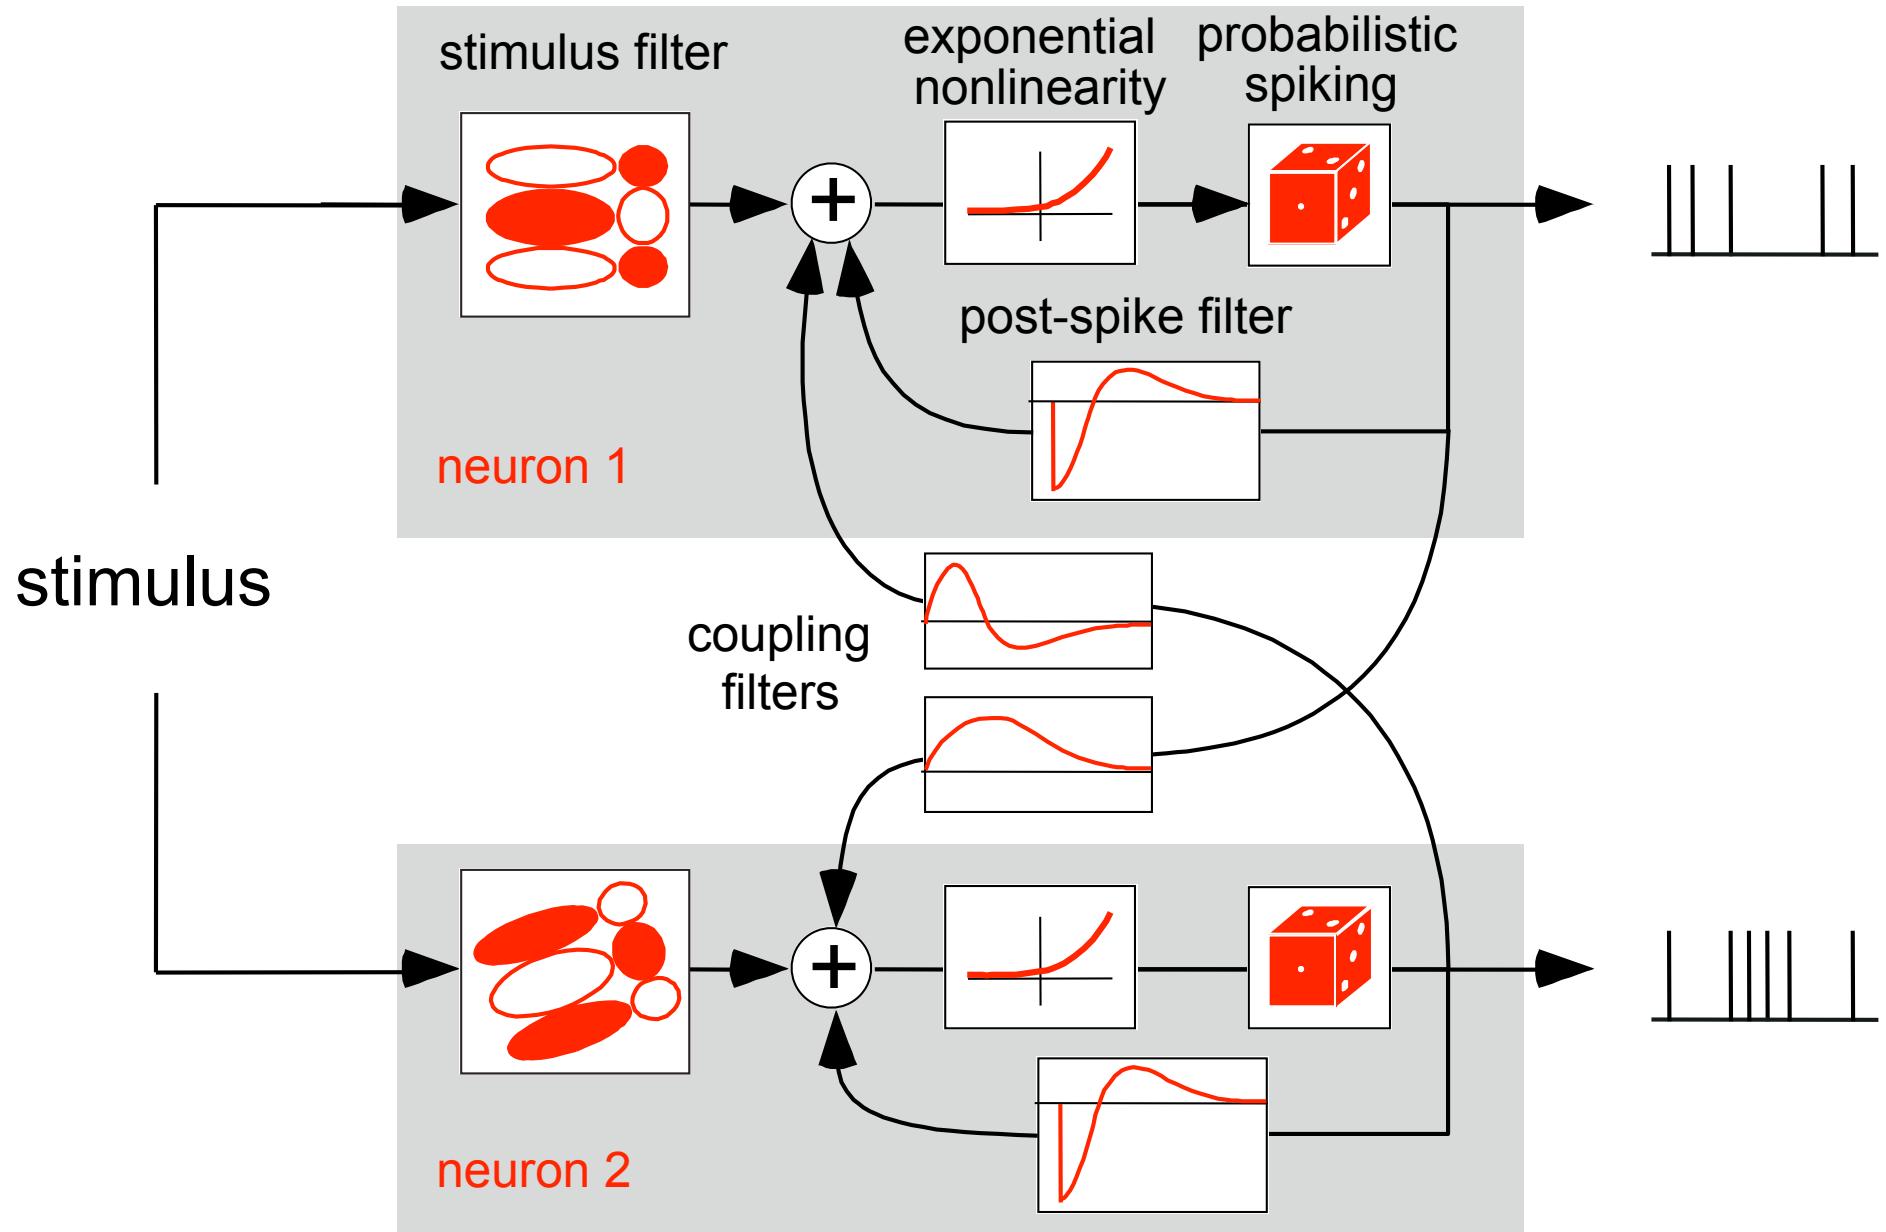
(Weber & Pillow 2017)



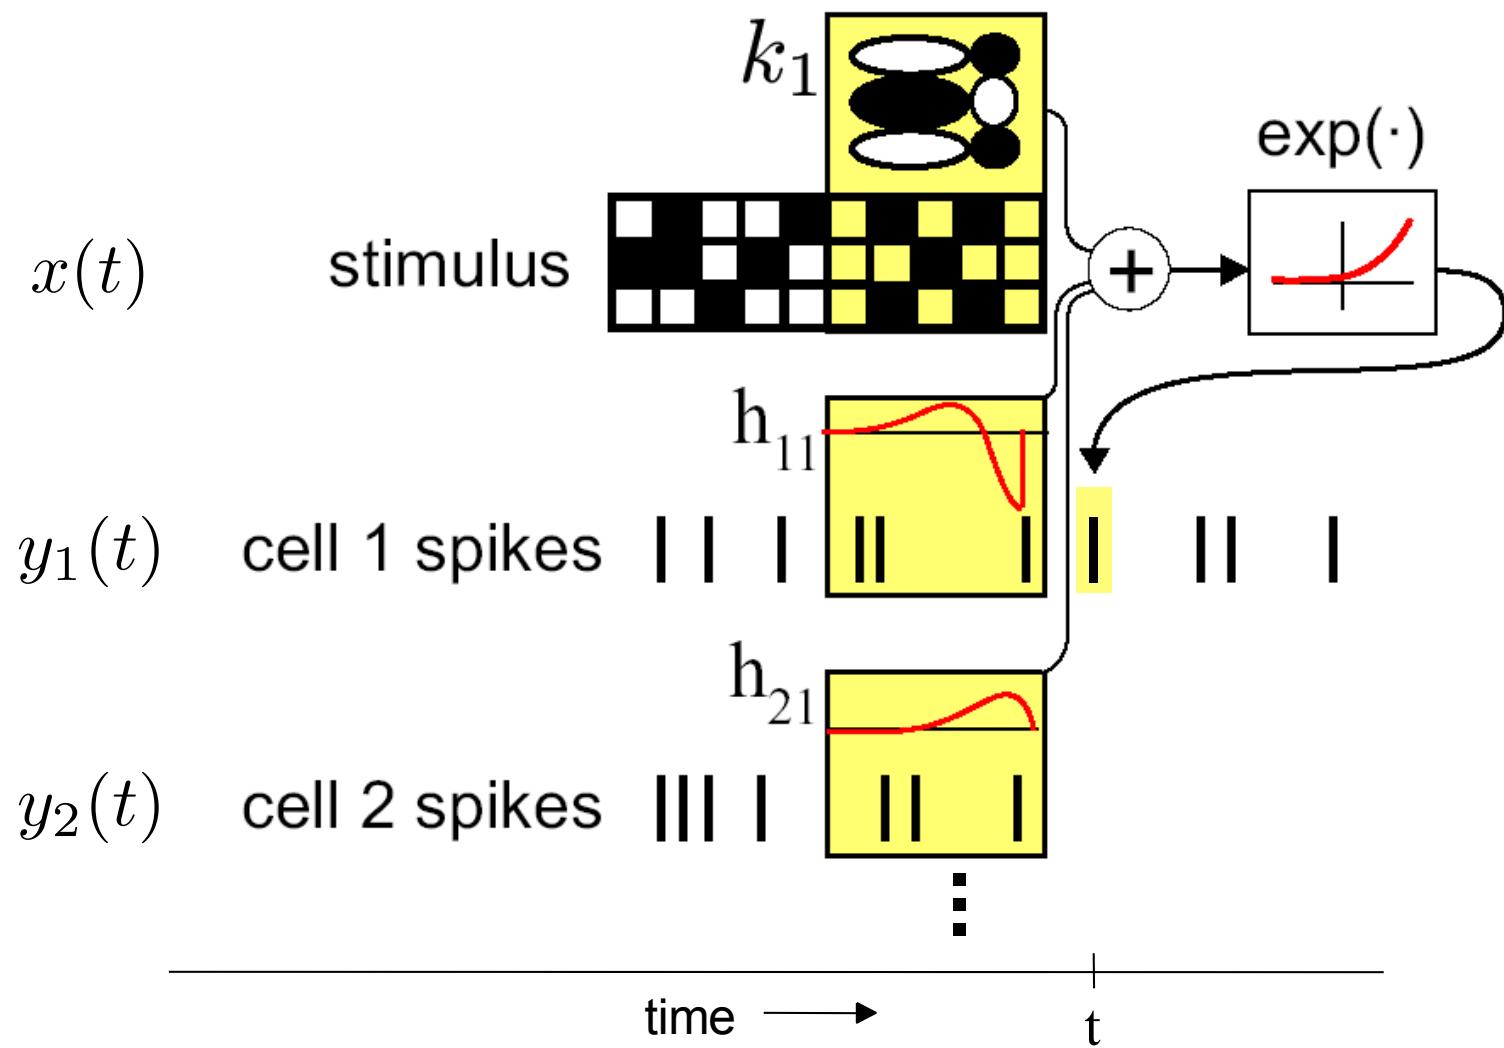
multi-neuron GLM



multi-neuron GLM



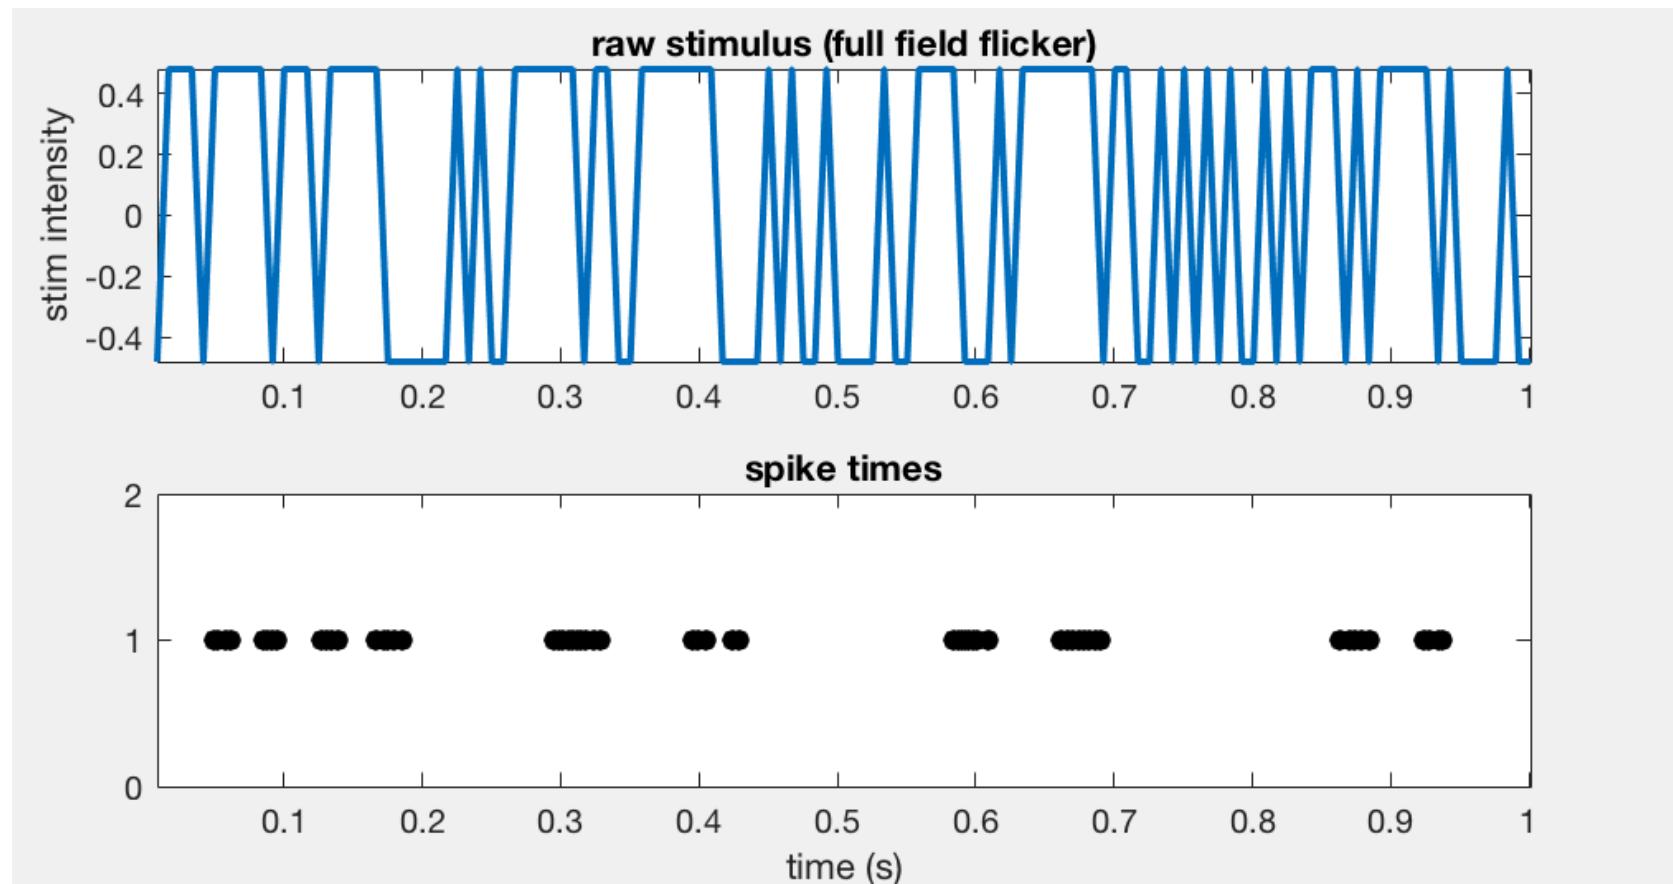
GLM equivalent diagram:



$$\text{spike rate} \quad \lambda_i(t) = \exp(k_i \cdot x(t) + \sum_j h_{ij} \cdot y(t))$$

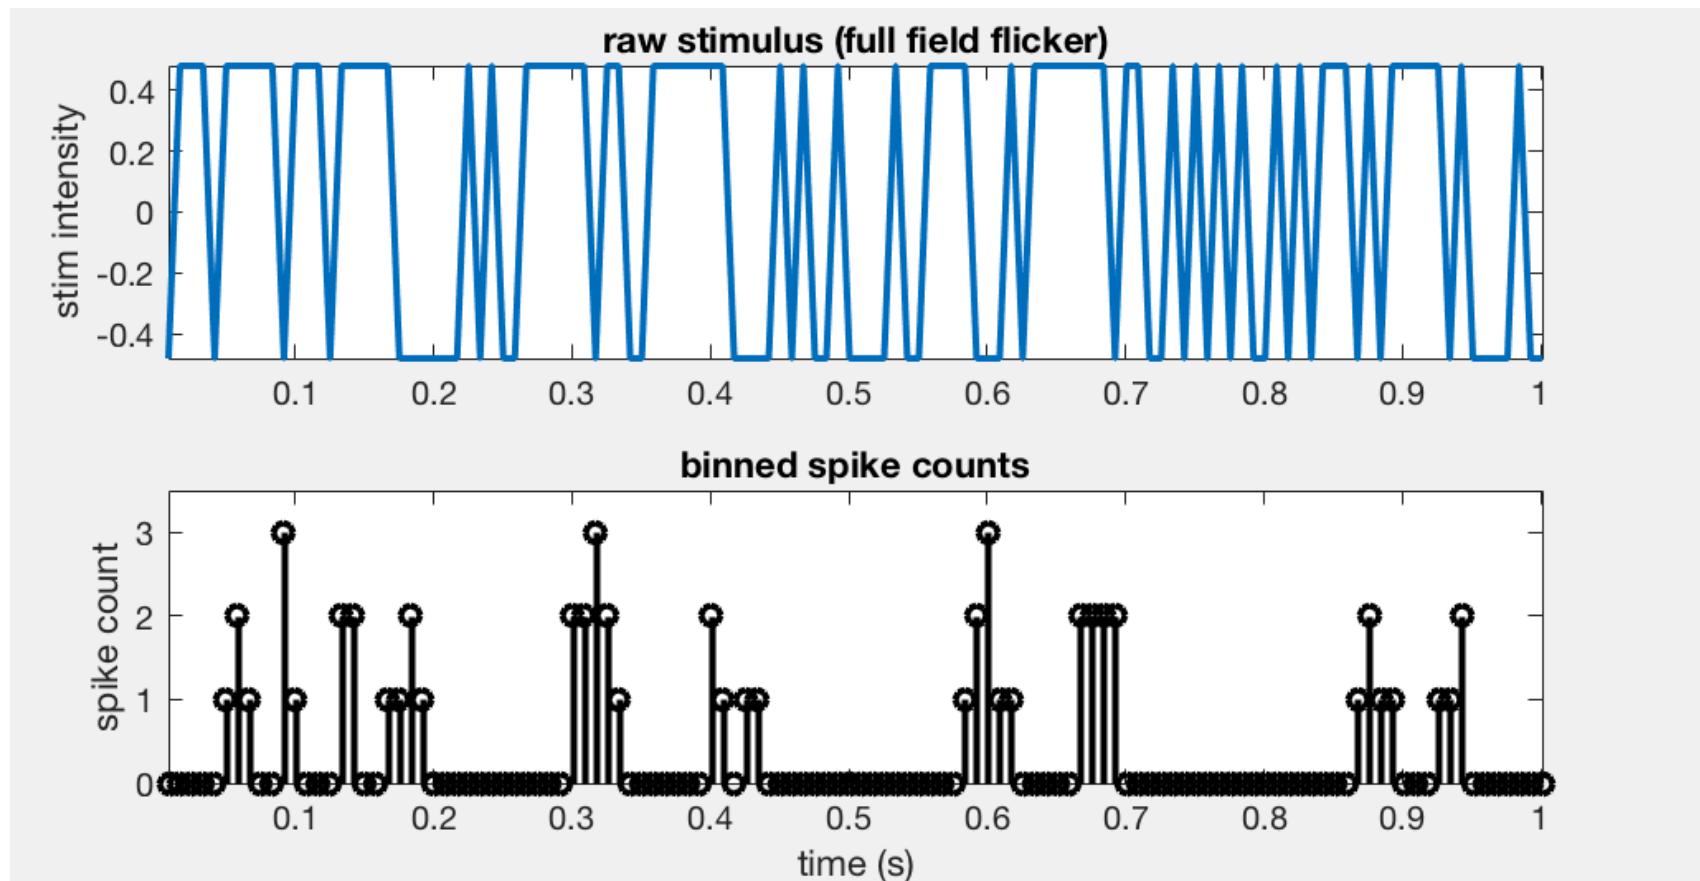
Example dataset

- stimulus = binary flicker
- parasol retinal ganglion cell spike responses



Example dataset

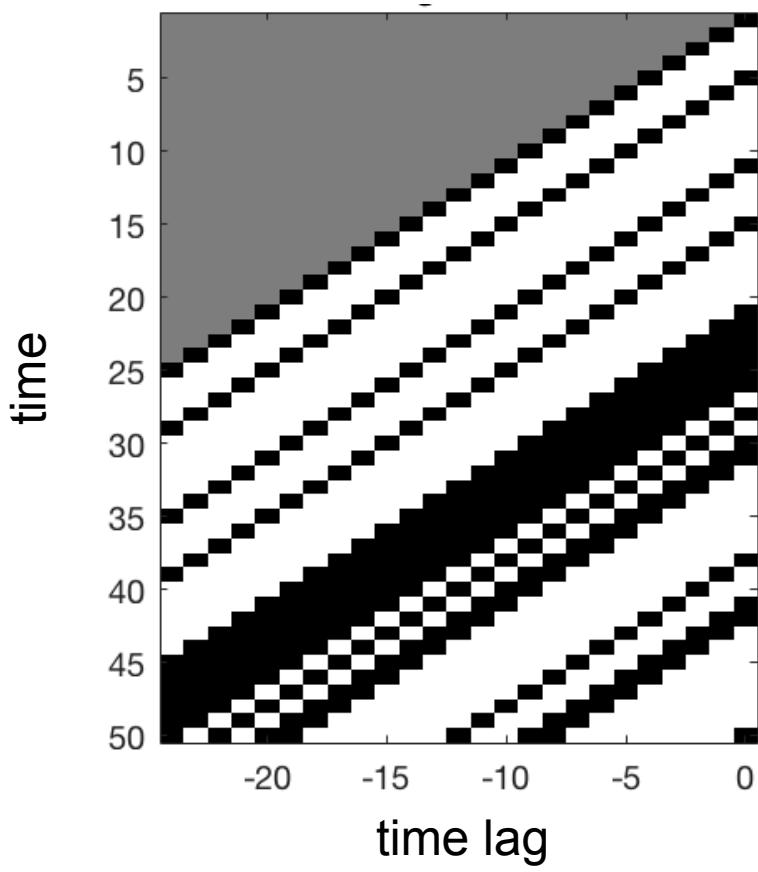
- stimulus = binary flicker
- parasol retinal ganglion cell spike responses



Stimulus-only GLM

design matrix

X



spike response

Y

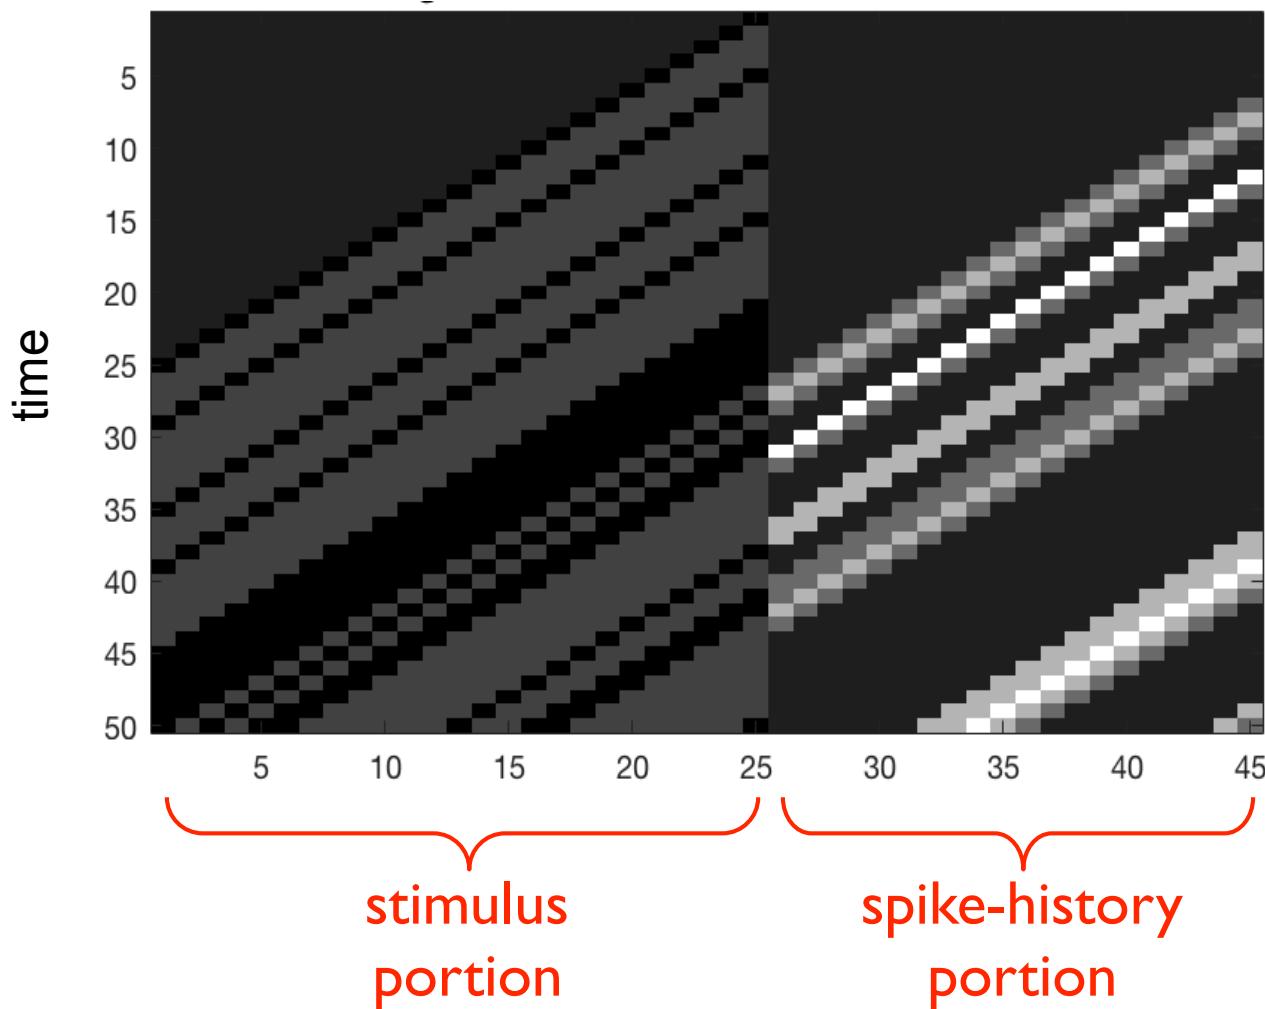
model
 $P(Y|X)$



Stimulus + SpikeHistory GLM

design matrix

X



spike response

Y

model

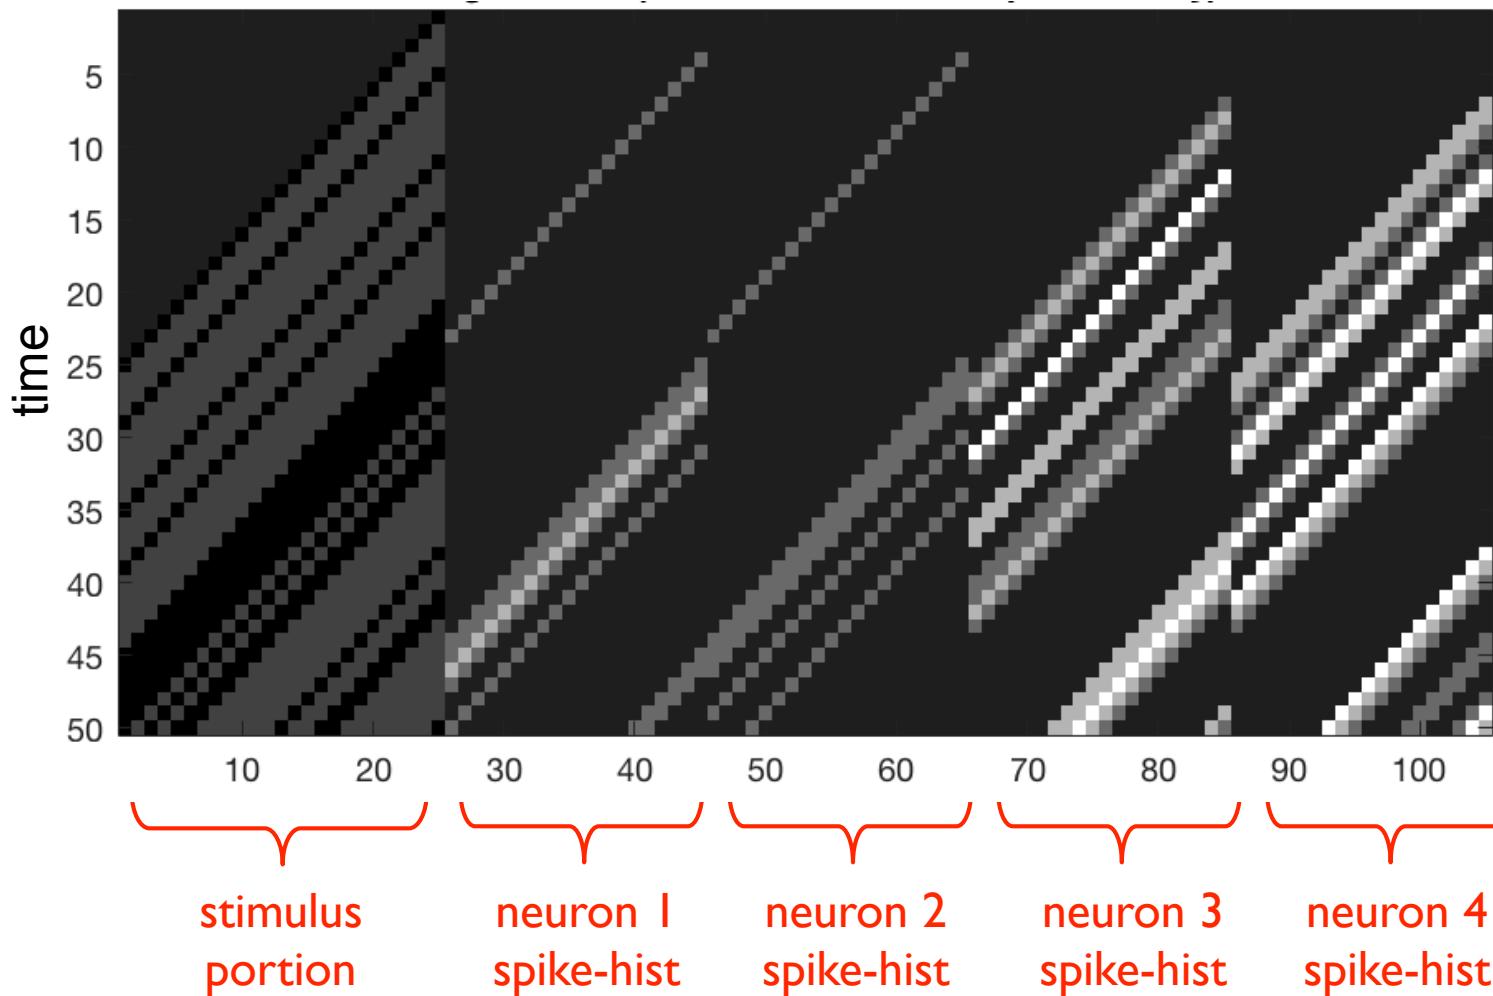
$$P(Y|X)$$



Stimulus + History + 3 Neuron Coupling GLM

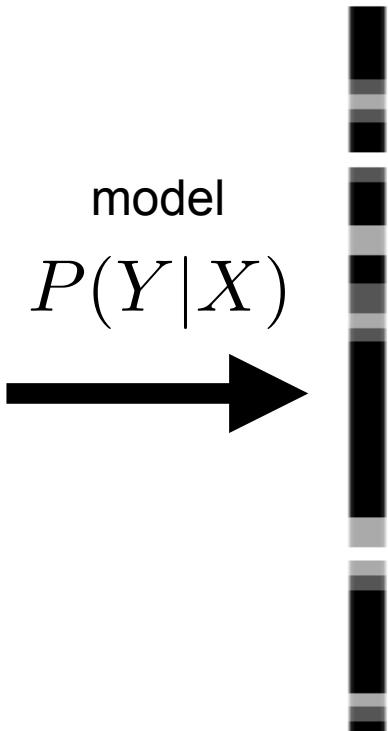
design matrix

X

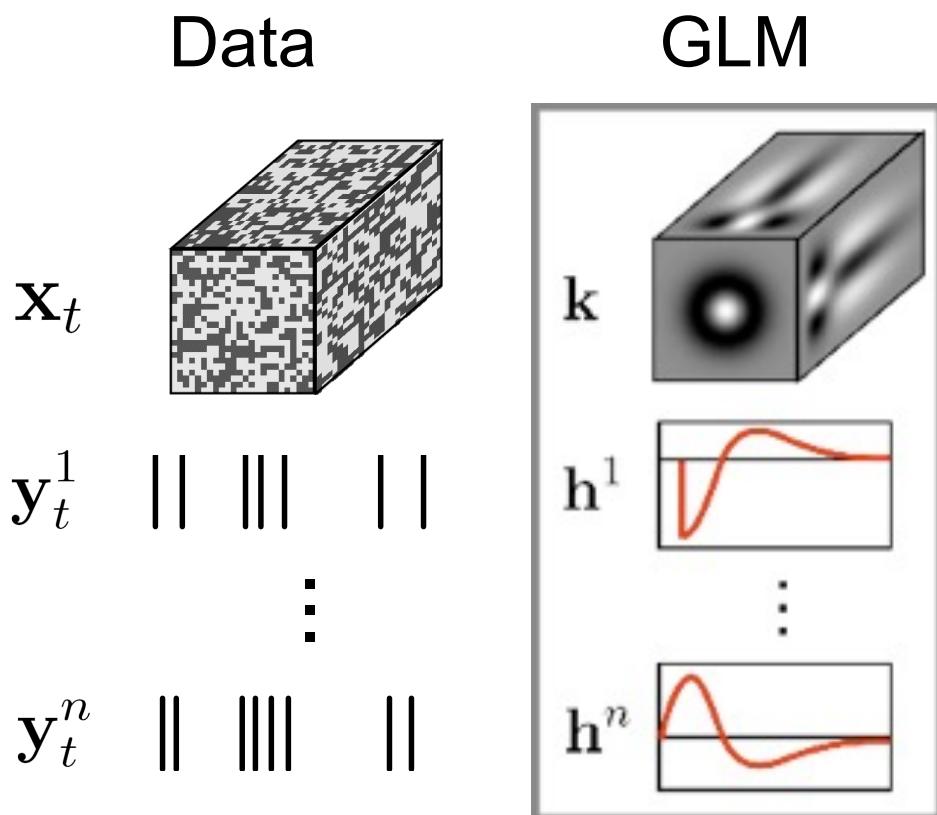


spike response

Y



Fitting: Maximum Likelihood

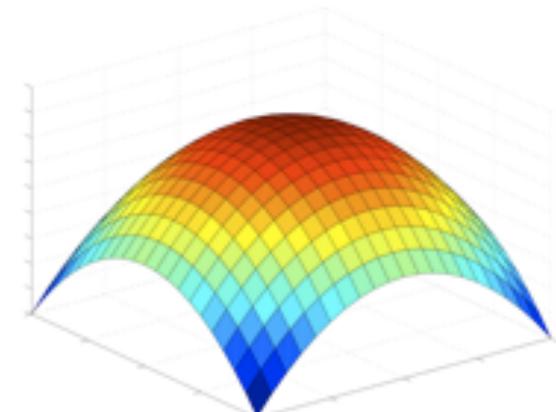


- maximize log-likelihood for filters $\{k, h_1, h_2, \dots h_n\}$

$$\text{firing rate: } \lambda_t = f(\vec{x}_t \cdot \vec{k})$$

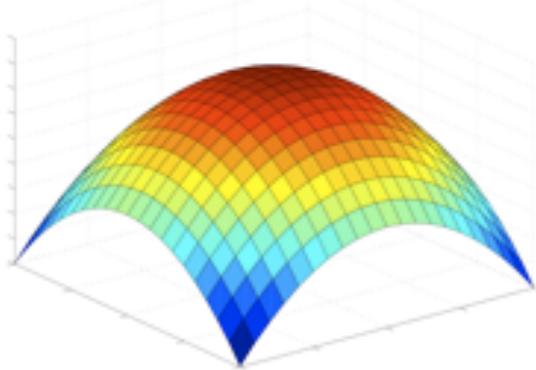
$$\log P(Y|X) = \sum_t y_t \log \lambda_t - \lambda_t$$

- log-likelihood is concave
 - no local maxima [Paninski 04]



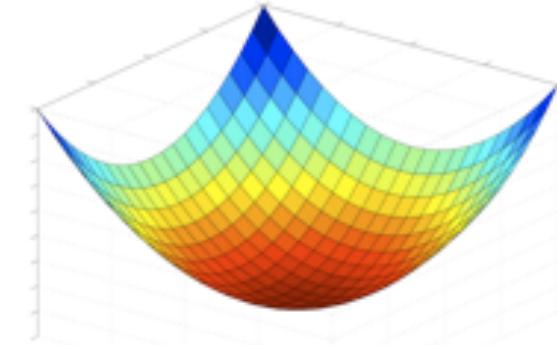
convexity and concavity

concave



- everywhere downward curvature

convex



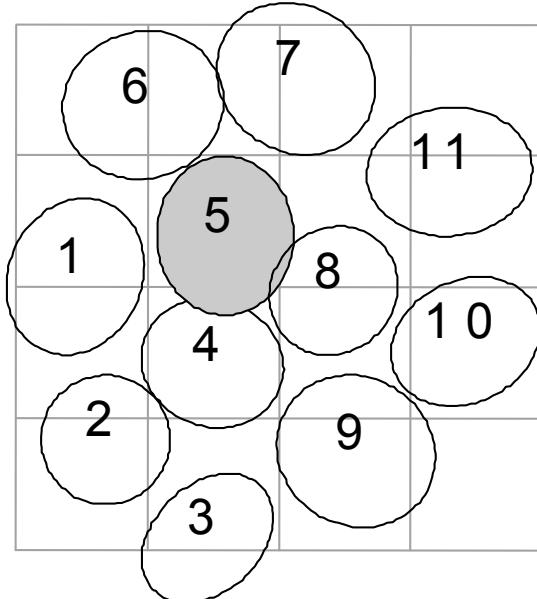
- everywhere upward curvature

- maximizing concave function \iff minimizing a convex function
- preclude existence of non-global local optima

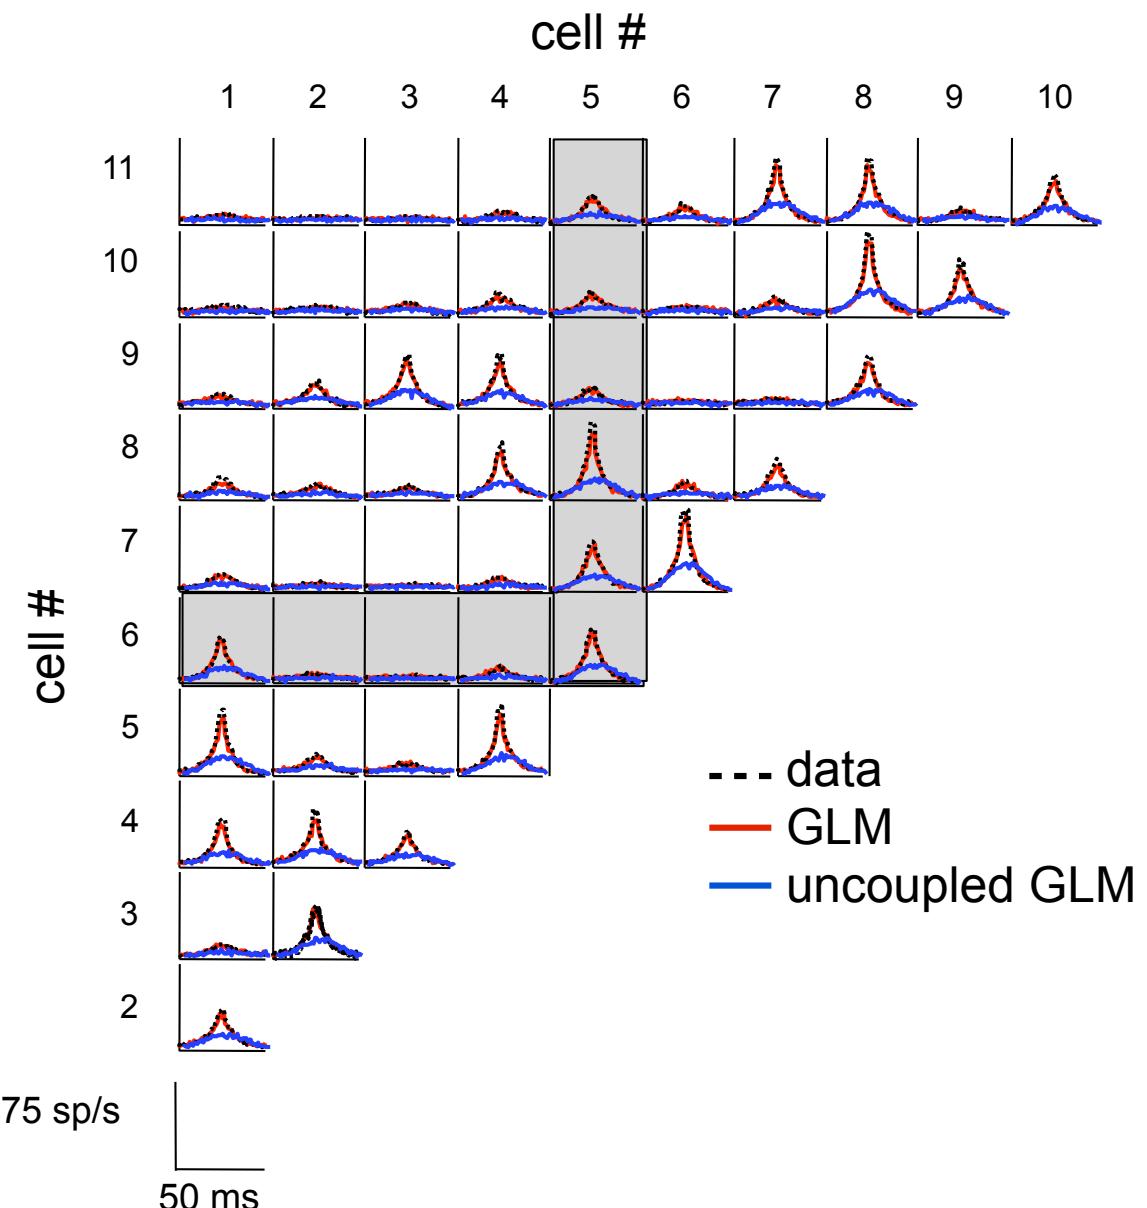
capturing dependencies in multi-neuron responses

[Pillow et al 2008]

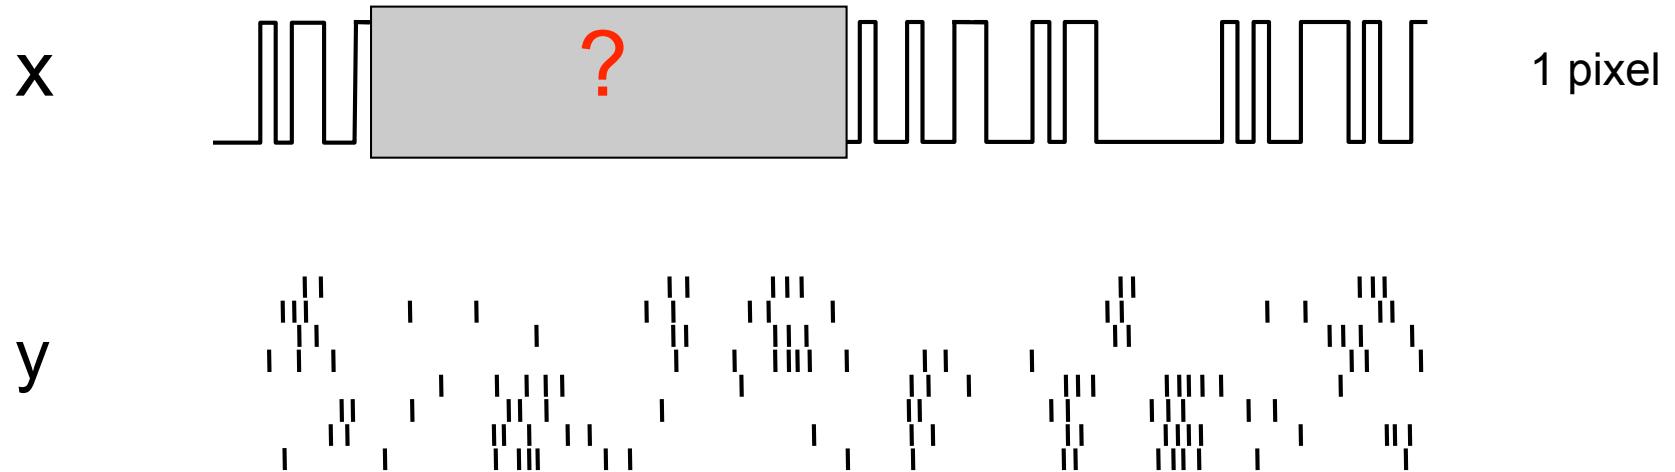
retinal receptive fields



cross-correlations

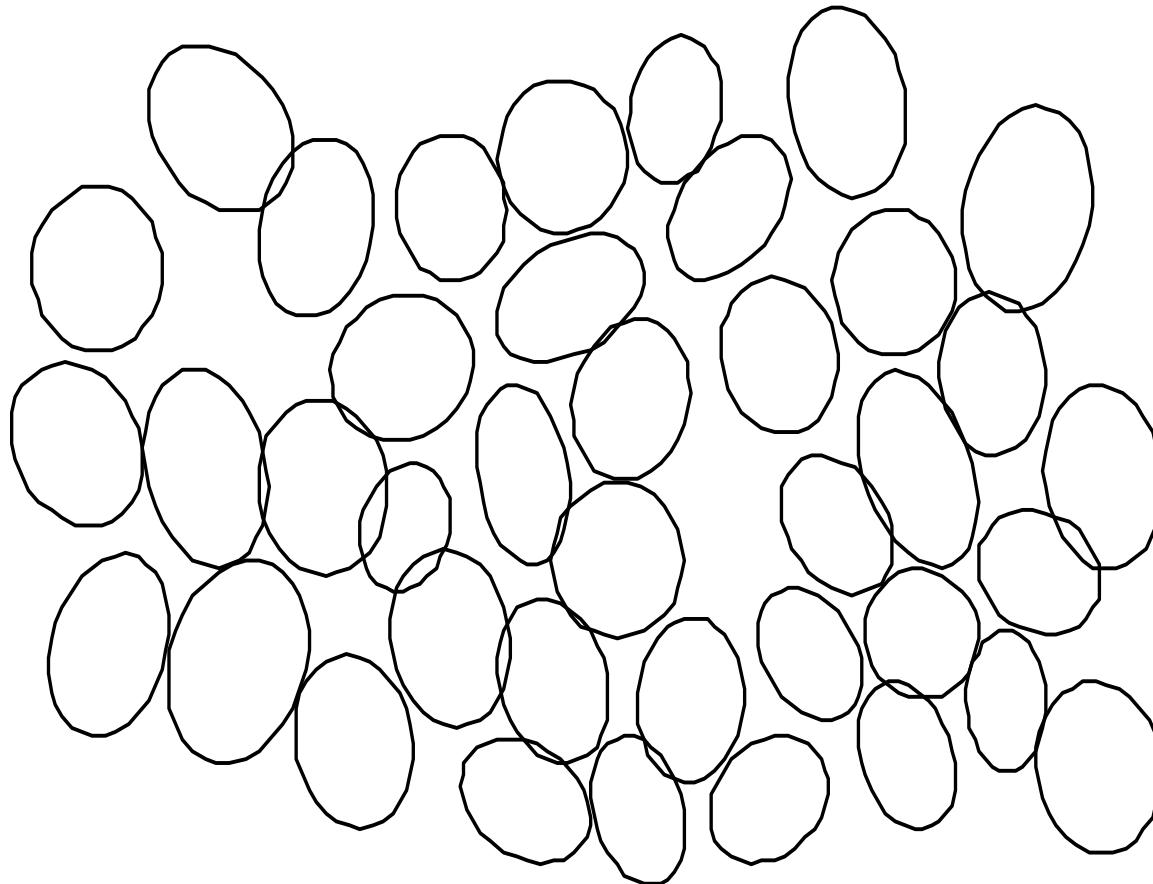


Decoding



- estimate stimuli from the observed spike times
- tool for comparing different encoding models

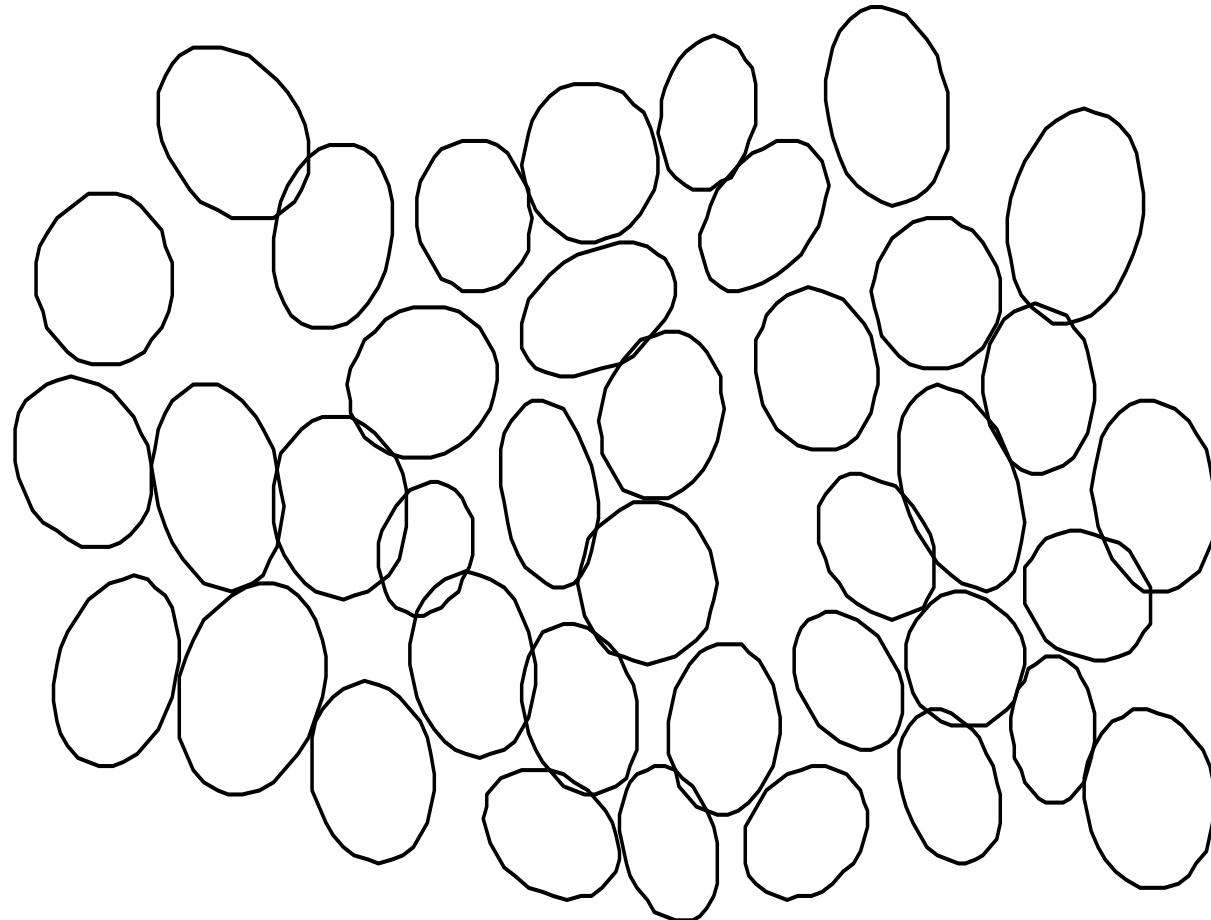
Decode: response 1



Q: what was the stimulus?

Frechette et al, 2005

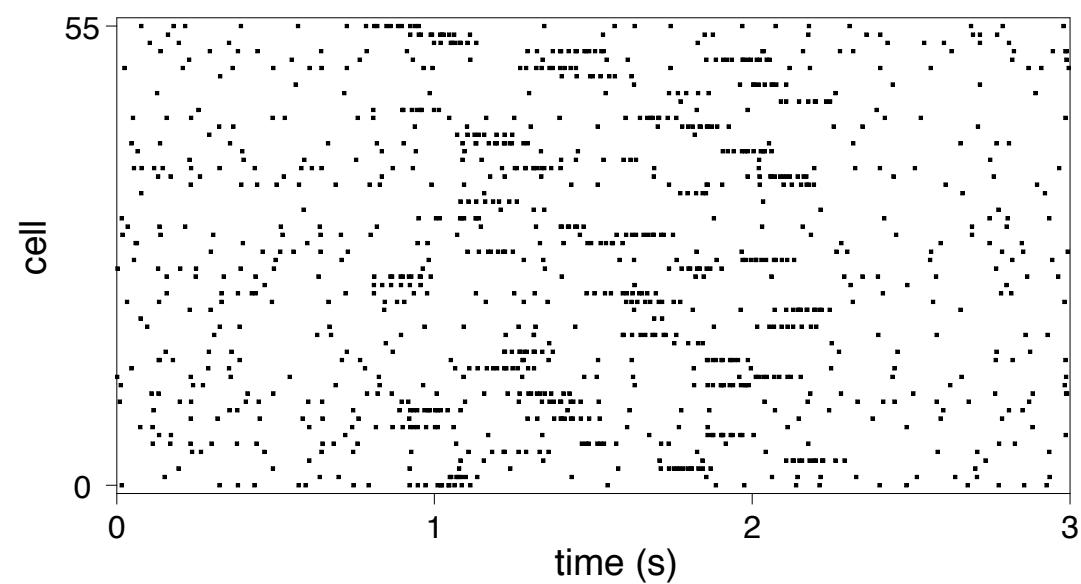
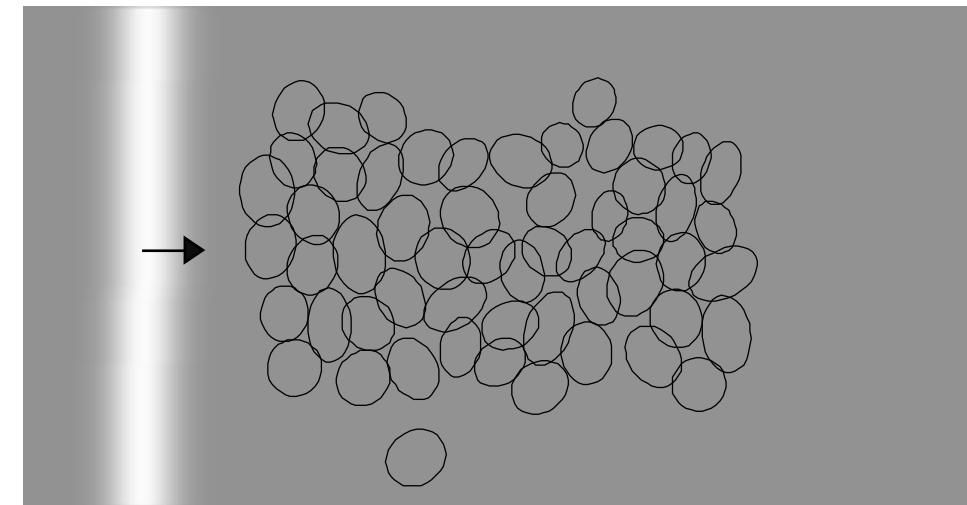
Decode: response 2



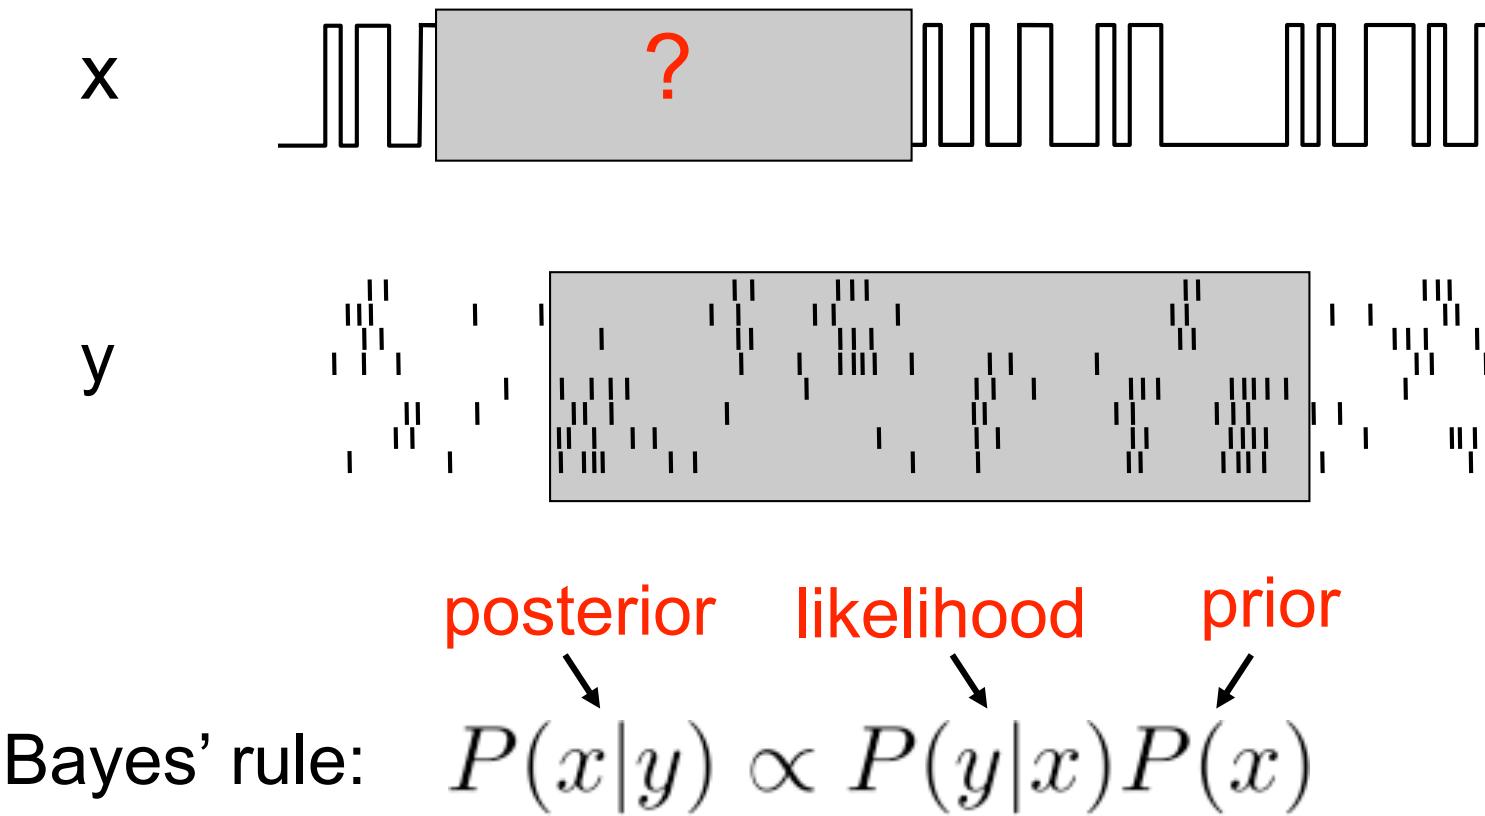
Q: what was the stimulus?

Frechette et al, 2005

Responses to Moving Bar



Bayesian Decoding



Bayesian Decoding



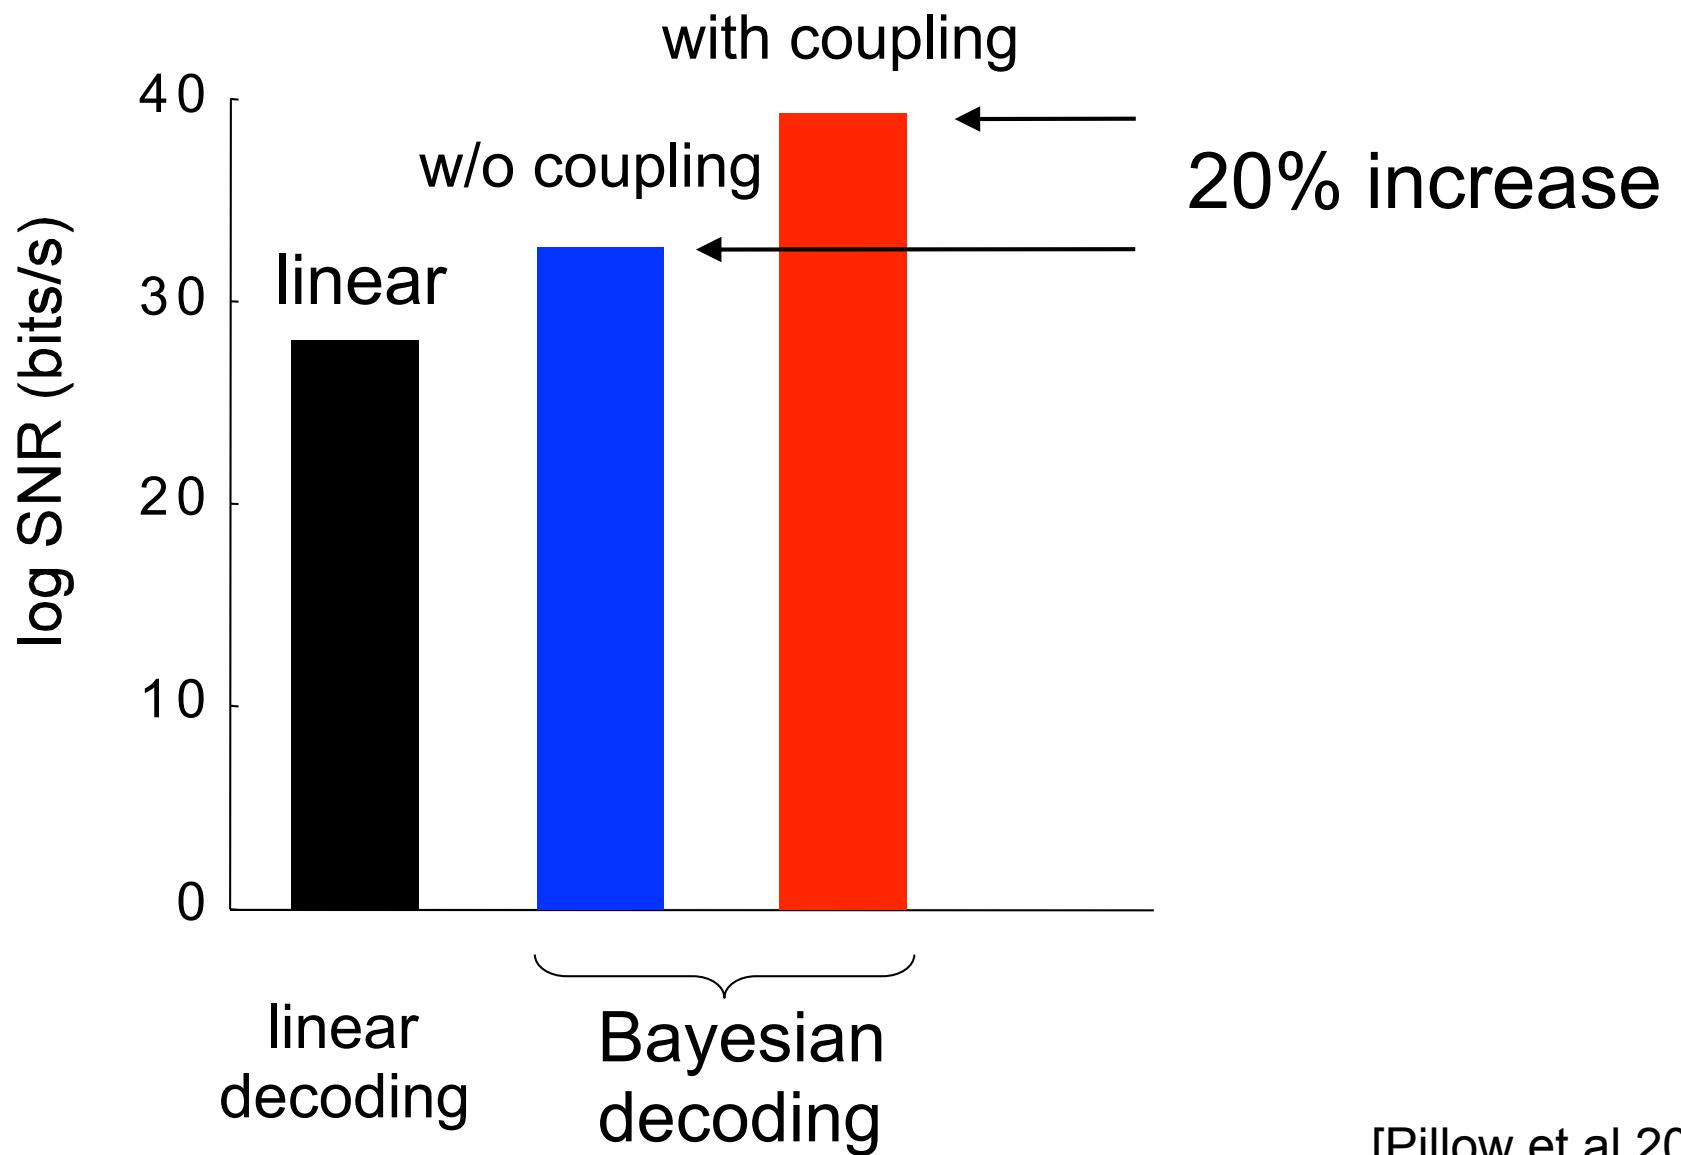
posterior likelihood prior

Bayes' rule: $P(x|y) \propto P(y|x)P(x)$

$P(y_1|x) \cdots P(y_n|x)$ vs. $P(y_1, y_2, \dots, y_n|x)$

“independent”
(uncoupled GLM) “joint encoding”
(coupled GLM)

Decoding Comparison



Regularization

Modern statistics

- more dimensions than samples $D \geq N$

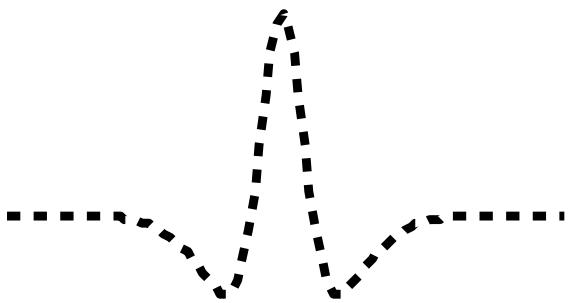
$\overbrace{\quad \quad \quad}^{\text{N observations}}$ {
$$\begin{bmatrix} y_1 \\ | \\ y_N \end{bmatrix}$$
 } =
$$\left[\begin{array}{c|c} \overbrace{\quad \quad \quad}^{\text{D regressors}} & \\ \hline \vec{x}_1 & \\ \hline \vec{x}_N & \end{array} \right] \begin{bmatrix} w_1 \\ | \\ w_D \end{bmatrix}$$
 + noise

- fewer equations than unknowns!
- no unique solution

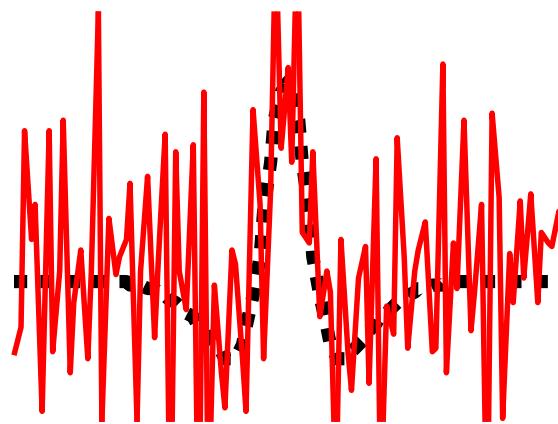
Simulated Example

- 100-element filter ($D=100$)
- 100 noisy samples ($N=100$)

true \mathbf{w}



maximum likelihood



maximize

$$\log p(\text{data}|\mathbf{w})$$

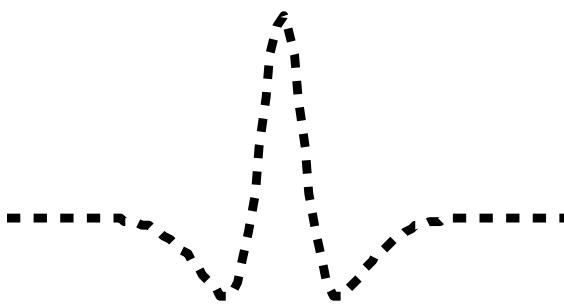
“overfitting” - parameters fit to details in the training data that are not useful for predicting new data

Simulated Example

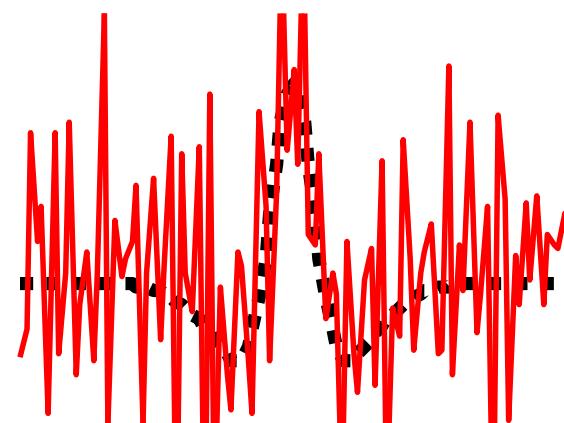
- 100-element filter ($D=100$)
- 100 noisy samples ($N=100$)

$$\hat{\mathbf{w}} = (X^\top X + \lambda I)^{-1} X^\top Y$$

true \mathbf{w}



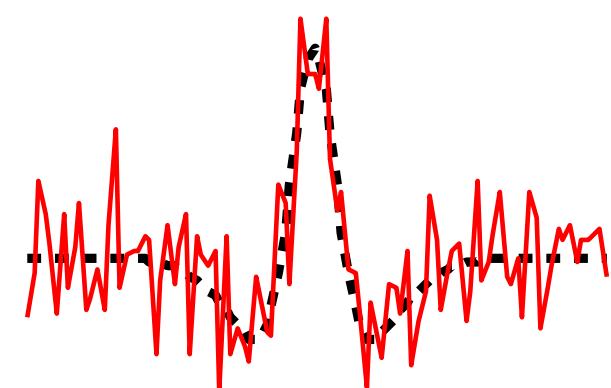
maximum likelihood



maximize

$$\log p(\text{data}|\mathbf{w})$$

“ridge regression”



maximize

$$\log p(\text{data}|\mathbf{w}) - \lambda \underbrace{\sum w_i^2}_{\text{penalty on big weights}}$$

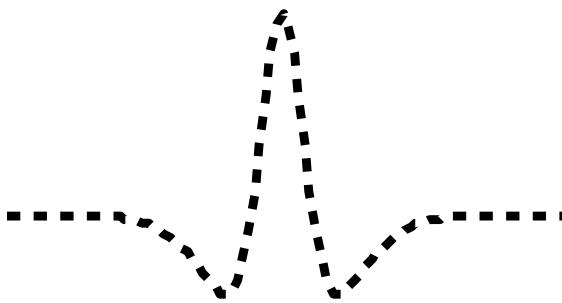
penalty on
big weights

- biased, but gives improved performance for appropriate choice of λ (James & Stein 1960)

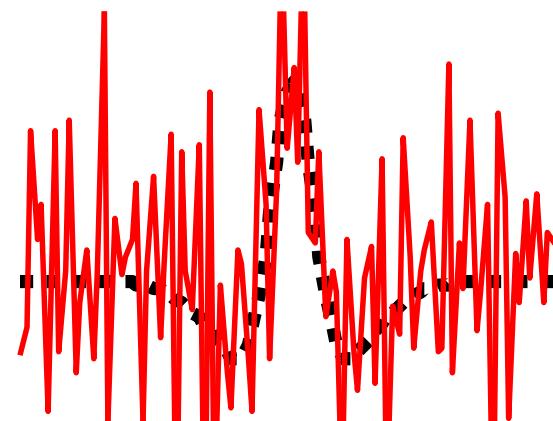
Simulated Example

- 100-element filter ($D=100$)
- 100 noisy samples ($N=100$)

true \mathbf{w}



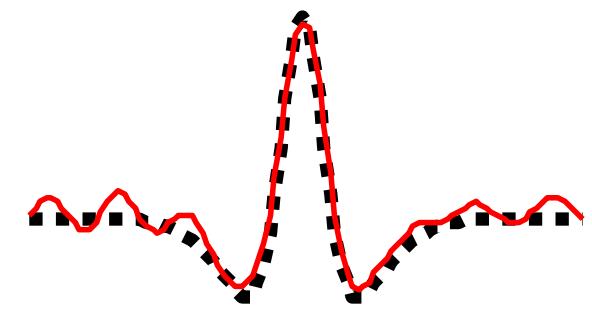
maximum likelihood



maximize

$$\log p(\text{data}|\mathbf{w})$$

“smoothed”



maximize

$$\log p(\text{data}|\mathbf{w})$$

$$-\lambda \sum (w_i - w_{i-1})^2$$

 smoothness

penalty

Q: how to set the regularization strength λ ?

Simplest answer: use cross-validation!

GLM tutorial (matlab):

code: <https://github.com/pillowlab/GLMspiketraintutorial>

data: available on request from pillow@princeton.edu

- **tutorial1_PoissonGLM.m** - fitting of a linear-Gaussian GLM and Poisson GLM (aka LNP model) to RGC neurons stimulated with temporal white noise stimulus.
- **tutorial2_spikehistcoupledGLM.m** - fitting of a Poisson GLM with spike-history and coupling between neurons.
- **tutorial3_regularization_linGauss.m** - regularizing linear-Gaussian model parameters using maximum a posteriori (MAP) estimation under two kinds of priors:
 - (1) ridge regression (aka "L2 penalty");
 - (2) L2 smoothing prior (aka "graph Laplacian").
- **tutorial4_regularization_PoissonGLM.m** - MAP estimation of Poisson-GLM parameters using same two priors as in tutorial3.

GLM summary

- linear (“dim reduction”) + nonlinear + noise
- incorporate spike-history via “spike history” filter
- rich dynamical properties: refractoriness, bursting, adaptation
- incorporate correlations between neurons via “coupling” filters
- flexible tool for encoding & decoding analyses
- regularize to reduce overfitting (essential w/ correlated stimuli)

Beyond GLM

polynomial models

Lee & Schetzen 1965
Marmarelis & Naka 1972
Korenberg & Hunter 1986

Volterra / Wiener Kernels

Taylor series expansion of a function $f(\vec{x})$ in n dimensions

$$y = k_0 + \vec{k}_1 \cdot \vec{x} + \vec{x}^t K_2 \vec{x} + K_3 \cdot \vec{x}^3 + \dots$$



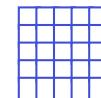
const



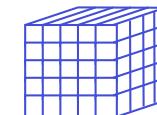
vector



matrix



3-tensor



parameters:



n
(20)

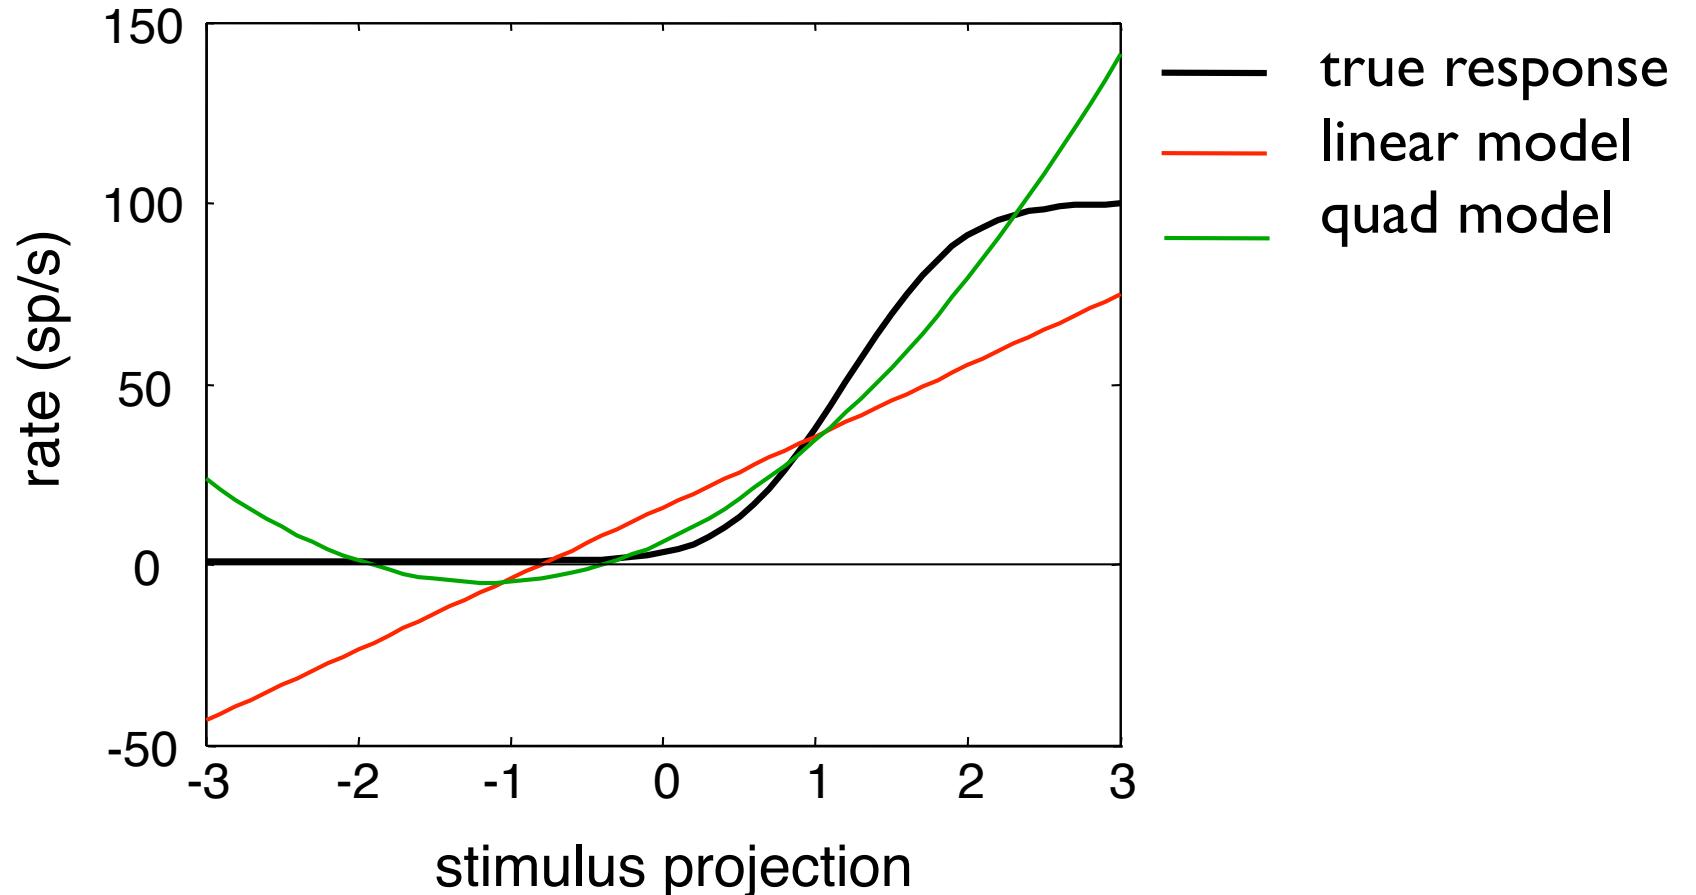
n^2
(400)

n^3
(8000)

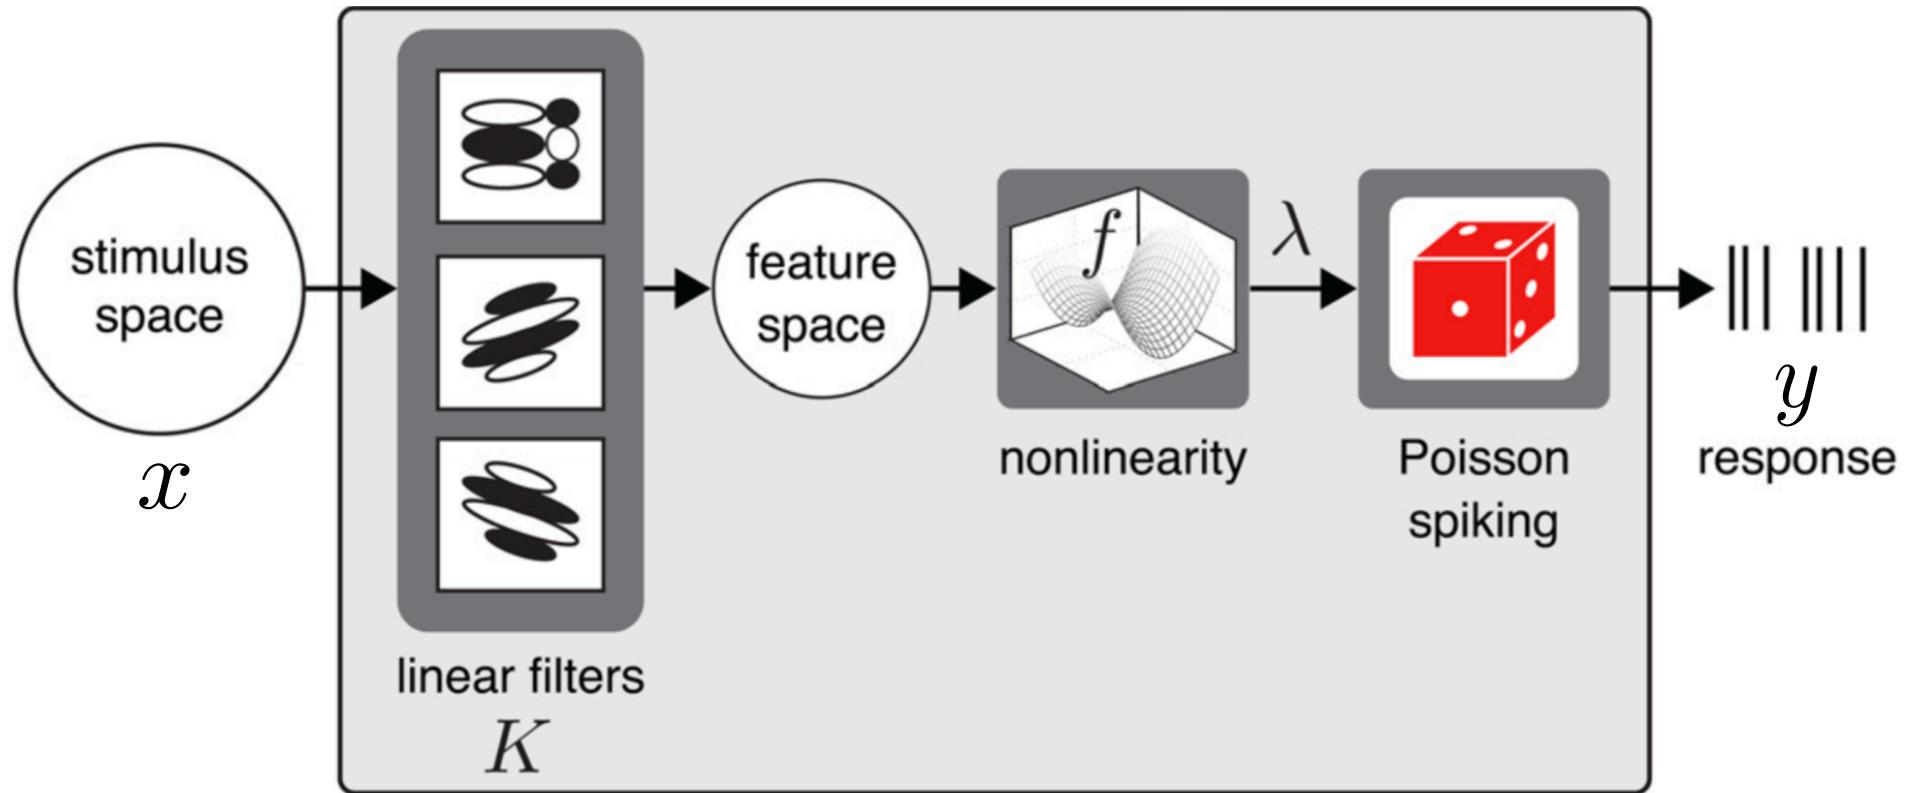
- from “systems identification” literature (1960s-70s)
- white noise stimuli
- estimate kernels using moments of spike-triggered stimuli

Why are Volterra/Wiener models (generally) bad?

- no output nonlinearity
- polynomials give poor fit to neural nonlinearities (e.g., rectifying, saturating)

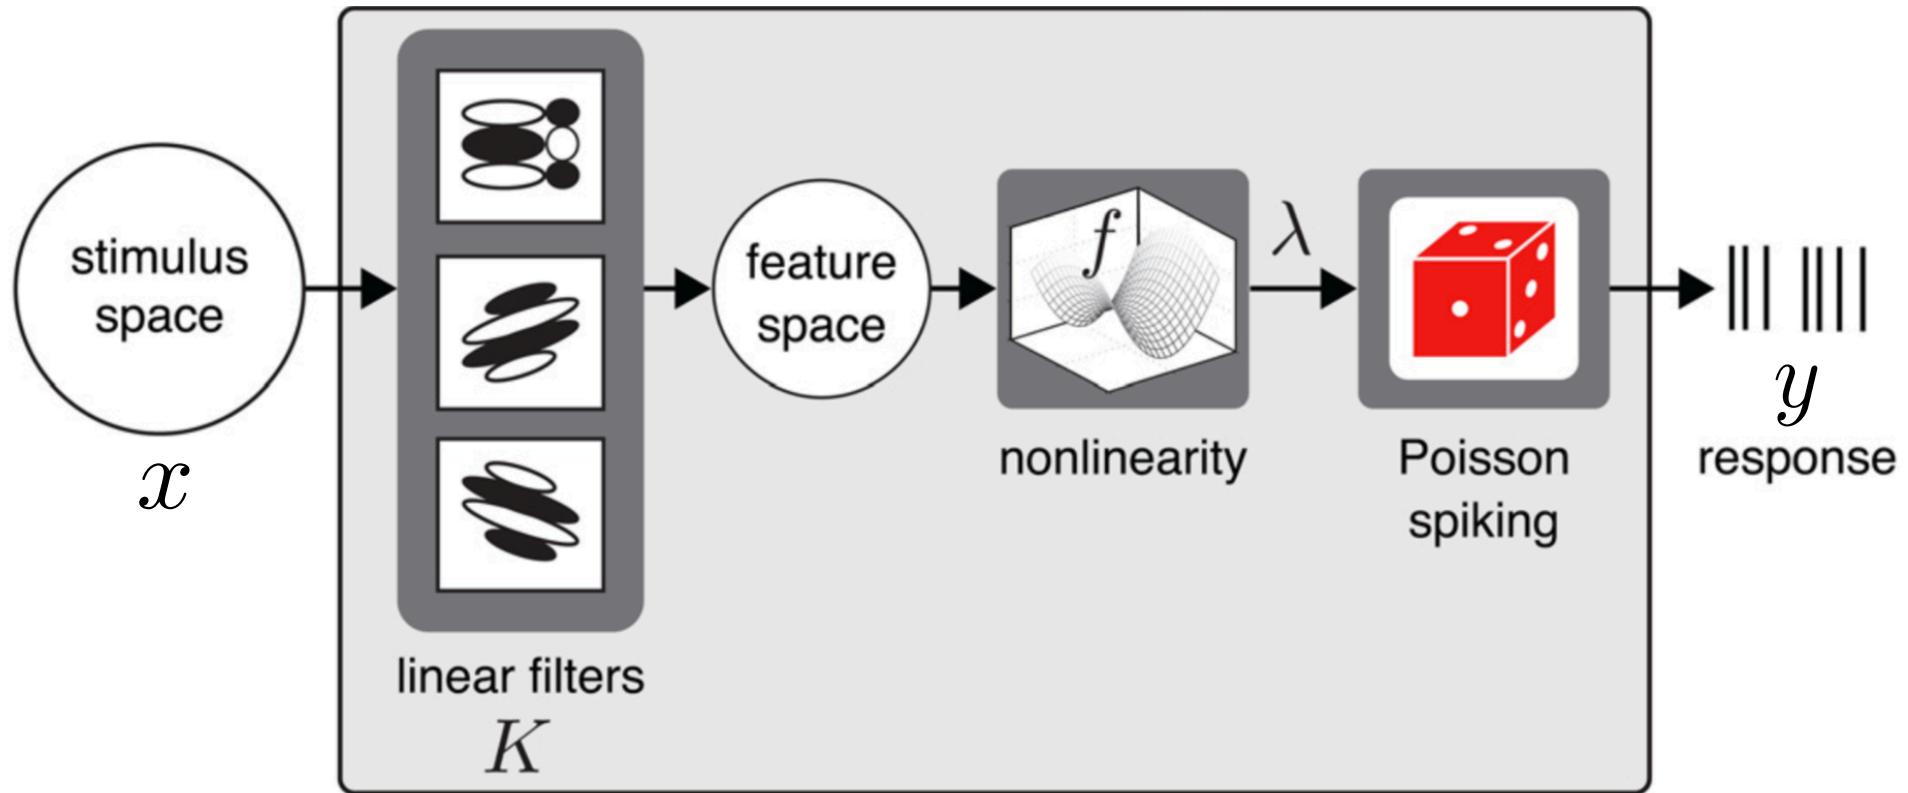


multi-filter LNP



- responses may depend on more than one projection of stimulus!
- emphasis on *dimensionality reduction*
- no longer technically a GLM if fitting nonlinearity f

multi-filter LNP



Estimators:

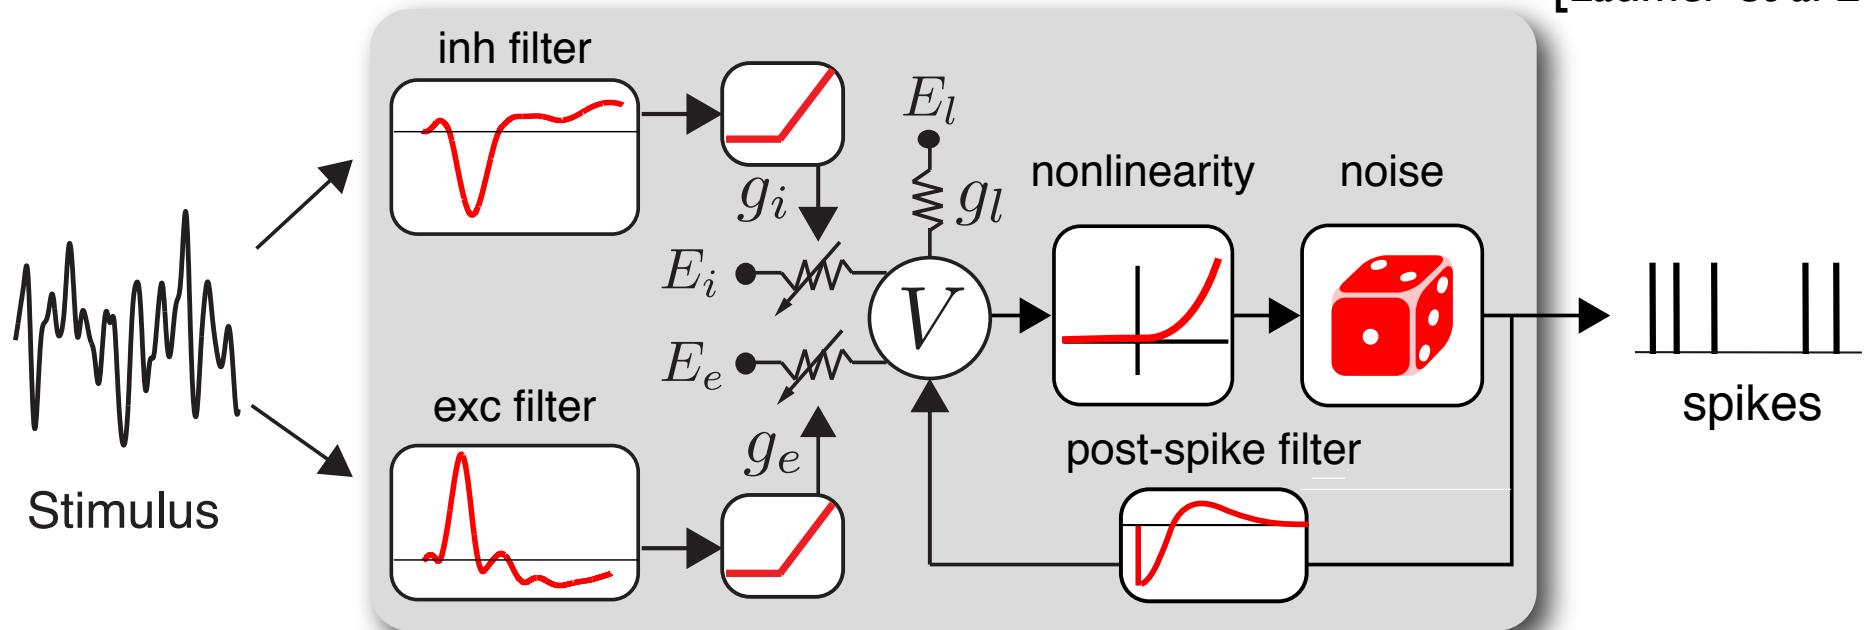
- Spike-triggered covariance (STC) [de Ruyter & Bialek 1998, Schwartz et al 2006]
- Generalized Quadratic Model (GQM) [Park & Pillow 2011; Park et al 2013; Rajan et al 2013]
- maximally informative dimensions (MID) / maximum likelihood

[Sharpee et al 2004]

[Williamson et al 2015]

extending GLM to conductance-based model

[Latimer et al 2014]



conductances

$$g_e(t) = f_c(k_e \cdot \mathbf{x}(t))$$

$$g_i(t) = f_c(k_i \cdot \mathbf{x}(t))$$

membrane
dynamics

$$\frac{dV}{dt} = g_l(E_l - V) + g_e(E_e - V) + g_i(E_i - V)$$

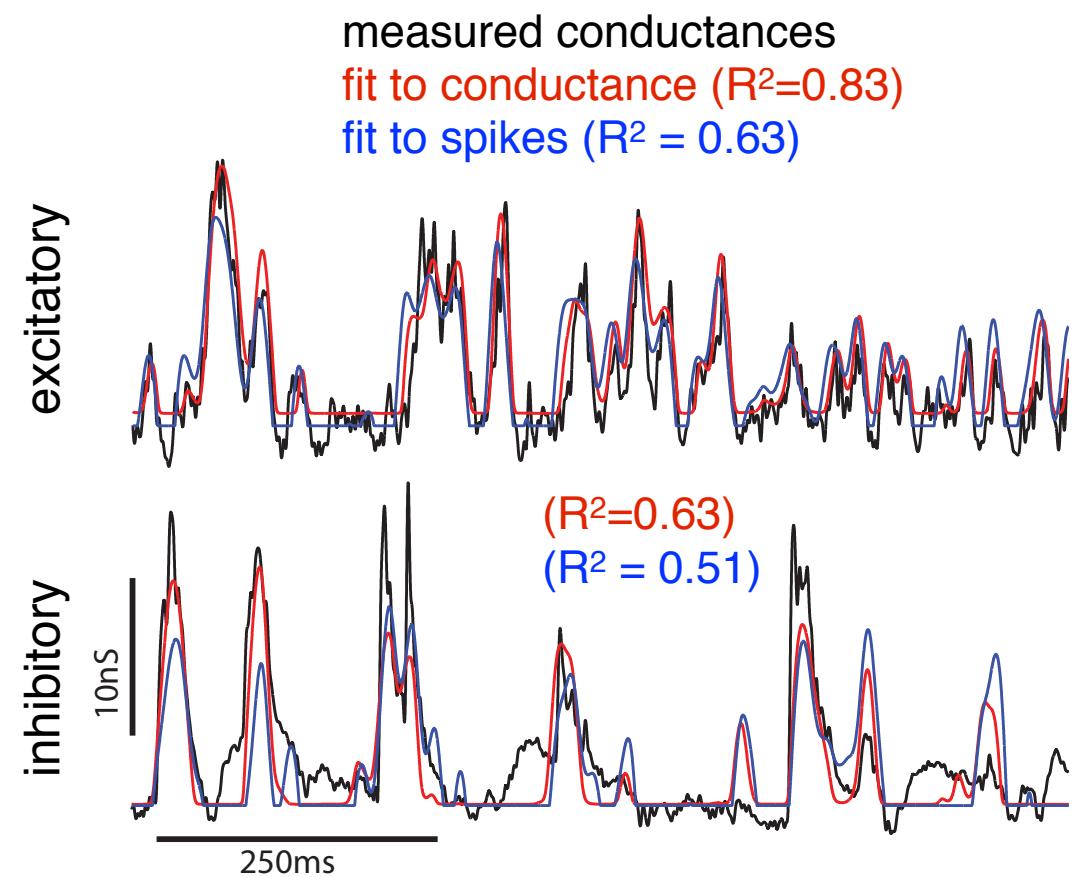
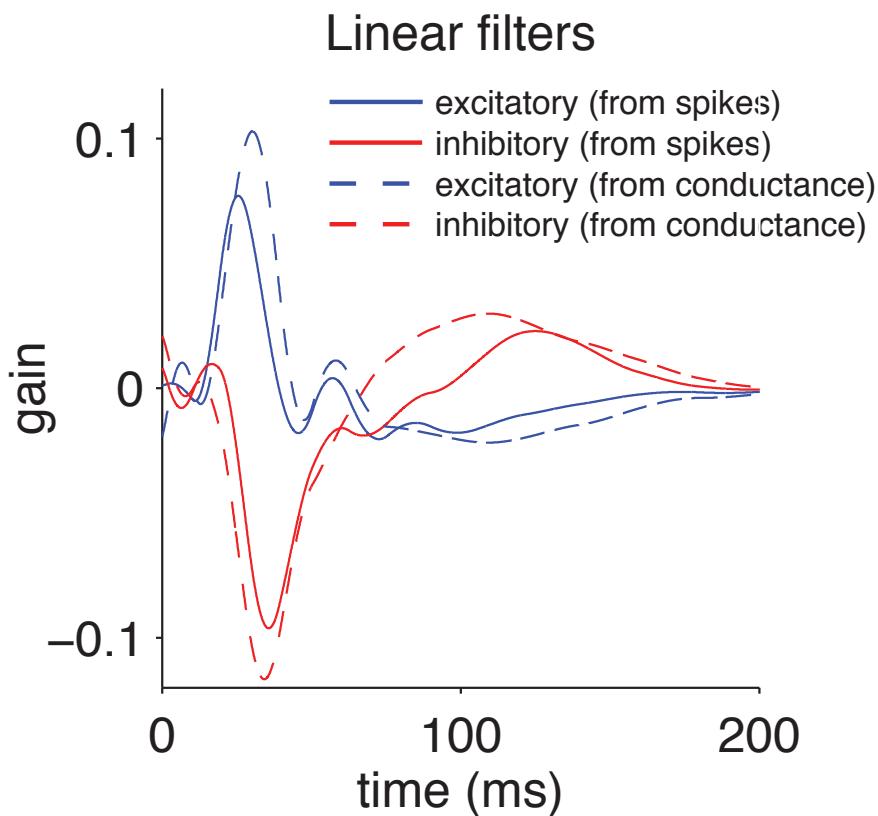
inst. spike rate

$$\lambda(t) = f(V(t))$$

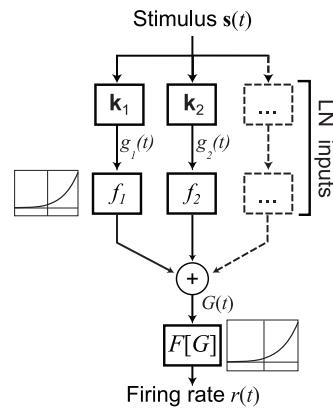
- shunting inhibition
- adaptive changes in dynamics

extending GLM to conductance-based model

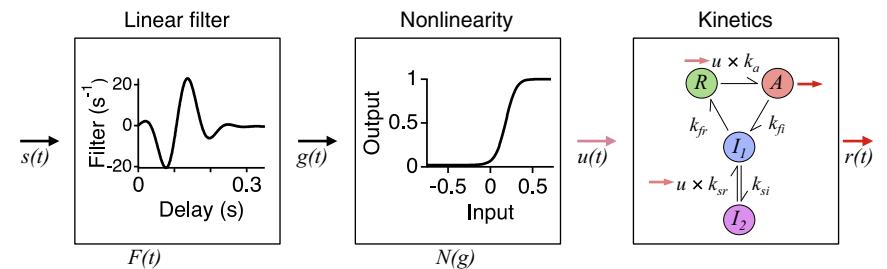
- intracellular recordings in macaque parasol RGCs (Fred Rieke)



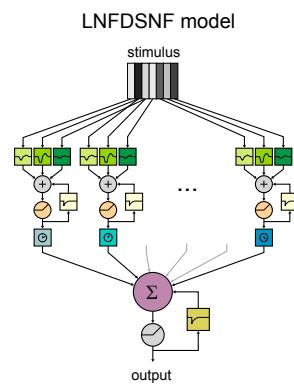
many other biophysically oriented extensions



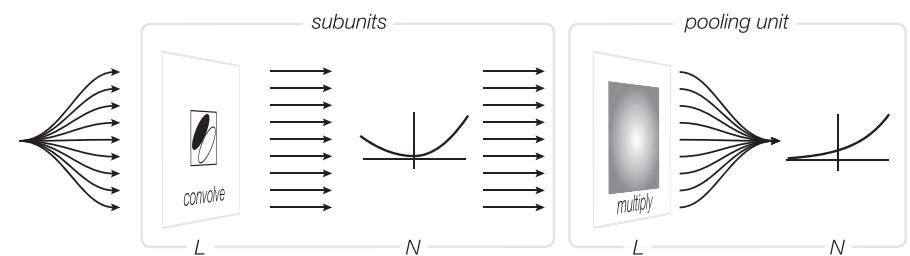
Nonlinear input model (NIM)
[McFarland, Cui, & Butts 2013]



Linear-Nonlinear-Kinetics (LNK)
[Ozuyal & Baccus 2014]

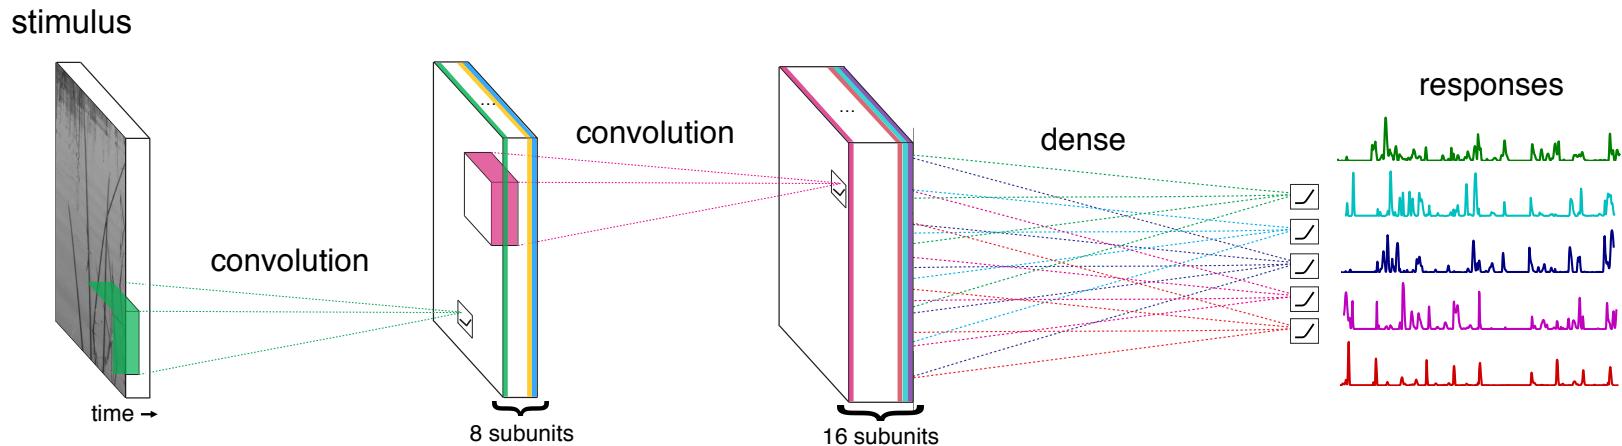


LNFDSNF model
linear-nonlinear-feedback-delayed-sum-nonlinear-feedback
[Real, Asari, Gollisch & Meister 2017]



convolutional subunit model
[Vintch, Movshon & Simoncelli 2015,
Wu, Park & Pillow 2014]

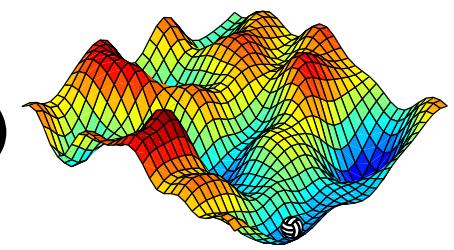
deep learning / deep neural networks (DNNs)



[Yamins et al 2014, McIntosh et al 2016, Maheswaranathan et al 2017, Benjamin et al 2017, ...]

If you understand GLMs... you understand DNNs!

- stack many LNs on top of each other: LN LN LN LN P
- use gradient ascent to maximize likelihood
- use software (tensorflow, theano) to compute gradients
(no more computing gradients by hand!)
- use a bunch of tricks (batches, noise, SGD, dropout,)
- do NOT worry about local maxima!



[credit: Jakob Macke]

Modern machine learning far outperforms GLMs at predicting spikes

Ari S. Benjamin¹, Hugo L. Fernandes², Tucker Tomlinson³, Pavan Ramkumar^{2,4}, Chris VerSteeg¹, Lee Miller^{1,2,3}, Konrad Paul Kording^{1,2,3}

macaque M1

Fig 2

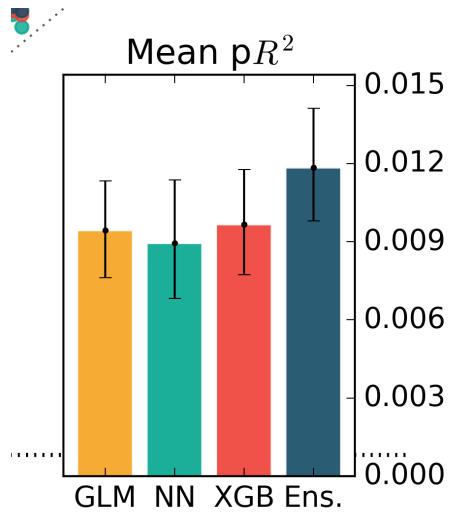


Fig 3

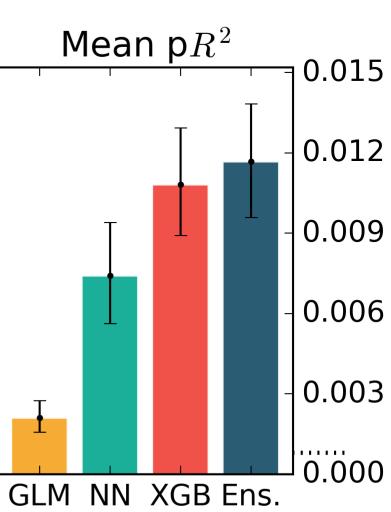


Fig 4

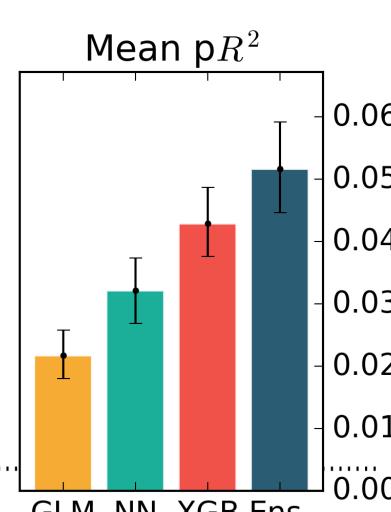


Fig 5

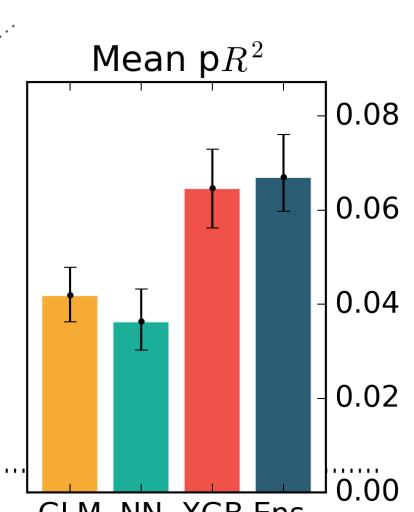
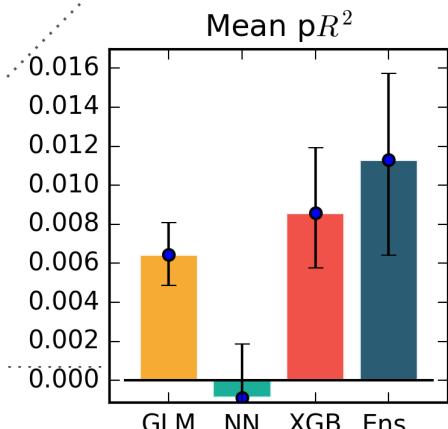
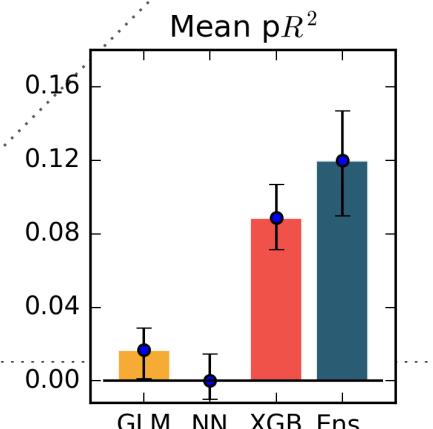


Fig 6



macaque S1

Mean pR^2



hippocampus

GLMs

NNs

~~Modern machine learning far outperforms GLMs at predicting spikes ?~~

Ari S. Benjamin¹, Hugo L. Fernandes², Tucker Tomlinson³, Pavan Ramkumar^{2,4}, Chris VerSteeg¹, Lee Miller^{1,2,3}, Konrad Paul Kording^{1,2,3}

Fig 2

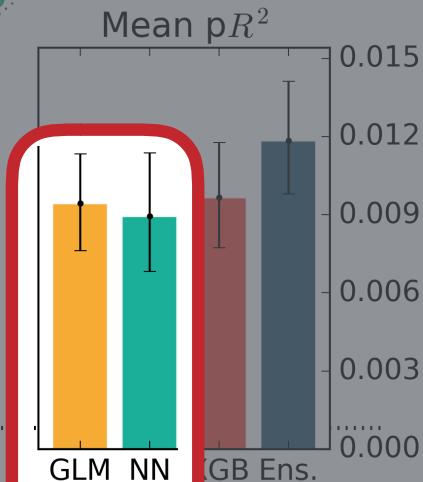


Fig 3

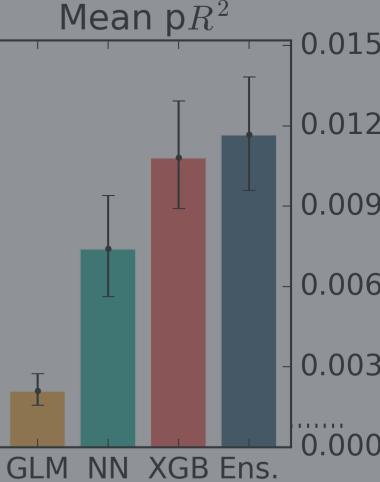


Fig 4

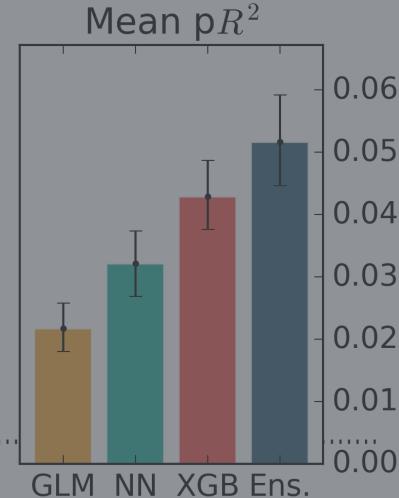


Fig 5

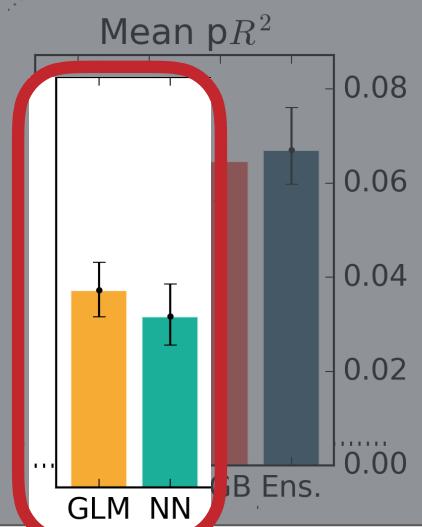
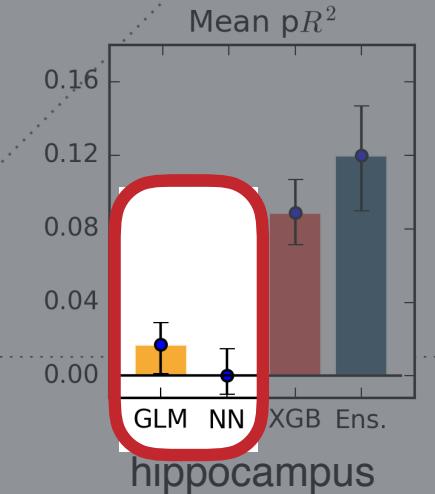
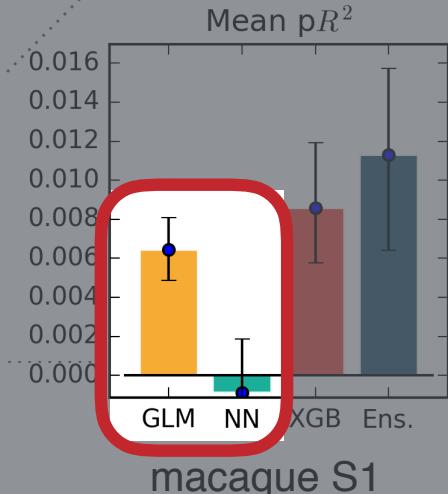


Fig 6



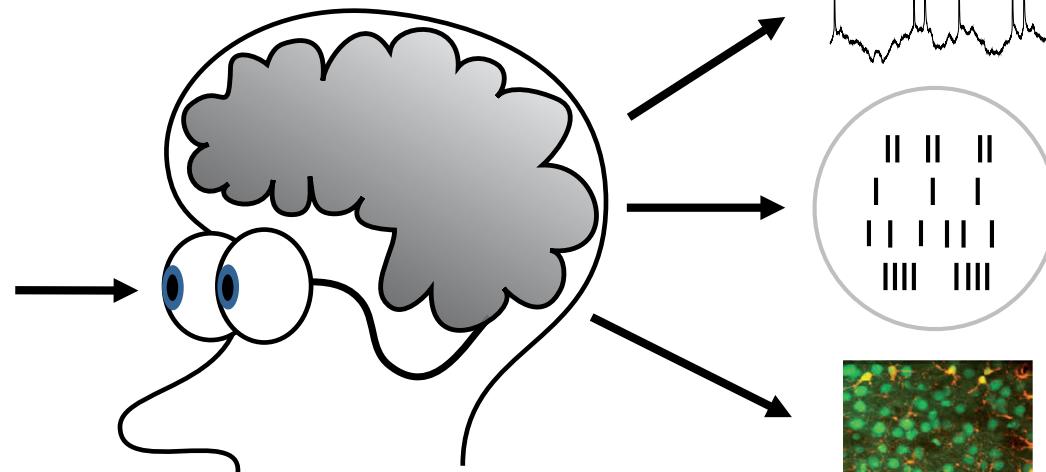
(No of course not!)

GLM is a special case of
NN!

encoding models



stimulus
 x

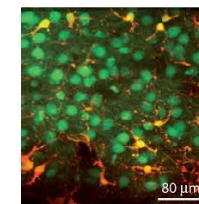


$P(y|x)$
encoding models

membrane
potential

spikes

imaging

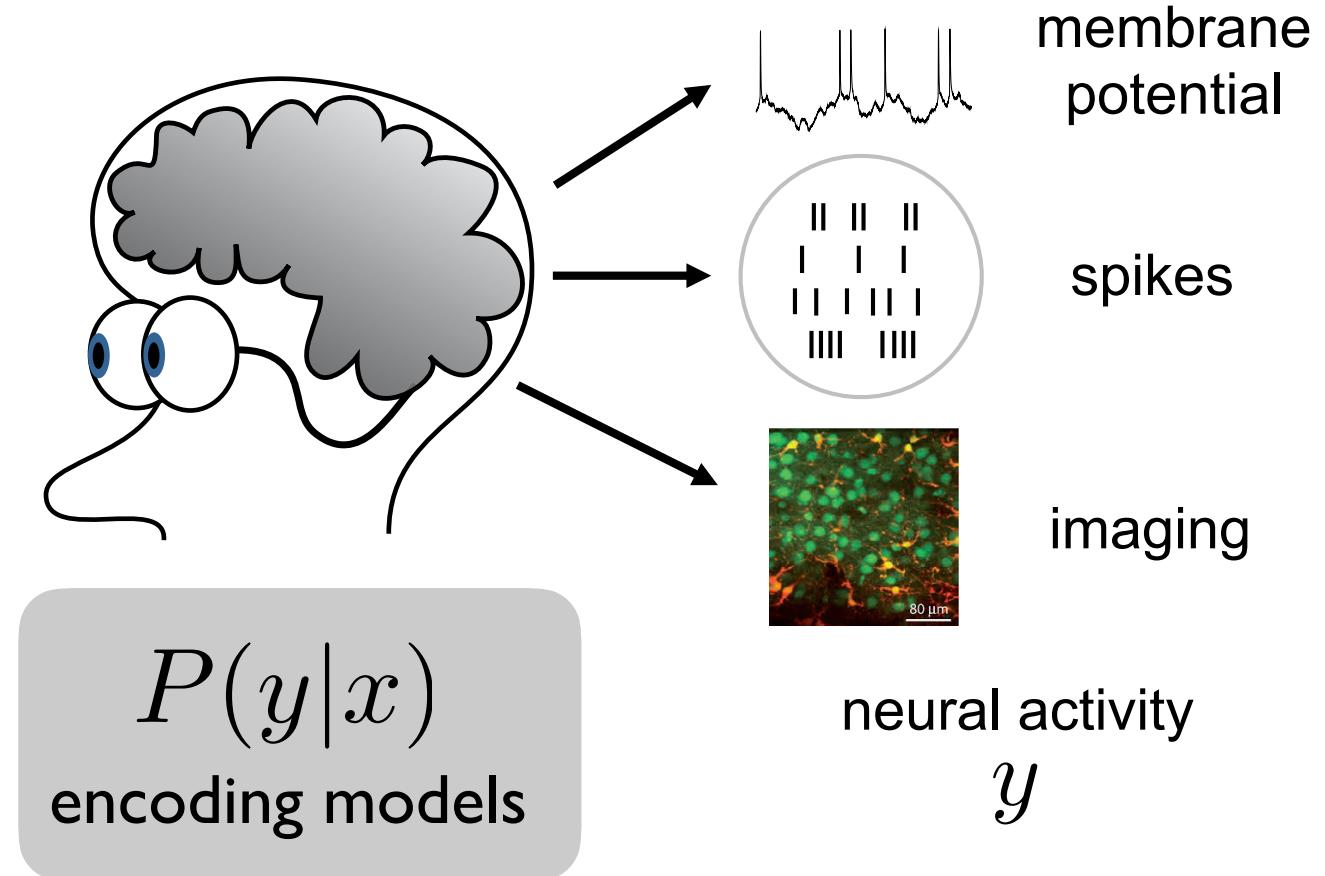


neural activity
 y

What if there's
no stimulus?

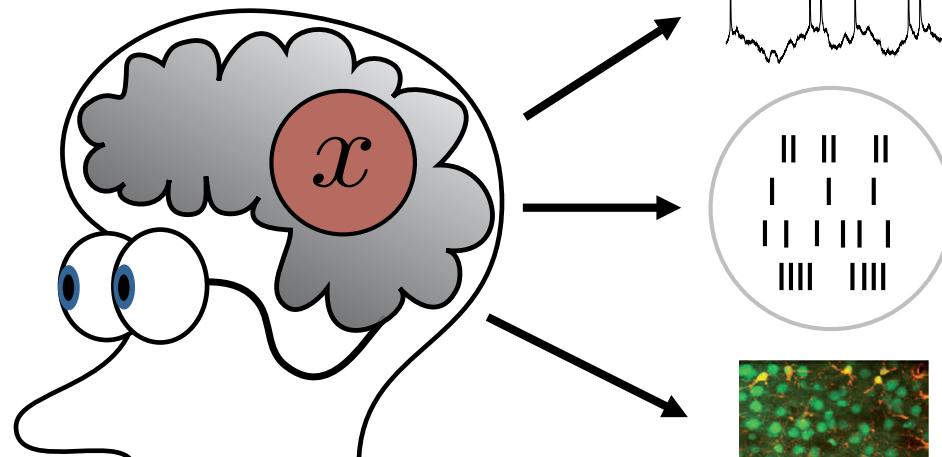
encoding models

**What if there's
no stimulus?**

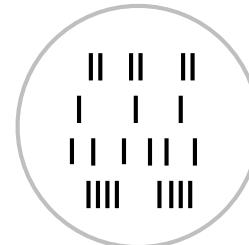


latent variable models

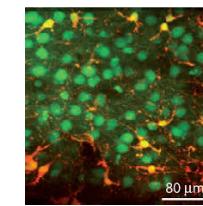
latent variable
(unobserved or
“hidden”)



membrane potential



spikes



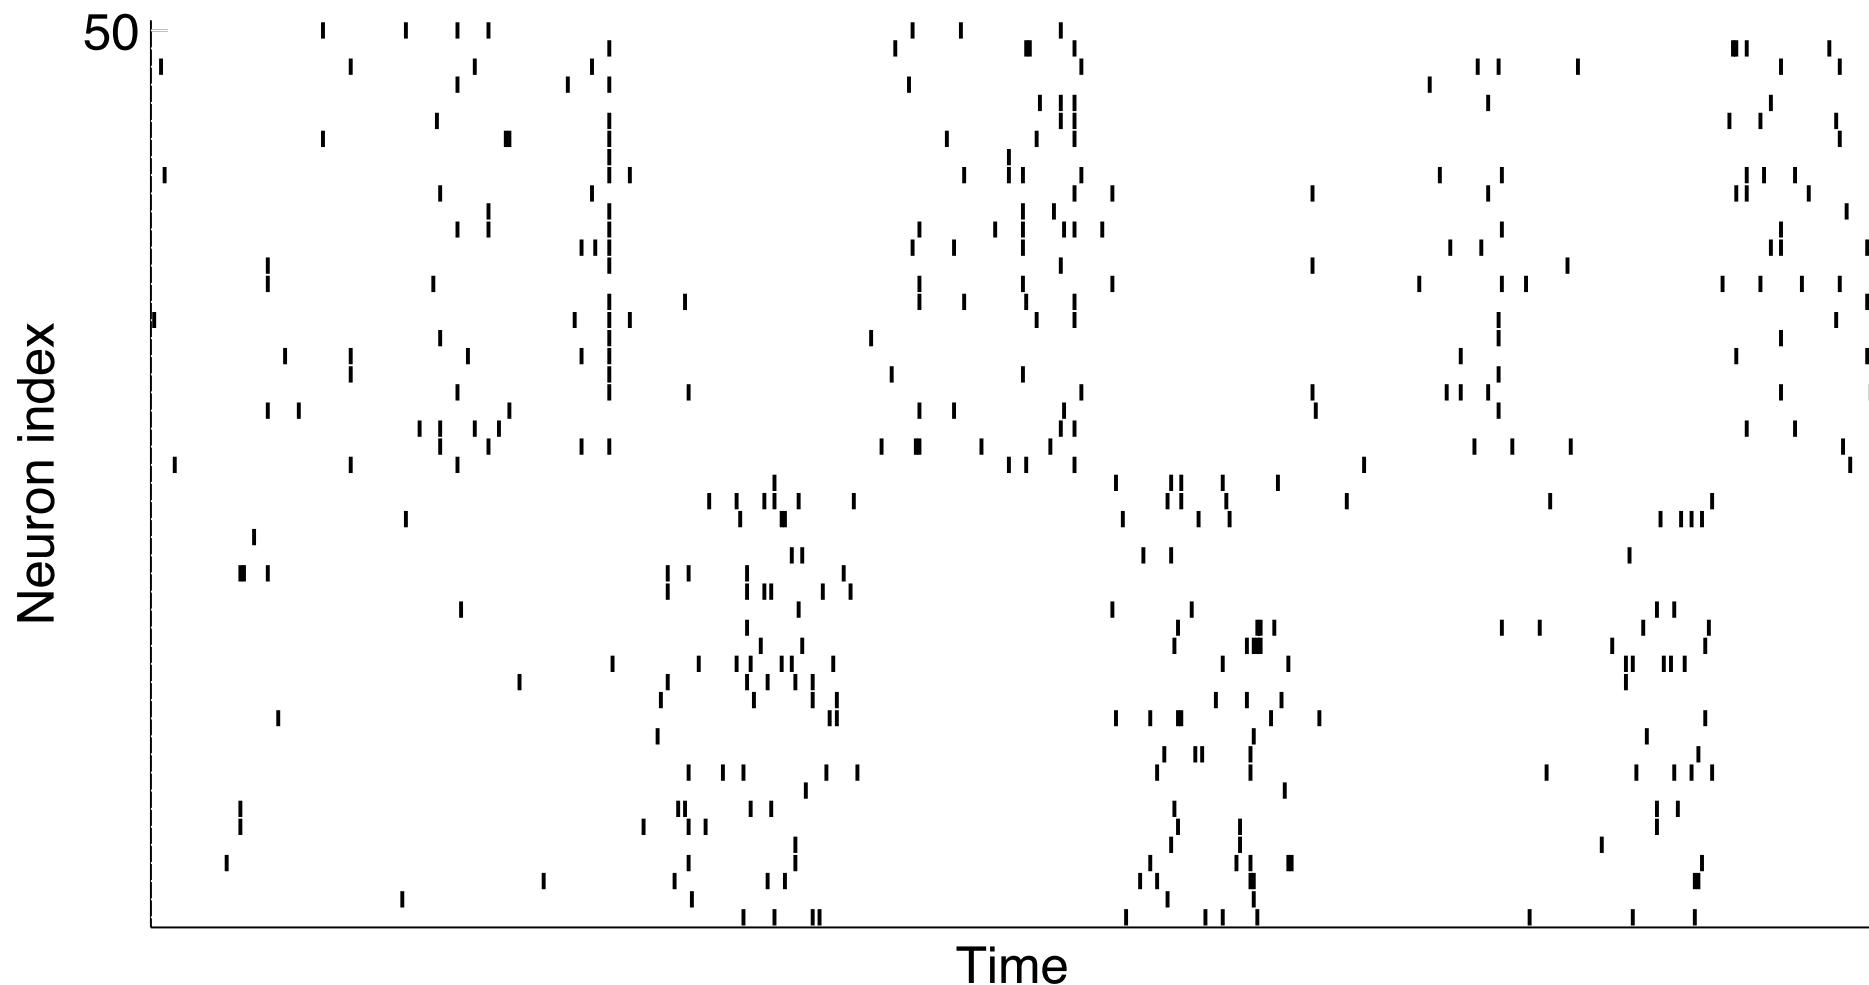
imaging

$$P(y|x) P(x)$$

latent encoding
models

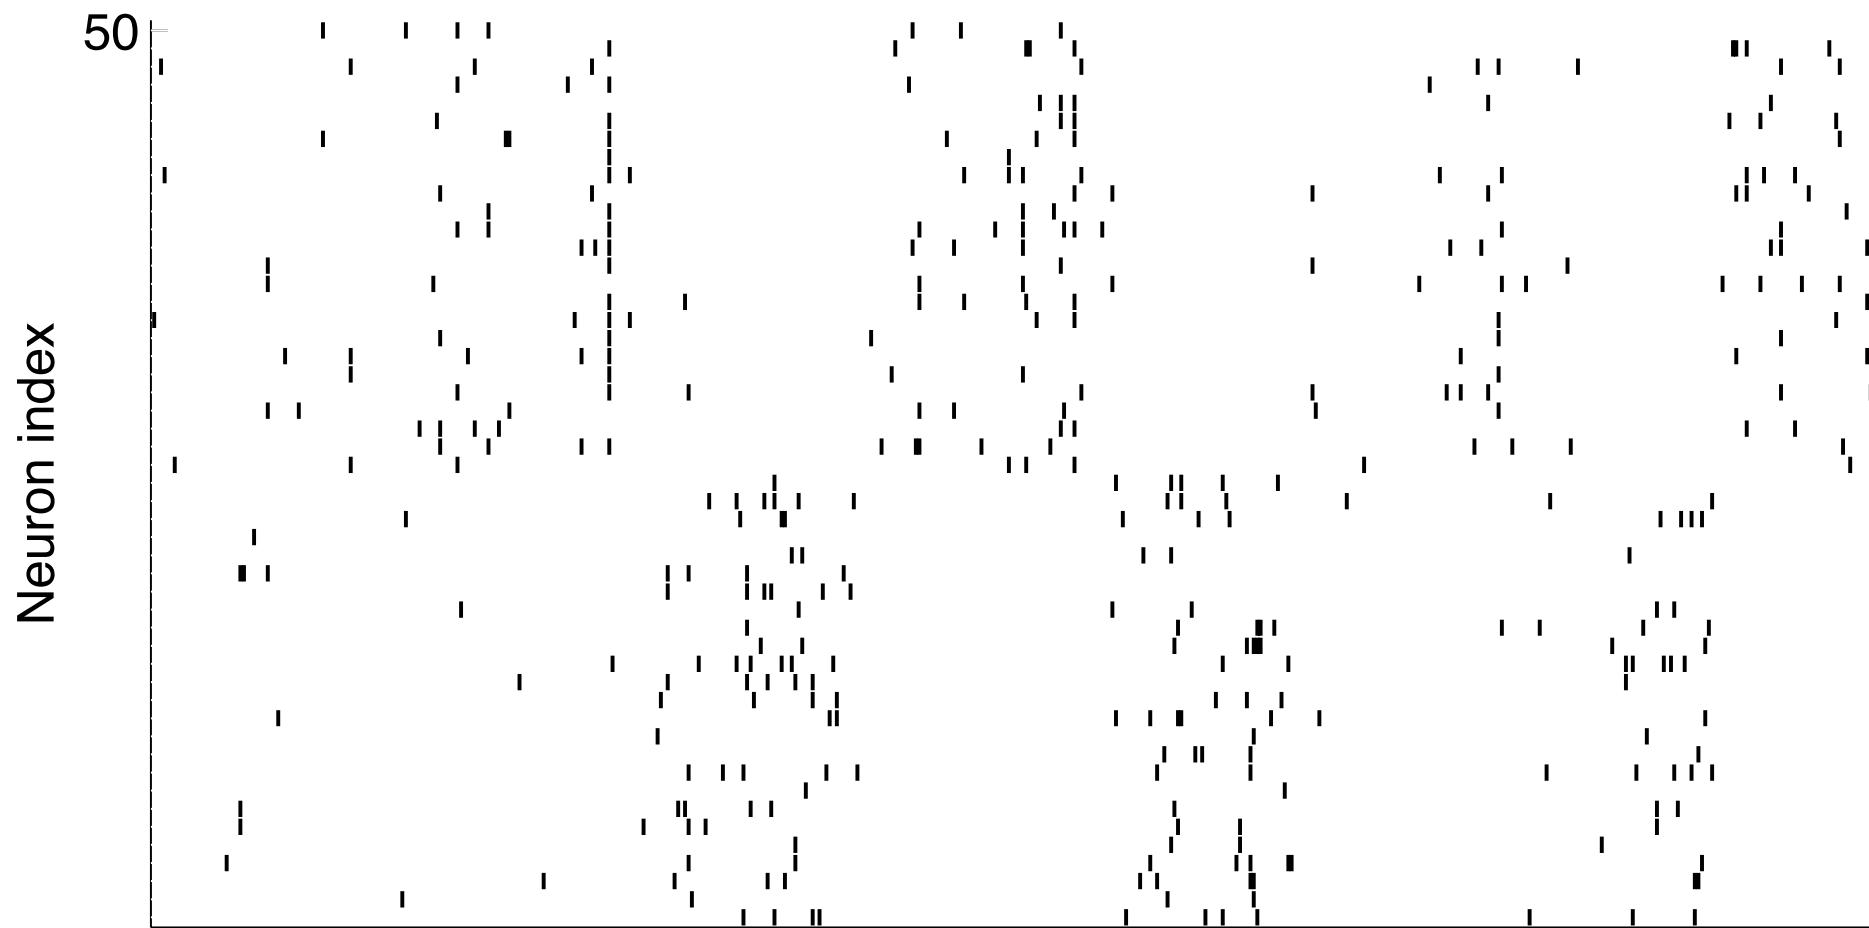
neural activity
 y

spike responses

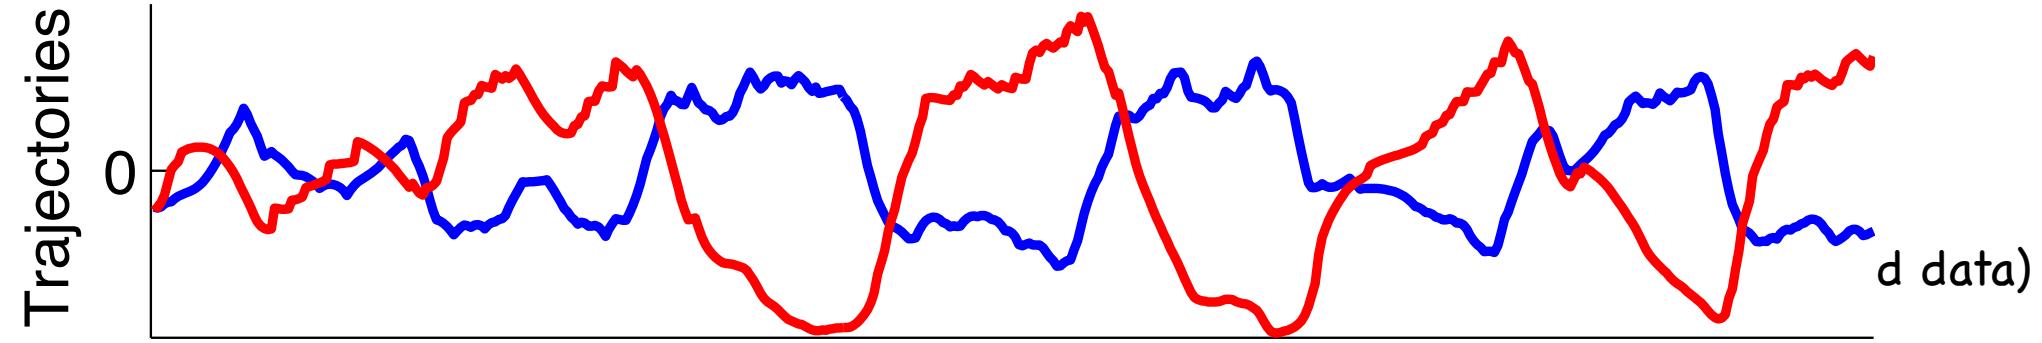


(simulated data)

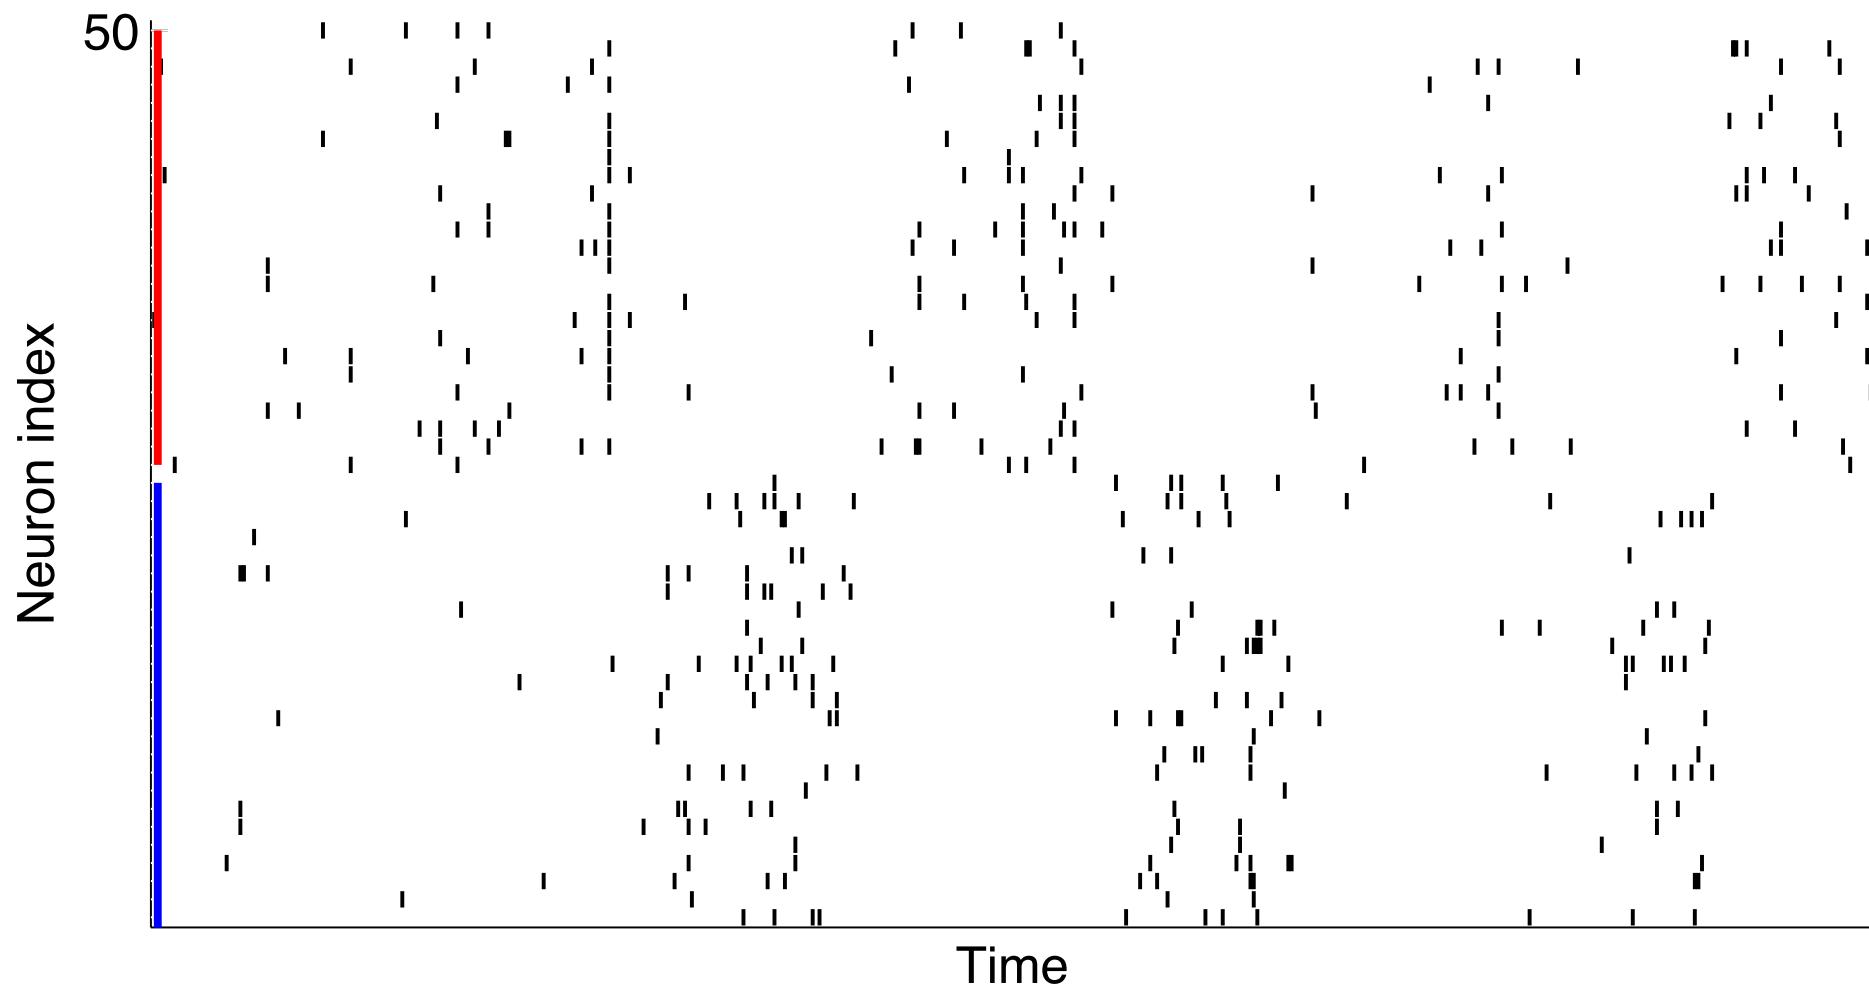
spike responses



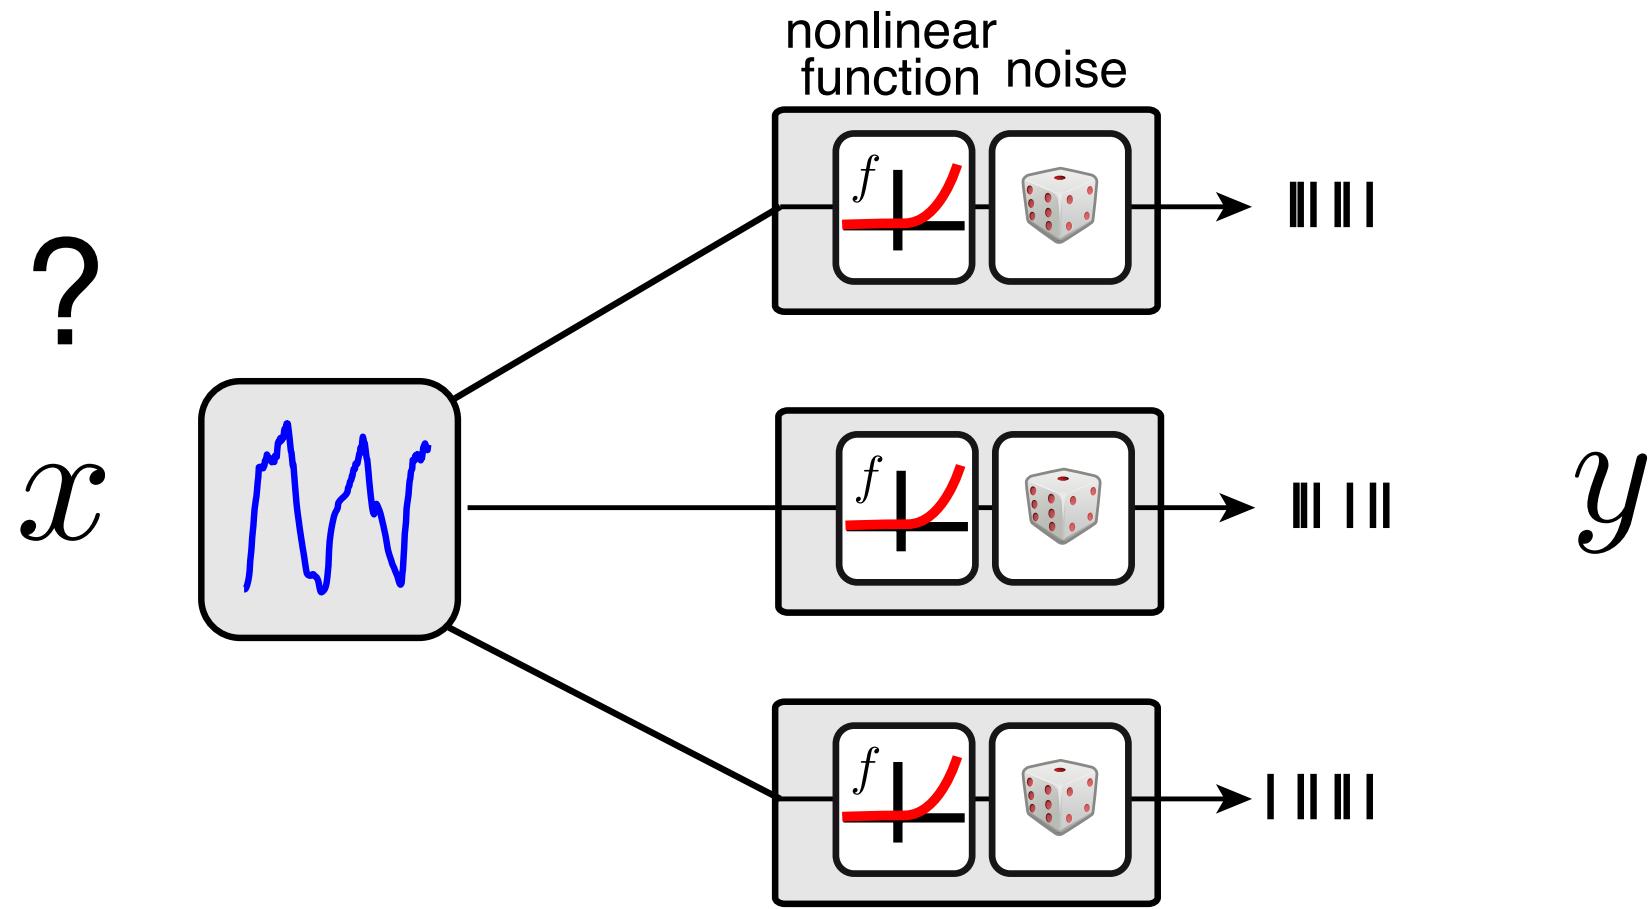
inferred latent variables



spike responses



latent variable models = GLMs where we don't know x



goal: find shared structure underlying y

chicken and egg problem

Why are latent variable models hard to work with?

- hard to compute likelihood!

Why are latent variable models hard to work with?

- hard to compute likelihood!

**sensory
encoding model**

x



y

Poisson GLM: fit using:

$$\log P(Y|X) = \sum_t y_t \log \lambda_t - \lambda_t$$

Why are latent variable models hard to work with?

- hard to compute likelihood!

**sensory
encoding model**

x



y

**latent
variable model**

x



y

fit using:

$$\log P(Y|X) = \sum_t y_t \log \lambda_t - \lambda_t$$

fit using: $\log P(Y)$

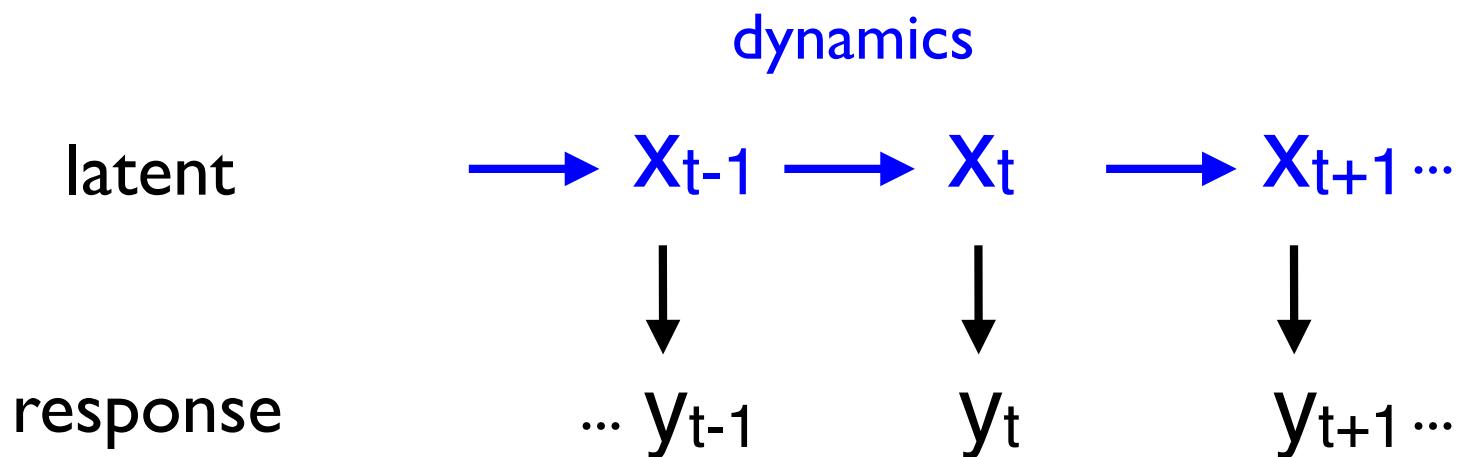
$$= \log \int P(Y|X)P(X)dx$$

requires an integral!

Why are latent variable models hard to work with?

- hard to compute likelihood!

latent dynamical model



fit using: $\log P(Y) = \int \prod_{t=1}^T \left(\underbrace{p(y_t|x_t)}_{\text{observations}} \underbrace{p(x_t|x_{t-1})}_{\text{dynamics}} \right) dx_1 dx_2 \cdots dx_T$

high-dimensional integral

Fitting Latent Variable Models

I. Sampling (“MCMC”) - fully Bayesian inference

- procedure for sampling joint distribution: $P(\theta, \{\mathbf{x}\} \mid \{\mathbf{r}\}, \{c\})$
- 1) sample latents: $\mathbf{x}^{(i)} \sim p(\mathbf{x} \mid \mathbf{r}, c, \theta^{(i)})$ conditional over latents
- 2) sample parameters: $\theta^{(i+1)} \sim p(\theta \mid \mathbf{r}, c, \mathbf{x}^{(i)})$ conditional over parameters

Fitting Latent Variable Models

I. Sampling (“MCMC”) - fully Bayesian inference

- procedure for sampling joint distribution: $P(\theta, \{\mathbf{x}\} \mid \{\mathbf{r}\}, \{c\})$
- 1) sample latents: $\mathbf{x}^{(i)} \sim p(\mathbf{x} \mid \mathbf{r}, c, \theta^{(i)})$ conditional over latents
- 2) sample parameters: $\theta^{(i+1)} \sim p(\theta \mid \mathbf{r}, c, \mathbf{x}^{(i)})$ conditional over parameters

2. Expectation maximization (EM)

Alternate updating parameters and posterior over latents.

Fitting Latent Variable Models

I. Sampling (“MCMC”) - fully Bayesian inference

- procedure for sampling joint distribution: $P(\theta, \{\mathbf{x}\} \mid \{\mathbf{r}\}, \{c\})$

1) sample latents: $\mathbf{x}^{(i)} \sim p(\mathbf{x} \mid \mathbf{r}, c, \theta^{(i)})$ conditional over latents

2) sample parameters: $\theta^{(i+1)} \sim p(\theta \mid \mathbf{r}, c, \mathbf{x}^{(i)})$ conditional over parameters

2. Expectation maximization (EM)

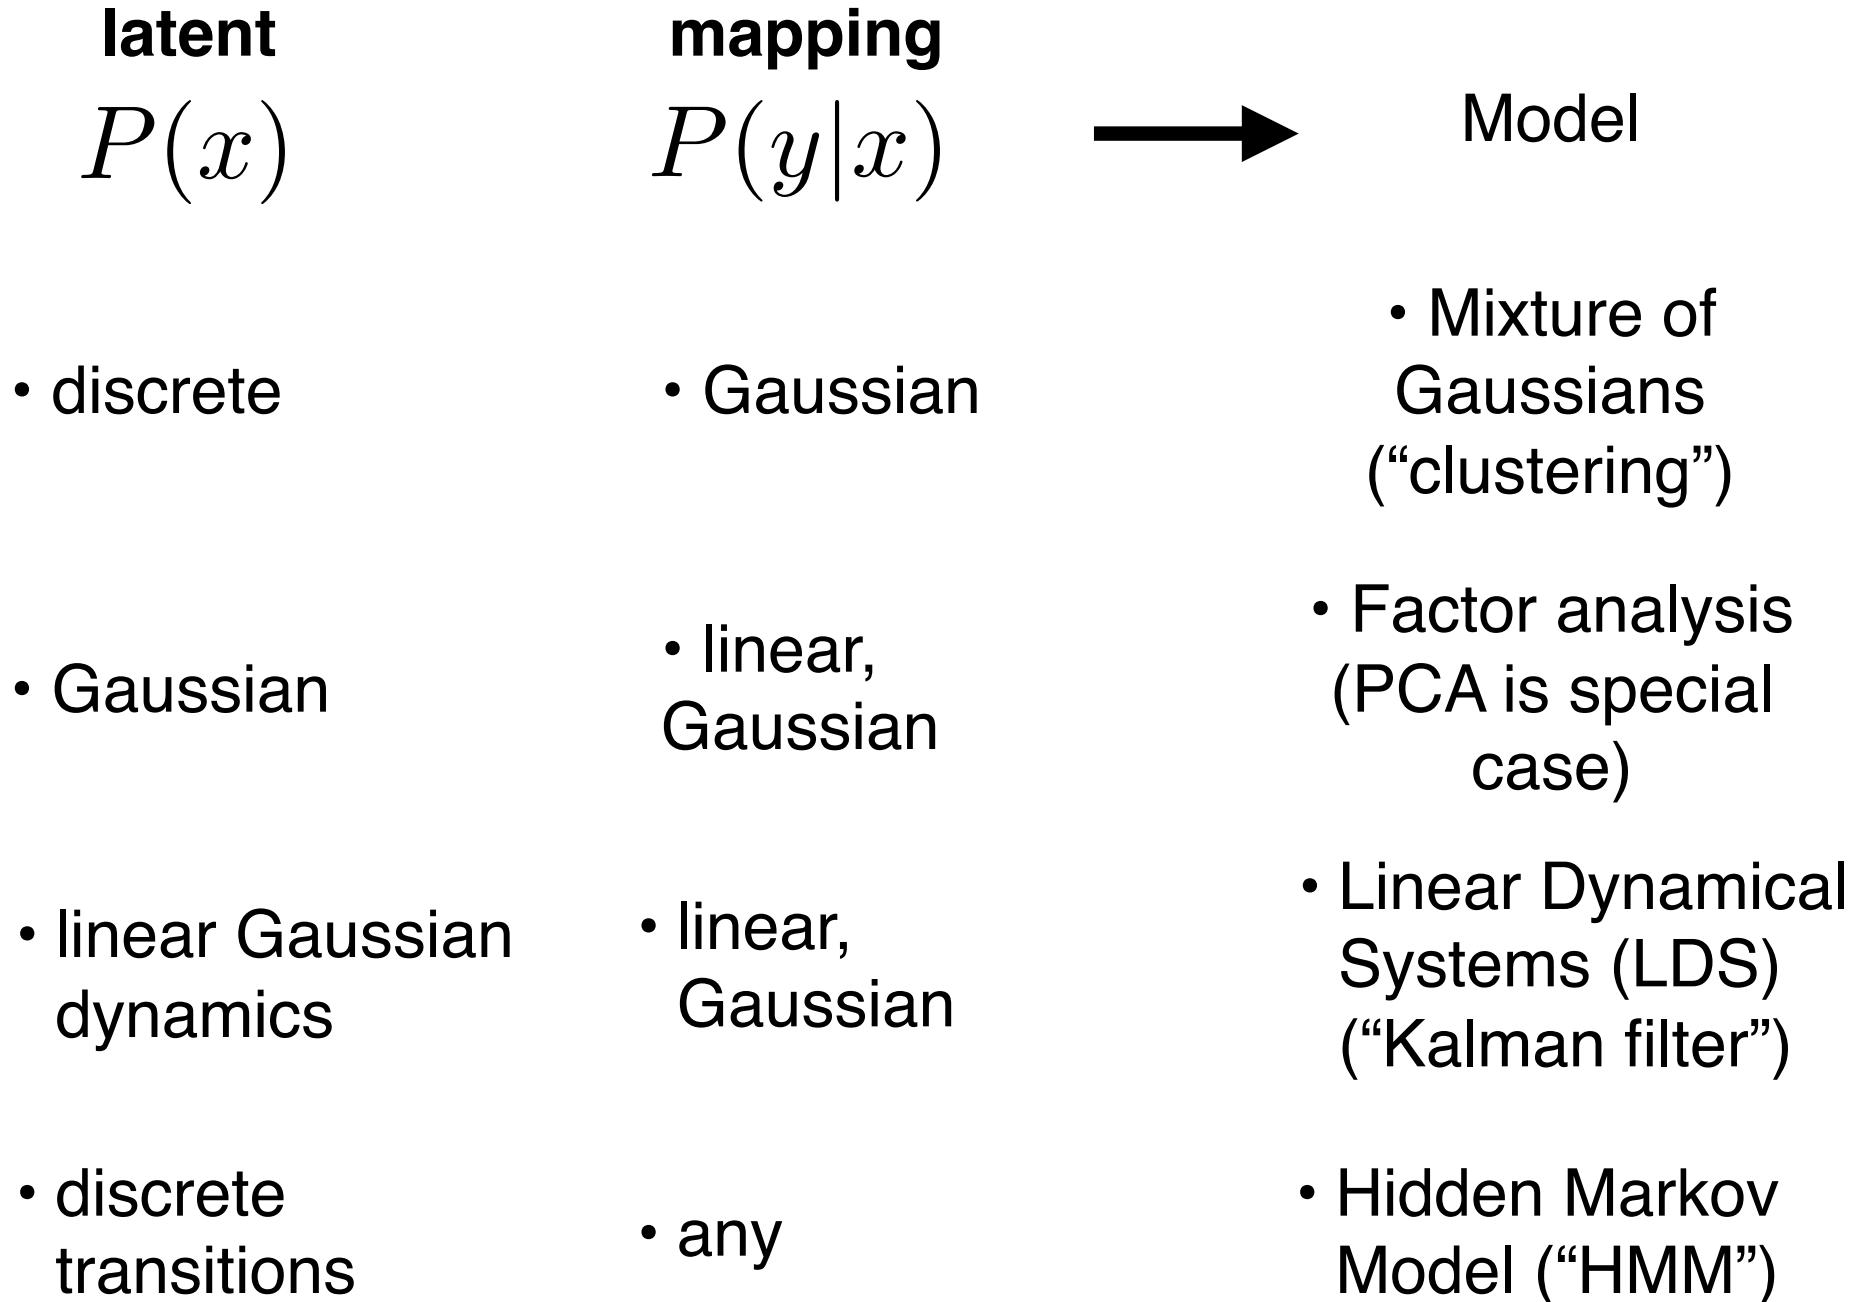
Alternate updating parameters and posterior over latents.

3. Variational inference

Optimize a lower bound on posterior over parameters

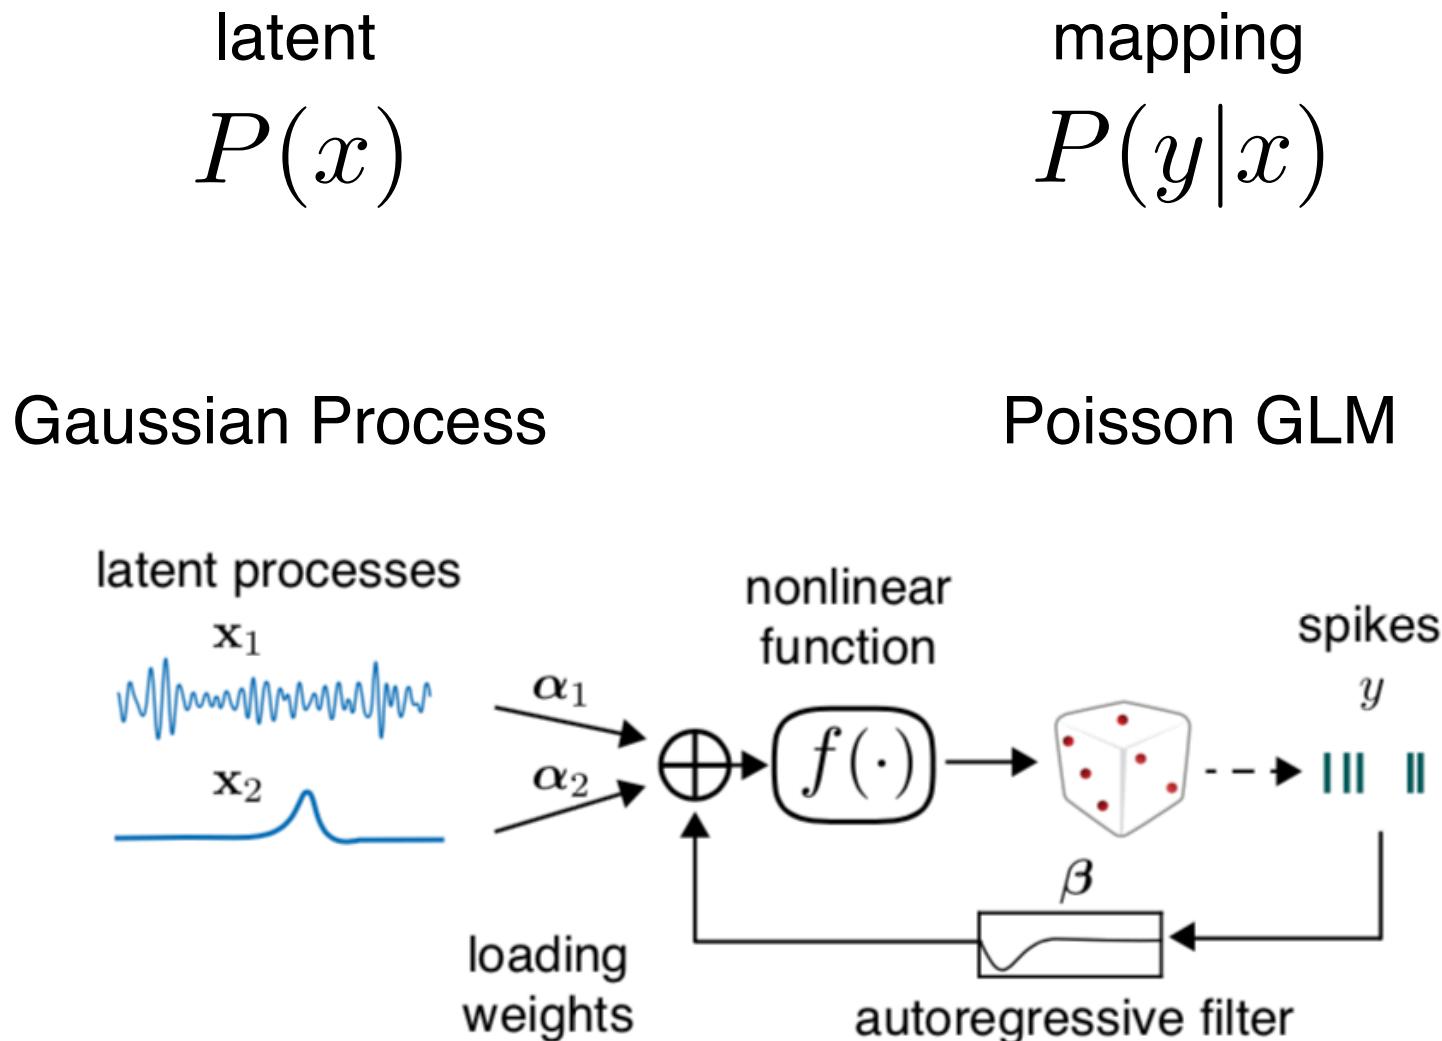
Easy with modern probabilistic programming languages
(STAN, Edward)

Latent Variable models are defined by two quantities:



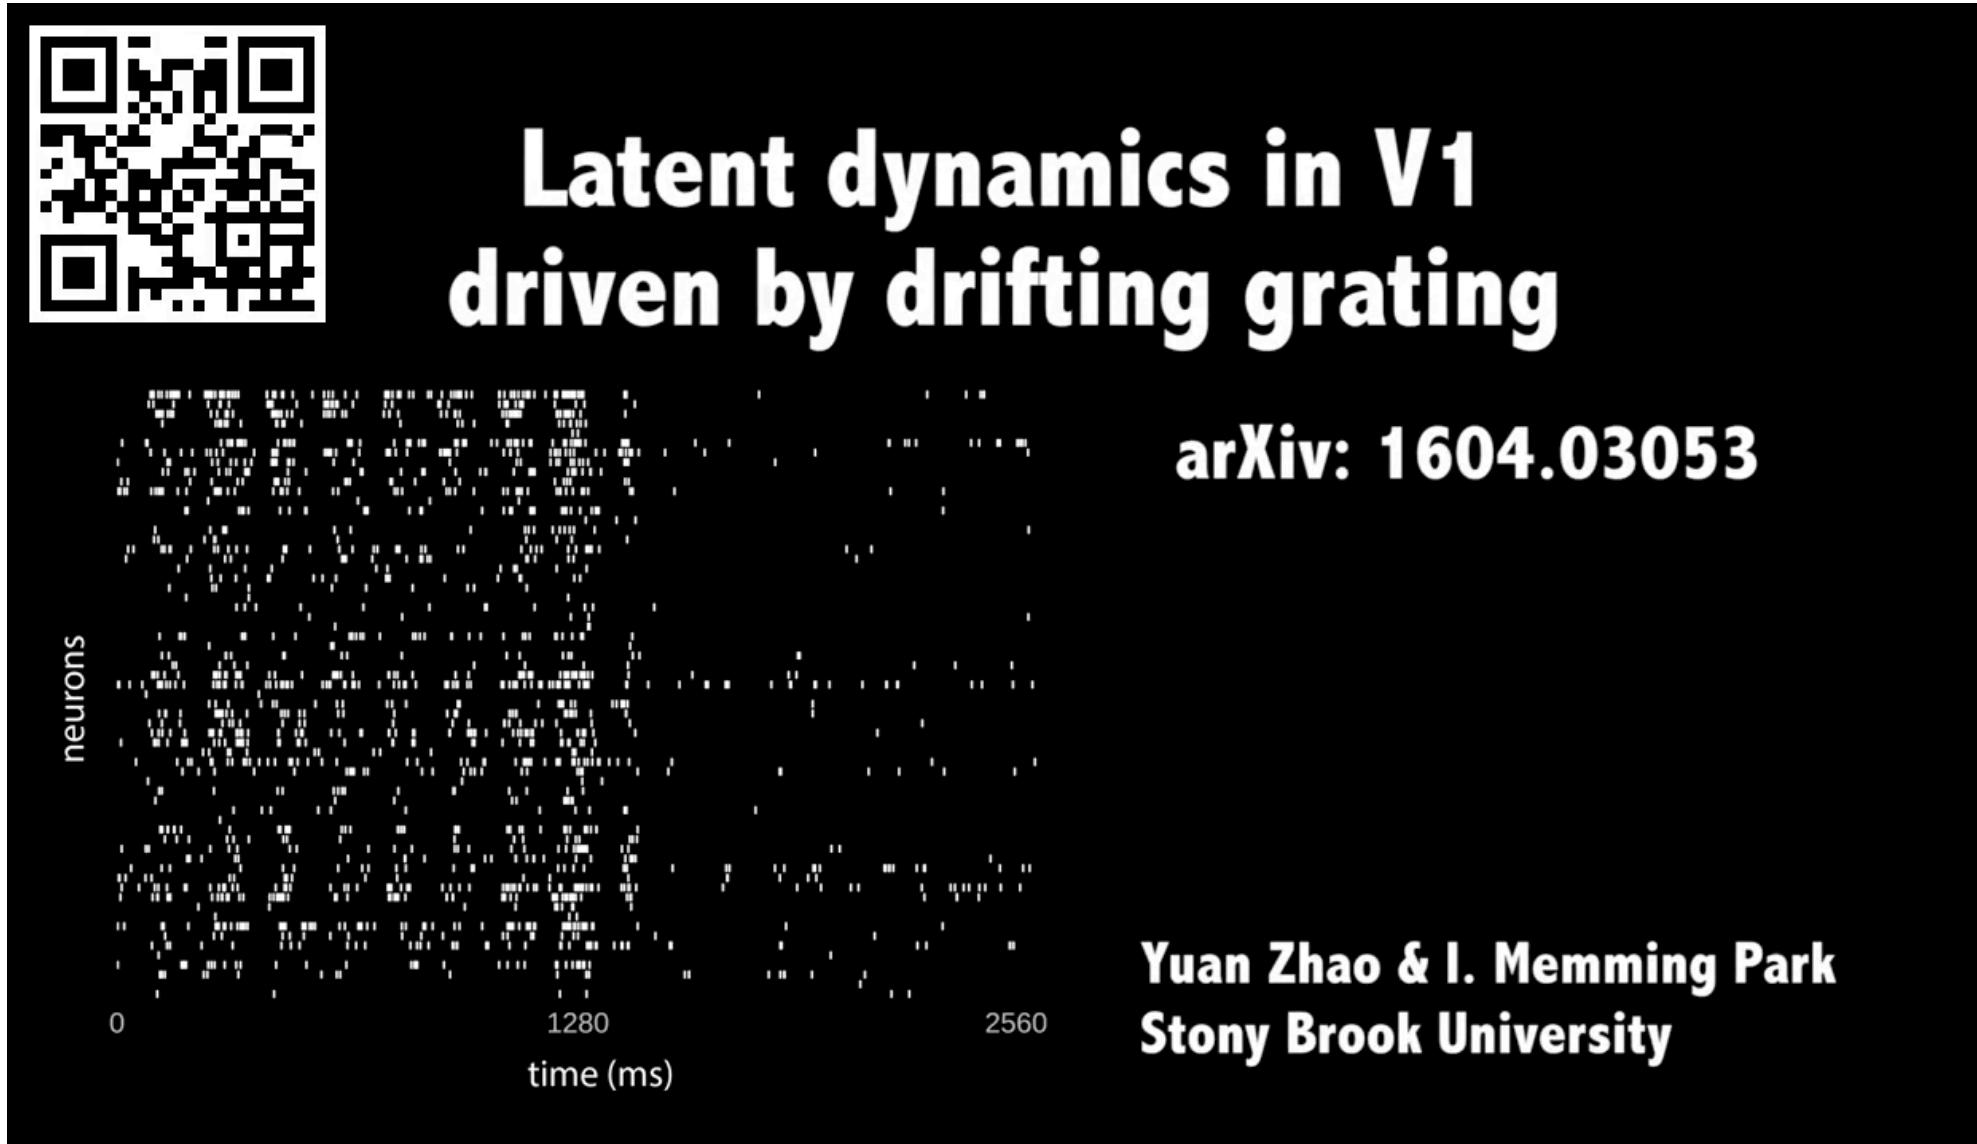
variational latent Gaussian process (vLGP)

[Zhao & Park 2016]



variational latent Gaussian process (vLGP)

- 63 simultaneously-recorded V1 neurons [Graf et al 2011]
- stimuli: drifting sinusoidal gratings



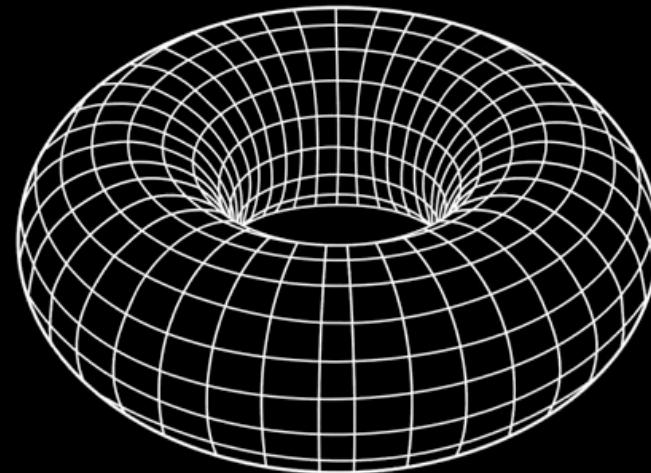
<https://www.youtube.com/watch?v=CrY5AfNH1ik>

variational latent Gaussian process (vLGP)

- 63 simultaneously-recorded V1 neurons [Graf et al 2011]
- stimuli: drifting sinusoidal gratings



**Latent trajectories
for 36 directions
form a torus**



Summary

- descriptive statistical “encoding” models
- seek to capture structure in data
- formal tools for comparing models
- encoding and decoding analyses via Bayes rule
- models are modular, easy to build /extend / generalize

Big Picture

- large-scale recording technology advancing rapidly
- lots of interesting structure in high-D neural data
- big opportunities in computational / statistical for developing new methods and models to find / exploit this structure!